

OpenCPI

Component Development Guide

Revision History

Revision	Description of Change	Date
1.01	Creation	2010-06-21
1.02	Add ocписca information and HDL application information	2010-08-05
1.03	Add more detail to HDL building, and general editorial improvements	2011-03-28
1.1	Editing, platform and device aspects of ocписgen, HDL details	2011-08-01
1.2	Update for latest HDL details	2013-03-12
1.3	Rename document, add VHDL coding details and ocписhdl utility	2013-06-15
1.4	Add new and complete assembly and container info, and add two new ocписhdl functions	2013-11-06
1.5	HDL coding practices, more HDL platform information	2013-12-10
1.6	Convert to ODT format, add more clarity for parameter properties and readonly properties	2014-03-31
1.7	Use standard template., minor updates, add generic OWD content.	2015-02-27
1.8	Incorporate the authoring model reference content into this document	2015-10-31
1.9	New section on projects, removed the HDL content to a different doc, removed authors cleaned up and refreshed all content.	2016-01-11
1.95	Processed additional clarity issues raised	2016-04-12
2.0	Clarify project dependencies in 13.5	2016-05-18
2.1	Update for 2017.Q2, major ocписdev update, minor unit test update.	2017-08-06
2.2	Update for 2018.Q1, projects, registries, package-IDs, pass/fail+remote containers in tests	2018-02-26
2.3	Update for 2018.Q3, ocписdev run/show	2018-09-28
2.4	Update for release 1.5. version 2, EOF, property accessibility, make variable deprecation	2019-04-16

Table of Contents

1	References.....	5
2	Overview.....	6
3	Introduction to Worker Development.....	8
3.1	A Simple Example Worker.....	10
4	Authoring Models.....	12
4.1	Requirements for All Authoring Models.....	13
4.2	Control Plane Introduction.....	14
4.3	Software Execution Model.....	20
5	Protocol Specifications (in OPS XML files).....	24
5.1	Protocol Element as Top-level Element.....	25
5.2	Protocol Specification (OPS) Examples.....	27
5.3	Message Payloads on Data Ports.....	28
5.4	End-of-File (EOF) Indications on Data Ports.....	29
6	Component Specifications (typically in OCS XML files).....	31
6.1	ComponentSpec Top-level Element.....	33
6.2	Properties Element of ComponentSpec Elements.....	34
6.3	Property Element of ComponentSpec or Properties Elements.....	35
6.4	Accessibility Attributes.....	38
6.5	Port Element of ComponentSpec Elements.....	42
7	Property Value Syntax and Ranges.....	44
7.1	Values of Unsigned Integer Types: uchar, ushort, ulong, ulonglong.....	44
7.2	Values of Signed Integer Types: short, long, longlong.....	44
7.3	Values of the Type: char.....	44
7.4	Values of the Types: float and double.....	45
7.5	Values of the Type: bool.....	45
7.6	Values of the Type: string.....	45
7.7	Values in a Sequence Type.....	45
7.8	Values in an Array Type.....	45
7.9	Values in Multidimensional Types.....	45
7.10	Values in Struct Types.....	46
7.11	Expressions in Property Values.....	46
8	Worker Descriptions in OWD XML Files.....	48
8.1	XML Attributes of the Top-level XYZWorker Element in the OWD.....	49
8.2	Property and SpecProperty Child Elements in the OWD.....	53
8.3	Attributes Allowed in OWD SpecProperty Elements.....	55
8.4	Built-in Parameters of All Workers.....	56
8.5	Port Elements of XYZWorker Elements.....	57
9	Worker Build Configuration XML Files.....	58
9.1	Build Configurations.....	59
9.2	Parameter elements in the <worker>-build.xml file.....	61
9.3	Configuration Elements in the <worker>-build.xml File.....	62

10	Component Libraries.....	63
10.1	The Component Library Makefile.....	65
10.2	The Library.mk File.....	67
11	Developing Workers.....	68
11.1	Creating Workers.....	68
11.2	The Worker Makefile.....	70
11.3	Editing Workers.....	72
12	The Worker Source Files.....	73
12.1	How Parameter Value Settings Appear in Source Code.....	73
12.2	Building Workers.....	73
13	Unit Testing of Workers.....	75
13.1	The Phases of the Unit Test Process.....	76
13.2	Unit Test Concepts and Terminology.....	78
13.3	Unit Test Description XML File.....	79
13.4	Unit Test Makefile Contents.....	89
13.5	Preparing Unit Test Inputs.....	90
13.6	Preparing for Unit Test Output Verification.....	91
13.7	Off-line One-time Tasks Prior to Test Execution and Verification.....	92
13.8	Testing on Remote Systems.....	93
13.9	On-line Tasks for Test Execution and Verification.....	95
13.10	Summary of Make Goals and Variables.....	97
14	Developing OpenCPI Assets in <i>Projects</i>.....	98
14.1	Managing Project Assets.....	100
14.2	Package IDs.....	101
14.3	The Project Registry: How Projects Depend on and Find Each Other.....	104
14.4	Project Makefiles.....	106
14.5	The Project.mk File for Project-wide Variable Settings.....	108
14.6	Exporting a Project to be used Elsewhere.....	114
14.7	Using Other Projects that Exist Outside the Project Being Developed.....	115
15	The ocpidev Tool for Managing Assets.....	116
15.1	Create/Delete Assets.....	118
15.2	Build/Clean Assets.....	123
15.3	Show Assets.....	126
15.4	Run.....	127
15.5	Refresh.....	130
15.6	Utilization.....	131
15.7	Register/Unregister.....	133
15.8	Set/Unset.....	133
15.9	Assets Managed by ocpidev.....	135
16	Environment Variables used in Component Development.....	137
17	Tools Used in Component Development.....	140
18	Appendices.....	141
18.1	Appendix A - ocpidev JSON.....	142

1 References

This document also refers to concepts and definitions in other documents, but does not depend on them.

Table 1: References to Related Documents

Title	Published By	Link
OpenCPI Overview	OpenCPI	https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_Overview.pdf
OpenCPI RCC Development Guide	OpenCPI	https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_RCC_Development.pdf
OpenCPI HDL Development Guide	OpenCPI	https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_HDL_Development.pdf
OpenCPI Application Development Guide	OpenCPI	https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_Application_Development.pdf

The OpenCPI Application Development Guide has introductory material on XML and the OpenCPI conventions for using XML. It also contains the description of the syntax for property values and expressions that is also used in XML files described in this document.

2 Overview

This document describes how to create OpenCPI component implementations (*workers*) in a component library, so that they are available for OpenCPI application developers and users. It introduces a kit of tools to specify and develop OpenCPI workers in any supported authoring model language and API. It also describes how to create, build, and manage libraries of heterogeneous components where the components may have multiple implementations.

This document describes tools and processes to development component libraries in general. Other documents describe the process of developing workers for specific authoring models, which currently include Resource-Constrained C Language workers **RCC** (C and C++ workers for software targets), Hardware Description Language workers **HDL** (VHDL or Verilog workers for FPGAs), and OpenCL workers, **OCL** (OpenCL workers for GPUs).

References to the kit of tools, scripts, documents and libraries used for developing components and workers in libraries are part of the **OpenCPI Component Development Kit (CDK)**. The CDK is not an integrated development environment (IDE), but rather is a set of commands, make and shell level tools, and scripts that support the development process. The CDK relies on several conventional tools, including GNU Make, and other basic POSIX command-line tools.

The CDK also includes tools specific to OpenCPI that support heterogeneous code generation and component testing. These tools are usually used indirectly, using the provided makefile scripts to build component libraries of workers (heterogeneous implementations), and, when applicable, building workers in each of the available authoring models. OpenCPI's code generation significantly reduces the code that needs to be hand-written in implementing heterogeneous components, applications, and FPGA bitstreams.

The OpenCPI CDK relies on technology-specific compilers (e.g. gcc), synthesis (for FPGAs) and simulation tools (e.g. Xilinx XST and Isim, Altera Quartus, Modelsim etc.). These tools are a mix of open-source/free and commercially available products. Specific supported tools and versions are found in the **OpenCPI Installation Guide**.

Several key concepts are described in the following sections, followed by the development process for creating component libraries.

Component Specification: an XML file that describes a component in such a way that it may be implemented using different languages and APIs for different processing technologies and environments. It specifies the properties and ports of the component.

Protocol Specification: an XML file that describes the allowable data messages and payloads that are used for communication between components.

Property: (or *configuration property*) are writeable and/or readable values that enable configuration, control and monitoring of workers by control software at run time.

Port: an interface of a component that allows it to communicate with other components using a protocol. Ports are unidirectional: input or output, consumer or producer.

Authoring Model: one of several ways of creating component implementations in a specific language using a specific API between the component and its execution environment. Existing models described below are RCC, HDL and OCL.

Worker: a specific implementation of a *component specification*, with the source code written according to an *authoring model*.

Component Library: a collection of *component specifications* and *workers* that can be built, exported and installed to support applications.

Project: a work area in which to develop OpenCPI components, libraries, applications, and other platform and device oriented assets.

3 Introduction to Worker Development

This section introduces the aspects of the worker development process that are common across all types of workers and authoring models. There are separate documents for each authoring model which describe their respective aspects in more detail, including languages and APIs. After this section introduces the general development process, following sections provide details for the contents of the various directories and files that are involved.

A worker is developed in its own directory, based on a component specification that typically exists in a file elsewhere. The component specification is the basis for multiple potential alternative implementations (workers). A component specification is an XML file called an **OpenCPI Component Specification (OCS)**, abbreviated as **spec** file, is described in detail in the [Component Specifications](#) section. The spec file also typically references one or more **OpenCPI Protocol Specification (OPS) files**, defined in the [Protocol Specification](#) section, to indicate the types of messages allowed to flow into and out of an implementation.

In addition to a worker having their own directory, they are typically developed in a component library (a collection of workers). The worker directories are then subdirectories of the component library's directory, and the OCS (and OPS) for a worker is typically found in the **specs** subdirectory of the component library.

Some authoring models (e.g. RCC) support creating a single binary artifact that implements multiple workers, but usually a single worker implementation is in its own subdirectory and when compiled, results in a single binary artifact.

The names of the worker directories have a suffix indicating the authoring model used for that implementation (e.g. `.rcc`, `.hdl`). For a component whose component specification file is named `xyz-spec.xml`, the RCC authoring model implementation of that component will typically be in a worker directory called `xyz.rcc`. The worker's directory must combine the name of the worker, before the “.”, and the authoring model used, after the “.”. A worker named `abc` written using the authoring model named `rcc`, would exist in a directory named `abc.rcc`.

The names of the spec file and the worker's directory do not have to match, but it is recommended and allows the use of more defaults to simplify the process. An HDL implementation of the component spec `xyz-spec.xml` would be in the subdirectory `xyz.hdl`. Note that these names “xyz” are not required to be the names that occur in the programming language source files (e.g. C, C++, Verilog, etc.), although that is usually the simplest.

An `xyz.test` directory, at the same level as the worker directories, should be created for unit tests common to all implementations of the xyz component's spec file. This means that tests in this directory apply to all workers that implement the same spec.

It is possible to have multiple workers implementing the same component specification, written in the same authoring model. In this case the worker names must be different and at least one of them must be different from the name implied by the component

specification. E.g. one might have `big_fast_xyz.rcc` and `small_slow_xyz.rcc`, both implementing the OCS in `xyz-spec.xml`.

Once an OCS is available, a worker can be created, usually in a library, by using the `ocpidev` tool, which creates a worker's directory and populates it with an initial version of several files that can be edited later. The `ocpidev` tool is described in the [ocpidev](#) section. The initial files created that are then edited as necessary include:

1. **OpenCPI Worker Description File** (the **OWD** XML file, `xyz.xml`)
2. Worker “make” file (named **Makefile**)
3. Worker initial source code file (named `xyz.<language-suffix>`)

All three of these file types have initial, automatically generated, skeletal contents that may be subsequently edited by the developer as required. Frequently only the source code files require editing. These files are described in detail below.

The OWD file is an XML file that describes the worker itself, by internally referring to an OCS and including implementation-specific attributes needed by the framework. The second file is a “make” file that describes how the worker is built, and the third is the initial source code file for the worker's actual logic.

The component specification (OCS) for the worker contains the description of the component's external behavior. These will exist in all implementations (workers) that reference the component specification. The OWD adds information about a particular implementation worker.

The OWD XML file has the name of the worker and the `.xml` suffix. The primary source code file has the name of the worker, with the typical suffix for the programming language used (`.c`, `.cc`, `.vhd` etc.). The primary source file for the `xyz.rcc` worker written in the C language would be `xyz.rcc/xyz.c`. A worker may also reference additional source files.

The worker building process invokes built-in scripts and makefiles, which automatically create and populate two types of subdirectories in the worker directory.

The first, called **gen**, holds automatically generated source code and XML files that are target-independent (architecture independent). The second type, with the name **target-`<platform>`** holds architecture-specific object/binary files usually generated for or by a compiler for a specific target. In this case **`<target>`** is the name of the compilation target being built, such as `centos6` for CentOS 6 Linux running on a 64-bit x86 processor. Both types of directories are files resulting from the build process and are removed by `make clean`, as they are always regenerated and should never manually edited. More details about these targets is in the [Developing Workers](#) section.

In the sections below, a simple example will be given, followed by a detailed description of the component specification files (OCS), followed by the three types of worker files just introduced.

3.1 A Simple Example Worker

Here is a simple example of a software worker written in C++. The component's function is to add a constant, called the `biasValue`, to each unsigned 32 bit integer at its input, and put the value at its output, one message at a time. The component specification XML file, OCS file, is named "`specs/bias-spec.xml`", and contains

```
<ComponentSpec>
  <property name='biasValue' writable='true' type='ulong' />
  <port name='in' protocol='u32-proto' />
  <port name='out' producer='true' protocol='u32-proto' />
</ComponentSpec>
```

The protocol specification XML file, OPS file, indicated by the `protocol` attributes in the OCS, would be found in the file `specs/u32-proto.xml`, and contains:

```
<Protocol>
  <Operation name='info'>
    <Argument name='values' type='ulong' sequenceLength='0' />
  </Operation>
</Protocol>
```

For the `bias.rcc` worker as created by `ocpidev`, which implements the above spec using the C++ language, the OWD XML file is named `bias.rcc/bias.xml`, and contains:

```
<RccWorker language='c++' spec='bias-spec'>
</RccWorker>
```

This OWD indicates that the authoring model is RCC, the spec is `bias-spec`, and the language is C++. The source file that implements this `bias.rcc` worker, simplified without header files or error checking, is in the file named `bias.cc`, and contains:

```
class BiasWorker : public BiasWorkerBase {
  RCCResult run(bool /*timedout*/) {
    size_t length          = in.info().values().size();
    const uint32_t *inData  = in.info().values().data();
    uint32_t *outData       = out.info().values().data();

    for (unsigned n = length; n; n--)
      *outData++ = *inData++ + properties().biasValue;
    out.info().values().resize(length);
    out.setOpCode(in.opCode());
    return length ? RCC_ADVANCE : RCC_ADVANCE_DONE;
  }
};
```

The `Makefile` in the worker's directory, automatically generated, would be:

```
include $(OCPI_CDK_DIR)/include/worker.mk
```

If the worker was written to the HDL model, in the VHDL language, its OWD would be:

```
<HdlWorker language='vhdl' spec='bias-spec' />
</HdlWorker>
```

For a detailed explanation for using HDL workers see the ***HDL Development Guide*** document for VHDL examples.

4 Authoring Models

This section introduces the concept of an **OpenCPI Authoring Model**, and defines aspects common to all authoring models. It specifies the concepts, lifecycle states and related operations, and XML metadata used and manipulated by OpenCPI tools and OpenCPI component developers.

The use of the term component is to encompass the functionality and abstract interface aspects of a model. The term worker is meant for particular implementation of a component written, authored, using a programming language source code.

The definition of a authoring model can be casually referred to as a way to write a worker. A key goal is to support different processing technologies available such as **General Purpose Processors (GPPs)**, **Field-Programmable Gate Arrays (FPGAs)**, **Digital Signal Processors (DSPs)**, or **Graphical Processing Units (GPUs)**.

Since there is no one language, or API, that allows all these processing technologies to be utilized with efficiency and utilization comparable to their native languages and tool environments, we define a set of authoring models that achieve native efficiency with sufficient commonality with other models to be able to:

- Implement an OpenCPI worker for a class of processors in a language that is efficient and natural to users of such a processor
- Be able to switch, replace, the authoring model and processing technology used for a particular component in a component-based OpenCPI application without affecting the other components of the application.
- Combine workers, component implementations, into an application using a multiplicity of authoring models and processing technologies.

An OpenCPI Authoring model consists of these specifications:

- An XML document, structure/schema/definition, to describe the aspects of the implementation that are specific to the authoring model being used and needed by tools and runtime infrastructure software.
- Three sets of programming language interfaces used for interactions between the worker itself and its environment:
 1. Control and configuration interfaces for run-time lifecycle control and configuration, referred to as the **control plane**.
 2. Data passing interfaces used for workers to consumer/produce data from/to other workers in the application (of whatever model on whatever processor), referred to as the **data plane**.
 3. Local service interfaces used by the worker to obtain various services locally available on the processor on which the worker is running.
- Each authoring model also specifies how a worker is built (compiled, synthesized, linked) and packaged, to be ready for execution in an application.

4.1 Requirements for All Authoring Models

- Enable/support well-defined data plane interoperability with other authoring models
- Define its OpenCPI Worker Description (OWD) XML format.
- Define programming language interfaces for control, data, and local services.
- Must define the packaging for delivering ready-to-execute workers.

The currently supported authoring models are:

RCC (for **R**esource-**C**onstrained **C**-language) is an authoring model used in the C or C++ language workers that execute on general purposes processors (GPPs). The C language model is a lean model well-suited to small resource-constrained processors such as embedded CPUs, DSPs or micro-controllers. The C++ variant is more powerful and more compact, carries a slightly higher resource footprint, and of course requires a C++ compiler. Developing workers according to the RCC authoring model (either C or C++) is fully described in the **RCC Development Guide**.

HDL (for **H**ardware **D**escription **L**anguage) is an authoring model using the VHDL (and less-supported Verilog) languages and is targeted at FPGAs. Developing workers according to the HDL authoring model is fully described in the **HDL Development Guide**.

OCL (for **O**pen**CL**) is an upcoming authoring model using the OpenCL (C subset/superset) language targeting graphics processors. It is fully described in the **OCL Development Guide**. This support is experimental as of the current release.

4.2 Control Plane Introduction

The material in this introduction is common to all authoring models. We use the term **control software** to describe software that launches and controls OpenCPI applications. This is either the standard utility program, `ocpirun`, or custom C++ programs that perform the same function embedded inside them. Such custom programs use the **Application Control Interface**, an application launch and control API described in the **OpenCPI Application Development Guide**.

We use the term **Control Plane** to encompass the various aspects of how control software, usually running in a centralized host processing environment, can control worker instances at runtime. The entity that is doing the controlling (or managing) is the *control application*, or simply *control software*. The control software uses all controllable worker instances the same, regardless of where they are running, on what type of processing technology, and with what authoring model they were written.

While control software sees a uniform view of how to control workers, each authoring model defines how this is accomplished from the point of view of the worker itself. In particular, each authoring model defines how the two key aspects of control are made visible to the worker's source code: *LifeCycle control* and *Configuration Property access*. The documents describing each authoring model give additional interface details of these interactions, but they all follow a common pattern which is defined here.

4.2.1 LifeCycle State Model

Most component-based systems have an explicit lifecycle state model, where workers are instantiated and then managed, according to a lifecycle state machine. Normally all components in the application are managed together and they all progress through the lifecycle together. However, there are cases where control software must control (start/stop etc.) some components in the application different than others.

The LifeCycle model is defined by the *control states* each worker may be in, and *control operations* which generally change the state a worker is in, effecting a state transition. The possible states and transitions are shown in the following diagram.

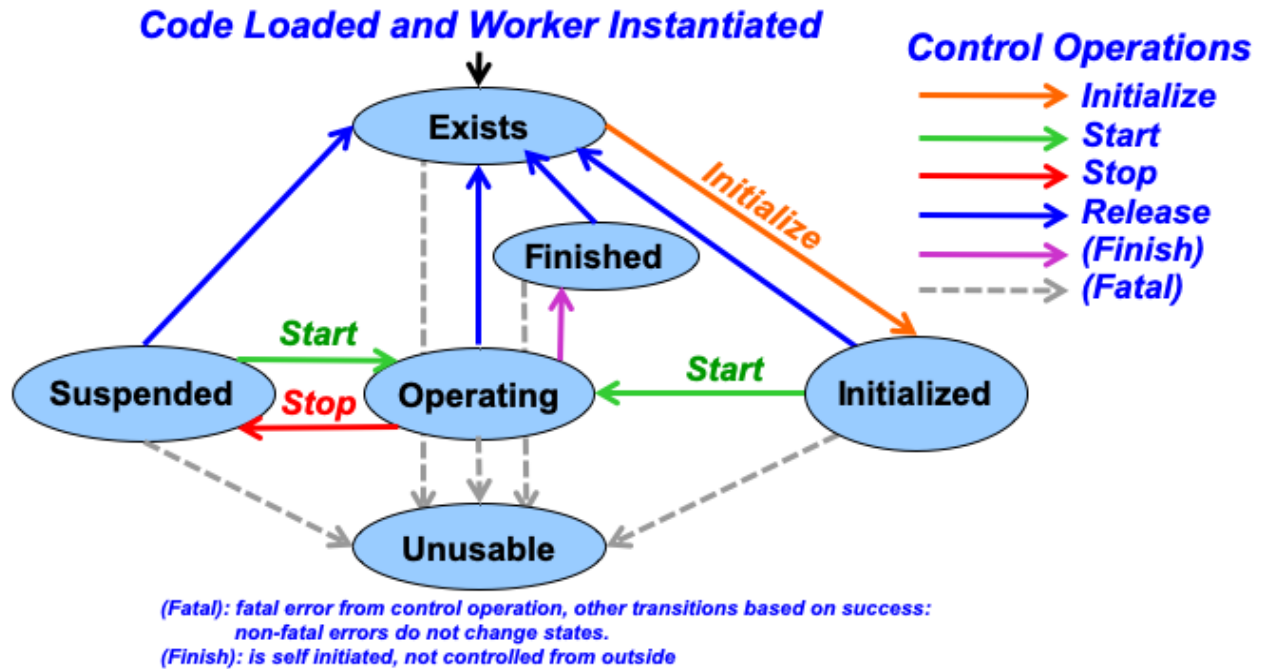


Diagram 1: Control States and Operations

4.2.2 LifeCycle Control Operations

Control operations have default implementations that only perform the state transitions: workers only need to implement the operations that have custom behavior. A good example is the *initialize* control operation. If the worker has no runtime initialization to perform, it can have no implementation of this operation and not even have an empty “stub”. Each authoring model describes which control operations must have implementations.

Control operations can have two error types: transient or fatal. Transient errors imply that no state change occurred and the operation can be retried. Fatal errors imply that the worker instance has become unusable and needs to be reloaded.

Control software is required to issue control operations correctly, in sequence, so workers can avoid checking for valid states and transitions. State descriptions are listed in the table below:

Table 2: Control States

Control State	Previous State(s)	Allowable Operations	Description
Exists	<i>Initial state or all (except unusable), after release</i>	Initialize	Follows instantiation or a successful release . Worker is loaded (if necessary), not fully initialized, with no properties valid, and property access not allowed.
Initialized	Exists	Start, Release	Initialization complete. Follows initialize . In a stable state “ready to start doing work”, not “operational”. Properties can be read and/or written.
Operating	Initialized, Suspended	Stop, (done), Release	Member doing normal work using properties and data at ports. Properties can be read or written. Follows start .
Suspended	Operating	(Re)start, release	Member not operating, will not produce or consumer at data ports. Properties can be read or written. Can be resumed via start .
Finished	Operating	Release	Member is finished and will not produce or consumer at data interfaces. Properties can be read or written. Entered autonomously.
Unusable	All (fatal errors)	none	Fatal error state, may not be ever reusable without reloading (container-dependent).

Configuration Properties

Configuration Properties are specified in the OCS or OWD, and are writeable and/or readable values that enable control and monitoring of workers by control software. They are logically the knobs and meters of the worker's “control panel”. All authoring models provide an interface enabling workers to access (read and write) these values.

Some authoring models define a flat/linear configuration address space where the configuration properties are accessed by accessing this memory space, roughly as a data structure or register file.

The component specification for the worker contains the description of the configuration properties that are part of the component’s external behavior. These will exist in all implementations (workers) that reference the component specification. However, each worker may also *add* to this set of configuration properties and define **implementation-specific** configuration properties. These can be useful for implementation debugging and testing, or in some cases to allow applications to configure properties specific to a particular implementation.

Each configuration property is defined with a name and data type from the data types listed in the table below. Each configuration property may vary in length (i.e. be strings or sequences) but must specify a maximum length. This enables components and workers to be compliant with a variety of component system standards, and enables authoring models for resource-limited embedded technologies.

Data types for configuration properties are based on the **scalar types** listed in the following table. A **property data type** can be one of the scalar types or a **structure** with typed members that are **property data types**. A property data type may have array dimensions, and also (after any array dimensions are applied) be defined as a variable length, single dimension **sequence**. This recursive definition allows for complex types such as: *a sequence of a two dimensional arrays of structures containing members that are themselves arrays or sequences or structures*.

Table 3: Scalar Data Types for Properties

Scalar Data Type	Support for Unsigned Version	Data Size	Notes
<i>bool</i>	N/A	8 bits	
<i>char</i>	Yes	8 bits	Unsigned version is uchar .
<i>double</i>	N/A	64 bits	Consistent with IEEE floating-point types
<i>enum</i>	N/A	32 bits	Types are represented by ulong values, but are associated with string names.
<i>float</i>	N/A	32 bits	Consistent with IEEE floating-point types
<i>long</i>	Yes	32 bits	Unsigned version is ulong .
<i>longlong</i>	Yes	64 bits	Unsigned version is ulonglong
<i>short</i>	Yes	16 bits	Unsigned version is ushort
<i>string</i>	N/A	fixed	Null terminated with a defined max length

When a property's type is a multidimensional **array** of the above types, the number of dimensions is fixed, and the length in each dimension is fixed.

When a property's type is a variable length **sequence** of the above types (or arrays of above types), it has a current length (number of valid values present), but is still required to have a defined maximum possible length (capacity). Sequences may have a current length of zero or any amount up to the specified maximum possible length. E.g. if a sequence is defined with a maximum length of 4, it means that it may hold zero, 1, 2, 3 or 4 values. Space is always reserved for the maximum number of values, but the current length is also recorded in the sequence and is set whenever a new value for the sequence is set.

Beyond data type, the configuration properties also have ***accessibility attributes*** indicating whether the value can only be:

- set at initialization time
- set any time during execution
- never set (read-only)

Similarly, a configuration property value can be described as volatile where the value may be changed by the worker during execution, or statically readable and will not change unless written by control software. These accessibility types are described in detail in the [Accessibility Attributes](#) section.

An OWD file allows additional properties to be defined unique to the worker, beyond those specified in the OCS. Additionally, the OWD may *add* to the accessibility of an OCS property. E.g., for debug purposes, the OWD may make a property readable that was not readable in the OCS. The accessibility added by the OWD results in the implementation having a superset of what was described in (and required by) the OCS.

4.2.3 Properties that are build-time parameters

While the OCS specifies properties and their initialization-time and run-time accessibility, an OWD can further declare that a property is a compile time parameter in this worker. This is not allowed for properties declared as writable at runtime, either in the OCS or OWD. When an OWD declares properties to be parameters, this means that the worker must be built compiled, synthesized, elaborated for specific values of such properties. This has three implications:

- An application can only use the worker if it is built for a property value that matches the value requested for an “initial” property in the application.
- Binary component libraries may have multiple binary artifacts for the same worker, but with different combinations of parameter values.
- The component developer must decide which combinations of parameter values to build, in order to make alternative settings of such parameters available.

This parameter feature allows implementations to have compile time optimization for certain values. It also allows a single worker's source code module to be optimized for different values.

Parameter property values are applied to the build process as per the language of the authoring model: e.g. by preprocessor symbol definitions for C, ***static const*** values for C++, generics for VHDL, and parameters for Verilog. A framework generated (built-in) parameter property is the `ocpi_debug` Boolean property that specifies the typical debug build vs. production build. Parameter properties are described in more detail in the [Parameter Attribute](#) OWD section.

Each authoring model specifies how properties, at runtime and compile time, appear to the worker code. In addition it specifies how the worker may read property values written by control software, and write values that will be read by control software.

Finally, different authoring models define how workers know when control software actively reads or writes these values.

4.3 Software Execution Model

The material in this section applies to most software-based authoring models (e.g. RCC and OCL). It does not apply to the HDL (FPGA) authoring model.

Execution of the model is based on a construct called **Containers**. Containers supply threads and execute software-based workers. This eliminates the need for workers to create or manage threads. This reduces the complexity of the worker code, eliminates any requirement to support a threaded API capability, and allows the container to determine the level of multithreading that is needed. The authoring model defines how the threading is provided and is detailed in the individual sections for each.

Execution of a worker occurs when the worker is transitioned into the operating control state. Workers are only executed when either a combination of its ports are ready, (port readiness), or an amount of time has elapsed. The combination of port readiness and the passage of time is referred to as the worker's **run condition**: the condition under which it should be run.

Every worker has an entry point called its **run method**. When a worker's run condition is satisfied, determined by its container, its run method is entered. Worker execution is a series of "runs" initiated by the container. The run method cannot block, but returns after doing some work, allowing the container itself to determine when the worker should be entered again: when its run condition is once again satisfied.

The container calls the worker's run method when the worker's run condition is satisfied. Run conditions are satisfied based on the availability of its input buffers, with data, or output buffers, with space, at a worker's ports or the passage of time/ The worker's run method executes some processing tasks and may:

- process some available messages at some of its input ports
- produce messages at some of its output ports
- indicate when messages are consumed as input or produced as output
- make any changes to its run condition

It then returns control to its container. Workers never block. The container conveys the messages in buffers between colocated workers as well as into and out of the container as required by the application assembly's connections.

The container determines when the worker should run, supplies it with buffers full of input messages, and buffers into which output messages may be placed.

4.3.1 Run Conditions

Workers declare a run condition which tells the container under what conditions the worker should run. The container evaluates the run conditions of all workers and runs them as resources and priorities allow.

The run condition object contains two aspects: *port readiness* and *time*. The worker is run when its port readiness requirements are satisfied, or a specified amount of time has passed. Either or both aspects can be specified.

While the worker is in the operating state, port readiness means that buffers are available at that port to be used by the worker. Input ports have available buffers when there is message data present that has not yet been consumed by the worker. Output ports are ready when buffers are available into which they may place new data. I.e. input ports are ready when the worker has data to consume, and output ports are ready when the worker has room to produce new data into a buffer. Workers may partially consume or partially produce messages in any given run.

This port readiness model implements simple data driven execution: code is run when data can flow. The **default run condition** specifies that the worker should run if data is available at *all* input ports and space is available at *all* output ports (or conversely, there are no ports that are not ready). Note that this default, for workers with *no* ports, means they are always ready to run.

The time aspect of run conditions, indicated by the worker, specifies the desired maximum time between invocations of the worker's run method. If no port readiness is also specified, this simply indicates periodic execution. If the time aspect is specified *with* port readiness, it indicates that execution should take place when either the indicated port readiness conditions are satisfied *—or—* the indicated amount of time has passed since the worker's run method was last entered.

The default time aspect of run conditions is: no such maximum time at all. In this case time passing does not affect worker execution, only port readiness.

A worker may change its run condition at any time during the execution of its run method by passing a new run condition to the container, to be considered after the run method returns to the container.

In summary, run conditions specify a combination of data-driven and time-driven execution. Most workers use one or the other, but both can be used together. The defaults allow most workers to never have to indicate any run condition at all.

4.3.2 Sending or receiving messages via ports

The worker indicates data flow to the container under two conditions. The first occurs when the worker has consumed the message in an input buffer at a port. Notification of which allows that buffer to be released and reused. The second happens when the worker has finished placing a message in an output buffer at a port. This allows the message in the buffer to be sent on.

4.3.3 Buffer management

The container provides and manages all buffers and provides references to buffers to the worker. Input ports operate by the container providing buffers to the worker filled with incoming messages. Output ports operate by the container providing buffers for the worker to fill with messages before being sent. Output buffers are either:

- obtained for a specific output port (since they may be in a special memory or pool specific to a particular output hardware path), *or*
- originally obtained from an input port and passed to output ports, with no copying by worker code.

Workers may modify data in input buffers, allowing input buffers to be used for temporary storage, to reduce overall memory requirements. When reuse occurs, the buffers must be annotated in the worker description XML. This ensures the container does not share the buffer with another consumer of the same data.

Several more advanced buffer management requirements are supported for certain situations:

- To support sliding window algorithms, workers are allowed to retain ownership of previous buffers by not releasing them while new ones are requested; i.e. allow explicit in order input buffer release, not just the most recent buffer obtained. The worker must still release the buffers in the order received.
- To support zero copy from input ports to output ports, workers are allowed to send a buffer obtained from an input port to an output port. This method does not require an empty current buffer to fill on the output port. Such buffers must be sent, or released, in the order received. This avoids copying data from input buffers to output buffers.

The features list previously are only needed in certain cases, and can be ignored for most simple workers. To support these more advanced modes, non-blocking interfaces for explicitly releasing, sending, and requesting buffers are available.

There are four operations performed with message buffers. These provide the basis for specific non-blocking functions in the APIs defined for each authoring model.

- **Request** that a new buffer be made available. For an input port, it will be filled by the container with a new input message. For an output port, it is to be filled by the worker with a new output message. In both cases the ownership of the buffer passes from container to worker when it becomes available. The new buffer may or may not be immediately available based on this request.
- **Release** a buffer to be reused, with its contents discarded. The ownership passes from worker to container. Input buffers must be released (or sent) in the order received, i.e. ownership of input buffers must be passed from worker to container in the order that ownership was given from container to worker.
- **Send** (enqueue) a buffer on an output port, to be automatically released after the data is sent. The ownership passes from worker to container. If the buffer was originally obtained from an input port, it must be sent or released in the order received.
- **Take** the current buffer from an input port such that it is no longer the current buffer of the port, but ownership is retained by the worker. This allows new input buffers to be made available while the worker holds on to previous buffers. A take implies a request for new buffers. This function allows workers to use previous buffers to hold history data for algorithms such as sliding window or moving average, without allocating any additional storage.

The concept of the **current buffer** of a port supports a model for workers that have no need to be aware of buffer management. A port is **ready** if it has a current buffer. A current buffer on an input port is available to read data from. A current buffer on an

output port is available to write data into. The concept of **advancing** a port, is simply a combination of releasing (input) or sending (output) the current buffer of the port, and requesting a new buffer to be made available on that port, to become the current buffer when it becomes available in the future:

- `advance = release_or_send + request`

Simple workers, using the default run condition, wait for all ports to be ready, process input buffers into output buffers, advance input and output ports, and return.

Worker APIs defined by the authoring model are designed to make this common pattern as simple as possible. Workers are run when ports are ready, and they advance ports after processing messages in current buffers.

5 Protocol Specifications (in OPS XML files)

An **OpenCPI Protocol Specification (OPS)**, describes, in one or more XML files, the set of messages that may flow between the ports of components. They are described separately from the OCS XML file as they are used by both sides of a connection. In a connection between component ports the specs of both components, in their Port elements, refer to the same OPS.

The OPS describes the set of messages defined in the protocol, as well as some top level attributes for the protocol.

In special cases the messages in a protocol are not specified individually, but rather a set of summary attributes is specified. This indicates the basic behavior of the ports using the protocol. The information is called a **protocol summary**. When messages are specified, protocol summary attributes are inferred from the messages. If protocol summary attributes are present when messages are specified, they override the attributes inferred from the message specifications.

As an example, a set of messages of different lengths and different payload formats might be bounded, having a maximum length, or unbounded, depending on whether any message has no maximum length. This **boundedness** attribute is normally inferred from the set of messages. Another example is the smallest unit of data in any message. If all messages in a protocol deal only with 64 bit integers, then the smallest unit of data for all messages is 8 bytes. This **minimum data granularity** attribute is inferred from examining all the messages specified for the protocol.

A protocol summary is the set of all summary attributes, whether inferred from the messages specified for the protocol, overridden by explicit values, or given directly when messages are not specified. The summary attributes are used by the OpenCPI code generation tools and runtime environment to determine certain behaviors.

In OPS files, messages are called **operations**, and fields of messages are called **arguments**. This is terminology based on the Remote Procedure Call (RPC), or Remote Method Invocation (RMI), model of communications. However, this concept does not apply to OpenCPI inter-component communications, as all communications are simply unidirectional connections conveying messages.

The term **opCode** is used to represent a zero-origin ordinal of operations within a protocol. In the runtime environment opCodes are used to indicate which operation of the protocol a given message represents. This if the opCode of a message is zero, then that message should be interpreted as the first operation in the protocol.

OPS files preferably carry the suffix `-prot.xml`, although `-protocol.xml`, `_protocol.xml` are also used.

5.1 Protocol Element as Top-level Element.

The **Protocol** element is a top-level element in a separate file whose name is the value of the **Protocol** attribute in a port element in an OCS. It specifies the message protocol used at a port. The protocol will likely be reused across a variety of components and interfaces since it specifies how two components talk to each other. The **Protocol** element has **Operation** subelements to indicate the different message types that may flow out of or into data ports using this protocol.

5.1.1 Name attribute of Protocol elements

When the **Protocol** element is the top-level element in a file, the optional name attribute is defaulted from the name of the file, with any **-prot**, **_prot**, **-protocol**, **_protocol**, **.xml** suffixes removed. Since protocols are usually defined in separate files, the names are usually not present in the XML and are derived from the file name.

5.1.2 Protocol Summary Attributes

These attributes are normally inferred from **Operation** elements in the protocol and are rarely used explicitly. They may be used to override the inferred values or they may be specified in the absence of **Operation** elements altogether. This table defines the protocol summary attributes normally inferred from examining **Operation** elements.

Table 4: Protocol Summary Attributes

Name	Type	Description
DataValueWidth	ulong	Size in bytes of smallest unit of data in any message
DataValueGranularity	ulong	Minimum number of data values in any message
DefaultBufferSize	ulong	Default buffer size for this protocol even if unbounded*
DiverseDataSizes	bool	Are there different size data values?
MaxMessageValues	ulong	Maximum number of data values in any message*
MinBufferSize	ulong	Minimum allowable buffer size for this.
MinMessageValues	ulong	Minimum number of data values in any message
NumberOfOpCodes	ulong	Number of message types
UnBounded	bool	Do any messages have unbounded length?
VariableMessageLength	bool	Can messages be different lengths?
ZeroLengthMessages	bool	Are any messages zero length?

* these attributes have no inferred value if protocol is unbounded.

Name	Type	Description
------	------	-------------

NumberOfOpCodes	ulong	Number of message types
DataValueWidth	ulong	Size in bytes of smallest unit of data in any message
DataValueGranularity	ulong	Minimum number of data values in any message.
ZeroLengthMessages	bool	Are any messages zero length?
MaxMessageValues	ulong	Maximum number of data values in any message, or zero if not bounded
VariableMessageLength	bool	Can messages be different lengths?
DiverseDataSizes	bool	Are there different size data values?
UnBounded	bool	Do any messages have unbounded length?
DefaultBufferSize	ulong	Buffer size for ports using this protocol even if protocol is unbounded. Default is zero when protocol is unbounded.

5.1.3 Operation element of Protocol elements

The term Operation is loosely associated with the analogous concept in RPC systems where the message is invoking an operation on a remote object. In the context of OpenCPI it simply describes one of the messages that is legal to send on a port with this protocol. It has two attributes and some number of argument child elements, which describe data fields in the message.

5.1.3.1 Name attribute of Operation elements

This string attribute is a case insensitive name of the operation/message within this protocol. It should be an appropriate identifier for programming languages.

5.1.3.2 Argument element of Operation elements

This child element indicates a data field in the message payload for the given operation. Its attributes are the same as a **Property** element and describe a configuration property with: **Name**, **Type**, **StringLength**, **ArrayLength**, **SequenceLength**, etc. If no **argument** elements are present under an Operation element, the operation defines messages with no data fields, referred to as a **Zero Length Message**. Argument elements are similar to **member** elements in **property** elements whose **Type** attribute is **struct**, but these arguments to an operation do not have to have bounded lengths. Here the **StringLength** attribute is not required for strings, and the **SequenceLength** attribute can be zero indicating no upper bound.

5.2 Protocol Specification (OPS) Examples

This protocol has one message type consisting of 1024 unsigned short values.

```
<Protocol
  <Operation Name="mess1">
    <Argument Name="val" Type="uShort" ArrayLength="1k">
  </Operation>
</Protocol>
```

5.3 Message Payloads on Data Ports

The message payload for each operation has a serialized format as a sequence of bytes that, when used in software, are laid out in byte-addressed memory. For example, if the operation element in a protocol contains:

```
<argument name='a1' type='uchar' />
<argument name='a2' type='ushort' arraylength='2' />
<argument name='a3' type='ulonglong' />
```

And the values of this payload are:

a1: 1, **a2:** {0x2345,0x6789}, **a3:** 0xfedcba9876543210

Then the byte sequence (with proper alignment, and encoded little-endian), would be:

Sequence # ►	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Contents (hex)	01	x	45	23	89	67	x	x	10	32	54	76	98	ba	dc	fe
Argument	a1		a2[0]		a2[1]				a3							
Contents	1		0x2345		0x6789				0xfedcba9876543210							

This layout and these values is the same for all types of workers in all (little endian) environments and over all data paths. The **x** values are padding for alignment.

Every numeric type is aligned for its size. Structure types are aligned according to the largest alignment requirement of any of its members. Sequences are aligned according to the alignment requirement of their members, and preceded by a 32 bit unsigned value indicating the number of elements. If the alignment requirement of the sequence is greater than 32 bits, padding is added after the length to align the first element, even if there are zero elements. Padding inserted to achieve alignment may be any value.

5.4 End-of-File (EOF) Indications on Data Ports

Regardless of protocol, the OpenCPI messaging system (its **data plane**) supports an **EOF** indicator in the data paths between workers of all types that indicates that no more messages/data will arrive at input ports or that no more messages/data will be produced on output ports. This indication is a persistent condition that will only be removed when the application and workers are restarted. It is not a message or a protocol operation.

Such an indication only happens when some worker asserts it on an output port, which then propagates it (down-stream) to the input port of the worker that the output is connected to. The most common usage of this EOF indication is when data is being read from a file (e.g. by a `file_read` component), and when that component encounters the end of the file it is reading. It typically then has three options:

- Assert an **EOF** on its output port after the last message/data is sent, and enter the **finished** control state.
- Do nothing on the port, and enter the **finished** control state.
- Start reading the file from the beginning again.

Thus in a typical execution of an application where end-of-data is a meaningful condition, the **EOF** will propagate through all workers to the final “sink” components (e.g. `file_write` or `capture`) that do not have output ports. In this way all workers know that the application is finished, and all can cleanly complete their work (perhaps flushing pipelines, closing output files etc.). In this way the EOF indication on the data plane is related to the action that a worker would take to enter the finished control state. I.e. the reception of EOF would commonly result in the worker entering that state.

For continuous applications where the concept of EOF is not meaningful, it will never be encountered by any worker since no worker will assert it on an output port.

The unit test framework described in [Unit Testing](#) makes use of EOFs to execute repeatable tests of components. Applications that process data sets of bounded size will also find the use of EOF to be convenient and useful, since it makes it clear to all components as well as the application as a whole that the application is “done”.

For most workers, the propagation of the EOF indication from its input port(s) to its output port(s) is automatically done without their involvement. Input data will simply stop arriving. However, some workers have more complex interactions with data ports, and perhaps more complex internal pipelining or state that requires special treatment when the last data arrives. They will recognize the EOF indications and perform final processing tasks accordingly. This latter type of worker declares “I will deal with EOFs”, (using the `WorkerEOF` attribute) and the OpenCPI infrastructure will not perform the automatic propagation.

Each authoring model has an interface for the worker to sense the EOF indication on input and assert it on output. These interfaces are described in the respective documents specific to the authoring models (e.g. RCC and HDL development guides).

The next section describes important issues of backward compatibility relating to EOF indications for workers implemented prior to release 1.5, which are still supported.

5.4.1 EOF Processing Prior to Release 1.5

Prior to release 1.5, the notion of EOF was overloaded with, and frequently confused with, the notion of a zero-length-message (a.k.a. ZLM). The pre-1.5 `file_read` component, upon encountering the end of the file it was reading, would (as a default), send a zero length message with the opcode of zero as an EOF indication.

This would directly conflict with protocols whose first operation happened to have no arguments. In general, users avoided the conflict by not defining protocols with the first operation having no arguments. And it was impossible to use the EOF concept with a “pure single event” protocol that only had one message with no arguments, which is useful for sending events (pulses) between components. These difficulties disappear with the introduction in release 1.5 of a cleanly separate EOF indication that has nothing to do with messages and protocols.

Workers written prior to release 1.5 that supported EOF for unit test or other reasons, did so by explicitly adding code to interpret zero-length messages with opcode zero (ZLM0s) as EOF, and passing them from input to output. A small number actually acted on the reception of ZLM0 to perform certain flushing operations. This ZLM0 recognition was required in order to make best use of the unit test framework and the extra worker code to process ZLM0s was considered “worth the extra code” in order to facilitate repeatable testing.

5.4.2 EOF-related Compatibility Measures for Pre-1.5 Workers.

In order to maintain backward compatibility with all such workers, in release 1.5 a “worker interface version” attribute was introduced in OWDs which indicates whether a worker expects to see the older ZLM0 messages as EOFs or not. Newer workers set the version attribute to 2 (or later) in their OWD and get the benefit of the new, simpler, orthogonal EOF model (and other interface improvements introduced in release 1.5).

Workers that do not have this attribute are assumed to be “pre-version2”, which means that if and when an EOF indication is asserted at its input port, a ZLM0 is injected at that input port so the worker sees the ZLM0/EOF as it did before. When such a worker sends a ZLM0 on its output port, it is interpreted as the assertion of EOF.

Workers that were written to not process ZLM0s as EOF (because they would never see them), and not have protocols whose first operation had no arguments, would never see or generate ZLM0, so they are also compatible with release 1.5.

These compatibility measures do not extend to workers that *do not* process ZLM0 as EOF but *do* use protocols whose first operation has no arguments. In this case when they produce a ZLM0 it will be incorrectly interpreted as EOF. A separate option to suppress this output-ZLM0-recognition-as-EOF may be developed in a patch release if it is discovered that this type of worker actually exists and deserves this additional compatibility treatment rather than requiring conversion to a new “version 2” worker.

There is a discussion in the various authoring model documents about such conversions/upgrades for migrating older workers to use version 2 worker interfaces.

6 Component Specifications (typically in OCS XML files)

An OpenCPI **component** is a functional abstraction with a specifically defined control and configuration interface based on configuration properties, and zero or more data interfaces (**ports**), each with a defined messaging **protocol**. An **OpenCPI Component Specification (OCS)** file describes both of these aspects of a component, establishing interface requirements for multiple implementations (workers) in any authoring model. Workers are developed based on an OCS.

In the unusual case where there is no expectation of multiple workers implementing the same OCS, the XML for the OCS may be embedded in the OWD, as a **ComponentSpec** element.

The OCS describes two things: (1) the configuration properties of the component (how it is initially and dynamically configured and controlled), and (2) the (data) ports of the component (how it talks to other components). Based on these all components can be configured and interconnected in an application, regardless of component implementations. An OCS does not contain a behavioral description of the component, but only its interfaces, for use by both implementations and for applications.

A OCS file is the first step in having a component implementations built and ready for use in an application. This file is the basis for all the implementations. An OpenCPI worker is an implementation based on an OCS and a particular authoring model. The worker consists of two things:

1. A separate XML description called the OpenCPI Worker Description (**OWD**) of the particular implementation, indicating the worker's authoring model the worker is based on and the OCS is it implementing
2. The source code in some programming language that does the actual computing function of the implementation, written according to the authoring model.

The OCS XML contains component-global attributes, configuration aspects and data port aspects. A component specification is contained in the XML element whose type is **ComponentSpec**” which should be a top-level element in a file, structured as:

```
<ComponentSpec
  ---attributes---
  >
  ---child elements---
</ComponentSpec>
```

The OCS XML file is called the **spec file** for the component, and has a **-spec.xml** suffix. The spec files for all components in a library are usually found the **specs** sub-directory of the library. When groups of properties or groups of message protocol operations, or message types, are shared between spec files they are placed in separate **-prot.xml** or **-prop.xml** files. This allows for references from multiple spec files. The suffixes and locations of the files are required for the component library

management scripts. This also enumerates what files must be exported when applications use components in the library.

The spec files, and if necessary property and protocol files, are used by two different processes:

- The implementations, in worker subdirectories, need these files to ensure the implementation matches the specification.
- Applications need these files to correctly *use* the components and connect them to each other. Reference the *OpenCPI Application Guide* document.

It is strongly encouraged to use a common spec, and common unit tests, between different implementations of the same functionality defined by an OCS.

6.1 *ComponentSpec* Top-level Element

Below are the attributes and elements of the **ComponentSpec** top-level XML element. **ComponentSpec** elements may have **name** and **noControl** attributes, and may contain **property** and **port** child elements.

6.1.1 *Name* Attribute of the *ComponentSpec* element

The optional **Name** attribute of the component specification provides a name that is unique within its name scope. The attribute is case insensitive within a library or application. This means two different component specifications cannot differ only in case. When the **ComponentSpec** element is the top-level element of a file, the component name attribute is defaulted from the name of the file before any suffixes. That means this attribute is optional when the **ComponentSpec** is the top level element of a file. Omitting this attribute and using this default is recommended since this eliminates any confusing mismatches between the name of the OCS file and the name of the component in the XML.

6.1.2 *NoControl* Attribute of *ComponentSpec* elements

The **NoControl** attribute of the component specification is a Boolean attribute that indicates, when true, that components using this specification have no lifecycle/configuration interface at all. This is generally not allowed for application components but is specified for certain infrastructure components.

6.2 *Properties Element of ComponentSpec Elements*

The **Properties** element of a component specification has no attributes but consists of a list of **Property** child elements. The Properties element may be in a separate file and referenced using the `<xi:include href="<file>" />` syntax. This is useful when groups of **Property** elements are shared among multiple component specifications. However, the most common usage is to have **Property** elements directly enumerated under the top level **ComponentSpec** element, without using the **Properties** element at all.

6.3 Property Element of ComponentSpec or Properties Elements

A **Property** element describes one configuration property. It occurs as a sub-element of either the **ComponentSpec** element or the **Properties** element. A **Property** element describes the name, data type and accessibility of a configuration property. Its data type can be a scalar type or a structure. Each property can also be an **array** and/or a **sequence** of its data type. The term **array** refers to a fixed number of data values of the specified type. The use of the term **sequence** refers to a variable number of data values, up to a specified maximum length. All variable length data types used for properties must be bounded. See the sequences and strings data types for more information.

Property elements as described here may also be present in the OWD for a worker, to specify worker-specific properties beyond what is specified and required by the OCS.

6.3.1 Name Attribute of Property (and Member) Elements

The **Name** attribute is the case insensitive name of the property. A set of properties cannot have properties whose names differ only in case. Mixed case property names can be used for readability. When a **Properties** element includes other **Properties** elements there is still only one flat case-insensitive name space of properties for the component.

Description Attribute/Element of Property (and Member) Elements

The **Description** attribute allows the author to add documentation of the property in XML. This text can also be supplied in a child element, which is more convenient if the test is multiple lines. To allow multi-line descriptions to be conveniently indented in an XML file, leading white space is removed from the start of all lines to the extent that it is common to all lines. Currently this description text is not used in OpenCPI, but it is intended to support automation for component data sheets. An example of the white-space removal is:

```
<property name='prop1'>
  <description>
    Hello this is some descriptive text
    Where we want to show
      - multiline features
      - white-space removal
  </description>
</property>
```

would result in the description test being:

```
Hello this is some descriptive text
Where we want to show
- multiline features
- white-space removal
```

6.3.2 *Type Attribute of Property (and Member) Elements*

The **Type** attribute specifies the data type of the property. The legal types are listed in table [Data Types for Properties](#). When the **Type** attribute has the **String** value, the **StringLength** attribute must also be supplied. This additional attribute indicates the maximum length of the string property values, *excluding* any terminating null character. If no **Type** attribute is present in the **Property** element, the type **ULong** is used as the default.

When the type is **Enum**, the actual values are zero-based **ULong**, but the named values are indicated by strings found in the **Enums** attribute described below.

The **ArrayLength** attribute is used when the property is a fixed-length one-dimensional array of the indicated type. The **SequenceLength** attribute is used when the property is a variable length sequence of the indicated type.

When the type is **Struct**, the **Property** element itself has **Member** sub-elements that indicate the types of the members of the **struct** property. No **struct** members can be of type **Struct**. The **SequenceLength** and **ArrayLength** attributes may apply to **Struct** properties. **Member** child elements are similar to **Property** elements in that they describe the name and data type information for the member.

All types have a maximum length and Properties cannot have unbounded length.

6.3.3 *StringLength Attribute of Property Elements*

The **StringLength** attribute is required when the **Type** attribute is **String**, and indicates the maximum length, excluding any null termination, string value that this property can hold.

6.3.4 *Enums Attribute of Property Elements*

This attribute is required when the **Type** attribute is **enum**, and its value is a comma-separated list of strings naming the enumerated values. The actual values are **ULong** and are zero-based ordinals based on the position of the names in this list.

6.3.5 *ArrayLength Attribute of Property Elements*

The presence of this attribute indicates that the property values are a fixed length one-dimensional array of the type specified in the **Type** attribute, and that fixed length is indicated in the value of this attribute.

This attribute is a convenience shorthand for using the **ArrayDimensions** attribute described below. E.g. if it is used like this:

```
ArrayLength='5'
```

This is really shorthand for:

```
ArrayDimensions='5'
```

6.3.6 *SequenceLength Attribute of Property Elements*

The presence of this attribute indicates that the number of property values is a variable, but bounded, sequence of the type specified in the **Type** attribute. The maximum length is indicated in the value of the **SequenceLength** attribute. This property has the specified maximum length, and always contains a current length, up to that limit. *When both **SequenceLength** and **ArrayLength** attributes are present, the meaning is sequence of arrays, not array of sequences.*

6.3.7 *ArrayDimensions Attribute of Property Elements*

The value of this attribute is a comma-separated list of array dimensions indicating an array whose number of dimensions is the number of values in the list. If this attribute is set, then the **ArrayLength** attribute should not be set. This attribute implies that values are multidimensional arrays of elements whose type is indicated by the **Type** attribute.

Property Accessibility

6.4 Accessibility Attributes

The accessibility attributes of the **Property** element specify how and when properties get their values. The boolean attributes that affect property accessibility are: **parameter**, **initial**, **writable**, and **volatile**. They all have the default value of **false**. When described as being “set”, it implies set to **true**. The attributes that specify property values in XML are **default** or **value**. How all these attributes are used is defined below and a summary of the valid combinations is in a table after that.

From the perspective of control-software (an ACI application or the **ocpirun** utility), the *initial* value of a property can come from four **sources**:

Set by control software itself via:

- *command line arguments to **ocpirun***
 - *the **PValue** argument to the C++/Python application constructor in the ACI*
3. Set in the application XML (OAS)
 4. Set in the the OCS XML property elements

Set by the worker

This list is in precedence order, with earlier mentioned sources overriding later ones (with one exception mentioned below).

Property values may be set/established at three different **times**:

At build/compile time, based on values in XML files.

At just before applications are started

35. After an application starts, during execution

The **parameter**, **initial** and **writable** attributes specify where the value comes from (its source) and when it may be set (its time). None or one of them may be set.

Parameter: the value is established in XML and has a constant, compile-time value, either determined in the OCS or determined by the worker's XML. Source is #1 and time is #1. No runtime changes are allowed.

Initial: the value is established when an application is initialized, before it is started, according to the precedence rules. Sources are #1, #2, or #3, and times are #1 or #2. No changes are possible *during* execution.

Writable: the value is established like **initial**, but may also be dynamically set it any time *during* execution (adds time #3 to **initial**).

<None>: If none of the **parameter**, **initial**, or **writable** attributes is set, the value is established by the worker (only source #4).

The **volatile** attribute specifies whether the worker may change the value *after* the application starts (time #3). The **volatile** attribute may not be set with **parameter**.

The **default** attribute may specify a default value in the XML, when **parameter**, **initial**, or **writable** is set. As an exception to the precedence rules, a *parameter* with a **default** value may be overridden by the underlying worker, in its XML.

Finally, for parameters only, the **value** attribute may be used instead of the **default** attribute when the value is fixed and may not be changed by precedence rules at all.

Reading Properties from the Control-Software Perspective

All properties are considered readable by control software. Parameter property values are read as the constant value set in XML (by OCS or worker XML). The worker is thus not involved in reading. Otherwise the following discussion applies.

If the property is indicated as **volatile**, the worker may change the value at any time and the value control-software reads is what the worker provides at that moment, with no caching. Even when the volatile property is specified as **initial** or **writable**, the value read by control software is what the worker provides at that moment and not necessarily what was written previously by control software.

If the property is *not* indicated as **volatile**, and it is specified as **initial** or **writable**, control-software simply reads back what it last wrote, using a cache and not involving the worker. If no boolean accessibility attributes are set, control software reads back the constant value that the worker provided after its own initialization, using a cache in the control-software after the first read.

The Order in which Properties are Written during Startup

This section only applies when properties are specified as **initial** or **writable**. Initial values are written prior to workers being started but after they have been initialized, with no inter-property ordering guarantees. To process multiple properties as a group, a worker must wait until the start operation or the first time the worker is run, since that is the only way to be sure all the properties that will be written, have been written. I.e. if some startup action depends on a combination of the values of multiple properties, this cannot be done prior to the start operation.

Property write notifications (telling the worker when a property has been written, different in each authoring model) will still be made prior to the start operation, to allow workers to know that a property has been written at the earliest possible time. However, no assumptions on ordering between these early property writes can be made. A worker may keep track of what properties have been written to determine when multiple properties have been written based on write notifications. Otherwise, it should wait until the start operation.

Initial Attribute of Property Elements

See the introduction above on the semantics of this boolean attribute. It precludes setting the value by control software during execution, allowing an initial value to be specified from any source but the worker itself.

The **initial** attribute setting in an OCS provides the most flexibility to workers implementing the specification, and is the most common property accessibility setting.

It allows workers to choose whether to implement the property as compile-time (setting **parameter** in the worker XML and perhaps building the worker for multiple compile-time values) or runtime (leaving it unchanged from the OCS, allowing any value). The **initial** setting allows there to be multiple workers that make this choice differently (compile-time/fast/small for a fixed set of values vs. runtime and flexible for all values).

The **initial** setting provides another benefit that makes workers easier to write: the worker does not have to deal with any changes in the property's value after the application starts.

allowing implementations to *either*:

- Have the property be configurable to any value at execution time
—*or*—
- Have the property value be fixed and compiled in to the implementation.

This enables some implementations to be flexible and allow different values at runtime, while other implementations can fix the value at compile time. When the application specifies a particular value, the worker with the compiled-in value can only be used if the requested value matches the compiled-in value. Workers with the non-parameter property can be used with any value. This second usage of the **parameter** attribute is discussed further in the section on the OWD XML files.

6.4.1 Writable Attribute of Property Elements

See the introduction above on the semantics of this boolean attribute. It enables setting the value by control software *during* execution.

Setting the **writable** attribute requires that the worker accept and implement dynamic runtime changes in the property's value. This setting should only be used if this is a true requirement of the component since it is more demanding on all workers that implement the specification.

6.4.2 Parameter Attribute of Property Elements

See the introduction above on the semantics of this boolean attribute. It indicates that the property's value is used at compile/build time when source code is processed into a binary artifact to be loaded and executed at run-time.

There are two uses for properties designated as **parameters** using this attribute:

1. A convenience variable for defining other attributes like string and sequence lengths or array dimensions, or as the basis for other property values when those values use expressions which use parameter properties as variables.

Making this property's value available for all workers as a compile-time constant.

The convenience usage (1) allows properties defined as parameters to be used in expressions for the value of **stringlength**, **sequencelength**, **arraylength**, **arraydimensions**, **value** or **default** attributes. For the allowable syntax of such expressions see the attribute expressions section below. An example is when multiple

properties are to have the same array dimensions, or to have array dimensions that relate to each other, e.g. one twice as long as the other. An example is:

```
<property name='nbranches' default='14' parameter='true' />
<property name='tree1' arraylength='nbranches' />
<property name='tree2' arraylength='nbranches*2 - 1' />
```

The second usage (2) makes the value for all implementations a compile-time value. If the property has a **value** attribute, that value is the fixed value for the property. If the property has a **default** attribute, that value is the value unless the worker specifies otherwise (see the OWD and build.xml sections).

These **parameter** values are also made available to worker source code during the compilation process, in a way specific to the authoring model and described in the document for the authoring mode.

Expressions in Numeric Attributes

For attributes that take numeric values, such as **StringLength**, **ArrayLength**, **SequenceLength**, and **ArrayDimensions**, the values can be non-negative numeric values, and can be expressions that may use properties defined as parameters as variables in the expression. They are parsed as the type **ulong**. The full expression syntax is described in detail in [Property Value Syntax](#) section. An example is above in the [Parameter Attribute](#) section.

6.4.3 Default Attribute of Property Elements

The name of this attribute is **default**. This attribute provides a default value for the property for all implementations and the initial value precedence rules described above. It is parsed based on the data type specified in the **Type** attribute, and if a numeric type, may be an expression (see the previous section). The format of the string value of this attribute is described in the [Property Value Syntax and Ranges](#) section.

6.4.4 Value Attribute of Property Elements

The **value** attribute is similar to the **default** attribute described above, except that its usage indicates that the value may not be overridden elsewhere (not by the application, not by the worker). Its primary use is simply as a convenient constant value that may be reused in other numeric expressions. It may only be used for **parameter** properties.

6.5 Port Element of ComponentSpec Elements

The component specification defines ports using the **Port** element. It specifies the direction/role of the port, producer or consumer, and the message protocol used at that port. For backward compatibility, **DataInterfaceSpec** can be used in place of **Port**.

The **Port** element has several attributes and one optional child element: the **Protocol**. The protocol is usually specified using the **protocol attribute**, but can also specified inline using the **protocol element** instead.

6.5.1 Name Attribute of Port Elements

This attribute specifies the name of this port of the component. The value of the **name** attribute is a string that is constrained to be valid in various programming languages. It must be unique and case insensitive within the component specification.

6.5.2 Producer Attribute of Port Elements

This Boolean attribute indicates whether this port has the role of a producer, when true, vs. the default of false for a consumer. *There is no Consumer attribute.* This attribute indicates whether the port acts as a consumer (input) when it is not set at all (which defaults to false), or is explicitly set to false (not needed since it is the default).

All ports are considered input/consumer ports unless this attribute is set to true.

6.5.3 Protocol Attribute of Port Elements

Not be confused with the **Protocol element**, described in the [OPS](#) section, this string attribute names an XML file containing the OPS for the port. The file is expected to contain a **Protocol** element at its top level. If the port being described is permissive, meaning it may produce/consume any protocol, then this attribute can be absent. An example of a permissive component is a file writing component that logs any messages as input, regardless of protocol.

As with all attributes that refer to an XML file, the **.xml** suffix is assumed if not present, and the file is sought using the search path for XML files.

When a protocol is not specified, the following protocol summary attribute values are implied:

Table 5: Protocol Summary Attribute Values when there is no Protocol

Name	Value
DataValueWidth	8
DataValueGranularity	1
DefaultBufferSize	system default
DiverseDataSizes	true
MaxMessageValues	none
MinBufferSize	0
MinMessageValues	0
NumberOfOpCodes	256
UnBounded	true
VariableMessageLength	true
ZeroLengthMessages	true

6.5.4 Optional Attribute of Port Elements

This is the attribute whose name is **optional** which is an optional attribute. This Boolean attribute indicates whether the data port may be left unconnected in an application. The default value of false indicates that workers implementing this component require that this port have a connection to some other worker in the application. When true, this port may be left unconnected and all workers implementing this specification must support the case when the port is not connected to anything.

6.5.5 Component Specification Examples

Here is an example of a component specification that declares one float property that can be set during initialization, but not during operation. It has one output producer port that uses the protocol defined in the `ushort_1k-proto.xml` file.

```
<ComponentSpec
  <Property Name="size" Type="float" Initial='true' />
  <Port Name="lvds_tx" Producer="true" Protocol='ushort_1K-proto' />
</ComponentSpec>
```

7 Property Value Syntax and Ranges

This section describes how property values are formatted to be appropriate for their data types. Property values occur in the `default` or `value` attributes of property elements described above. This syntax is also used when property values are specified in the worker Makefile described below. The type names presented are those acceptable to the **Type** attribute of the property element in the OCS file.

Remember that attribute values in XML syntax are in single or double quotes. The syntax described here is used inside these quotes. To have quotes inside attribute values the other type of quotes is used to delimit the attribute value. In either case, inside the quoted attribute value, the `&` and `<` characters must be escaped using the official XML notions: `&` for `&`, `<` for `<`. If *both* types of quotes must be in an attribute value, then the official XML escape sequences can be used: `"` for double-quote, and `'` for single quote.

Property values are also used when running applications. That usage is described in the *Application Development Guide*, but the format is as described here.

7.1 Values of Unsigned Integer Types: *uchar, ushort, ulong, ulonglong*

These numeric values can be entered in decimal, octal with leading zero, or hexadecimal with leading 0x. The limits are the typical ranges for unsigned 8, 16, 32, or 64 bits respectively.

The **uchar** type can also be entered as a value in single quotes, which indicates that the value is an ASCII character, with backslash escaping as defined in the C language. The syntax inside the single quotes is as described for the **char** type below.

7.2 Values of Signed Integer Types: *short, long, longlong*

These numeric values can be entered in decimal, octal with a leading zero, or hexadecimal with a leading 0x, with an optional leading minus sign to indicate negative values. The limits are the typical ranges for signed 16, 32, or 64 bits respectively.

7.3 Values of the Type: *char*

This type is meant to represent a character, i.e. a unit of a string. In software it is represented as a signed char type, with the typical numeric range for a signed 8-bit value. The format of a value of this type is simply the character itself, with the typical set of escapes for non-printing characters, as specified in the C programming language and IDL:

`\n \t \v \b \r \f \a \\ \? \' \"`

A series of 1-3 octal digits can follow the backslash, and a series of 1-2 hex digits can follow `\x`.

OpenCPI adds two additional escape sequences as a convenience for entering signed and unsigned decimal values of type char. The sequence `\d` may be followed by an optional minus sign (`-`) and one to three decimal digits, limited to the range of -128 to

127. The sequence `\u` can be followed by one to three decimal digits, limited to the range of 0 to 255.

These escapes can also be used in a string value. Due to the requirements of the arrays and sequence values, the backslash can also escape commas and braces, i.e.:

`\, \{ \}`

7.4 Values of the Type: float and double

These values represent the IEEE floating point types with their defined ranges and precision. The values are those acceptable to the ISO C99 `strtof` and `strtod` functions respectively.

7.5 Values of the Type: bool

These values represent the Boolean type, which is logical true or false. The values can be case insensitive: `true` or `1` for a true value, and `false` or `0` for a false value.

7.6 Values of the Type: string

These values are simply character strings, but also can include all the escape sequences defined for the `char` type above. Due to the requirements of the arrays and sequence values, the backslash can also escape commas and braces (`\,` and `\{` and `\}`). Double quotes may be used to surround strings, which protects commas, braces, and leading white space. To be interpreted this way, the first character must be a double quote. Two double quotes can represent an empty string.

7.7 Values in a Sequence Type

Values in a sequence type are comma-separated values. When the type of a sequence is `char` or `string`, backslash escapes are used when the data values include commas.

7.8 Values in an Array Type

When a value is a one-dimensional array, the format is the same as the sequence, with the number of values limited by the size of the array. If the number of comma-separated values is less than the size of the array, the remaining values are filled with the *null* value appropriate for the type. Null values are zero for all numeric types and the type `char`. Null values for string types are empty strings.

7.9 Values in Multidimensional Types

For multidimensional arrays or sequences of arrays, the curly brace characters (`{` and `}`) are used to define a sub-value. For example, a sequence of 3 elements, each consisting of arrays of length 3 of type `char`, would be:

`{a,b,c},{x,y,z},{p,q,r}`

This would also work for a 3 x 3 array of type `char`. Braces are used when an item is itself an array, recursively.

7.10 Values in Struct Types

Struct values are a comma-separated sequence of members, where each member is a member name followed by white space, followed by the member value. A struct value can be “sparse”, i.e. only have values for some members. If the struct type was:

```
struct { long e1[2][3]; string m2; char c; }; // C pseudo code
```

A valid value would be:

```
e1 {{1,3,2},{4,5,6}}, c x
```

This struct value would not have a value for the `m2` member.

7.11 Expressions in Property Values

Both numeric and string typed scalar values can be specified using an expression syntax and operator precedence from the C language, where any parameter property with a value can be accessed as a variable. All C expression operators can be used except the comma operator, assignments or self-increments/decrements. The conditional operator using `?` and `:` is supported. Expressions can be used as elements of arrays or sequences, or as structure member values.

For example, if the `nbranches` property was a parameter, a valid expression elsewhere might be:

```
nbranches == 0x123 ? 2k-1 : 0177
```

7.11.1 Numeric Values

The numeric constant syntax is typical C language syntax (integer and floating point), with the following additions:

- Integers with explicit radix after a leading 0 can use `0t` for base 10 and `0b` for base 2, in addition to the normally used `0x` for base 16 and no letter for base 8. All these prefixes can be applied to the fraction and exponent for floating-point syntax.
- Integers can use a letter suffix of `K`, `M`, or `G`, upper or lower case, indicating 2^{10} , 2^{20} or 2^{30} respectively. E.g. `2k-1` is 2047.
- All arithmetic is done using a numeric data type exceeding the range and precision of `uint64_t`, `int64_t` and `double`, and then assigned to the actual data type of the property whose value is being specified.
- The `**` binary operator (pow) from the python and FORTRAN languages is also supported.

When the value of the expression is assigned to the property value or numeric property attribute (e.g. `ArrayLength`), it is range checked for validity. Boolean properties are set to true if the value is non-zero. Fractions are discarded when assigning values to integer types.

7.11.2 String Values

Within expressions, string constants (using double quotes as in C) and string-valued parameter properties can be used. All comparison operators are case sensitive and result in boolean numeric values (0 or 1). All operators requiring boolean values (`!`, `||`, `&&`, `?:`) use the length of the string (zero being false, otherwise true). The `+` operator concatenates strings. There is no implicit or explicit conversion between string values and numeric values. E.g. if `sparam` is a string-typed parameter with the value `abc`, then this expression has the numeric value of 1:

```
sparam == "abc"
```

This expression would have the string value `xyz_abc`:

```
"xyz_" + sparam
```

Expressions used for string values

To distinguish when a string value is an expression, as opposed to a string that might look like an expression, the string value must have a special prefix of `\:` (backslash followed by colon). So if a property (or sequence/array element or structure member) is a string type, any value assigned to it is considered not to be an expression by default. To make the string value interpreted as an expression, use the `\:` prefix.

So if `sparam` is a parameter with value `abc`, then specifying the property's string value

```
sparam+"xyz"
```

would simply define that exact string (with plus sign and double quotes included) but if the string value was:

```
\:sparam+"xyz"
```

then it would be interpreted as an expression, and the actual string value would be:

```
abcxyz
```

8 Worker Descriptions in OWD XML Files

Each worker directory contains an XML file describing the worker and references the spec file typically located in the component library's **specs** sub-directory. This XML file is referred to as the **OpenCPI Worker Description (OWD)**. The generic, common across authoring models, aspects of these implementation description files are described in this section. The OWD files are specific to different authoring models. These differences are described in the respective documents for each authoring model. Some authoring models allow multiple workers to be implemented in one worker directory. In these cases multiple OWDs may be in a single worker directory.

The worker description file adds non-default implementation information to the basic information found in the spec file. Each authoring model defines what the worker-specific information might be. An example would be the width of an FPGA data path for a port. All OWDs have as the top-level XML element an **XYZWorker** element, where **xyz** is the authoring model of the worker.

If the worker has no non-default behavior, there is no need for an edited OWD. In this case the framework (**ocpidev**) will generate a default one. This default OWD simply contains a reference to the spec file and specifies the authoring model and language. For example, if an RCC worker based on the spec file **search-spec.xml** only had default implementation attributes, the OWD file would be:

```
<RccWorker spec='search-spec' />
```

This section describes aspects common to the OWDs for all authoring models. Further OWD content specific to authoring models is described in their respective documents.

The top level element must refer to an OCS by either:

- indicating an OCS file by using the **spec** attribute (preferred)
- containing a ComponentSpec child element, not shared with any other OWD (rare)

Below is an example OWD for an HDL worker, found in **fastcore.hdl/fastcore.xml**. The **fastcore** implementation of the **core-spec** specification is using the HDL authoring model. It references the component specification found in the **core-spec.xml** file, probably in the **specs** directory of the component library containing this worker:

```
<HdlWorker Spec='core-spec'
  ---other attributes---
  >
  ---other child elements---
</HdlWorker>
```

OWD files contain the information (metadata) that the OpenCPI build process and runtime environment need to know about the worker that is not in its source code.

A second, optional, XML file in a worker's directory is its

8.1 XML Attributes of the Top-level XYZWorker Element in the OWD

8.1.1 Name Attribute of the XYZWorker Element.

The **Name** attribute defaults to the name of the OWD XML file itself without the directory or extension and is normally omitted. The **Name** attribute of the component implementation is constrained to be an identifier in several contexts. It is sometimes called the worker name or implementation name.

Worker names may include both upper and lower case for stylistic or programming language purposes. The OpenCPI framework identifies workers in a case insensitive manner. There should not be two workers using the same authoring model in the same package namespace whose names differ in case.

The name of the implementation may be the same as the name of the OCS. It is not required to have a unique name for the OWD unless there are multiple implementations of one OCS that use the same authoring model. I.e. OWD names are implicitly scoped by authoring model.

8.1.2 Spec Attribute of the XYZWorker Element

This string attribute specifies the name of the file for the OCS for this worker. The build scripts and makefiles automatically place the specs subdirectory, in the component library's top-level directory, into the search path when these worker description files are processed. The spec files need only be referenced by their name and not any directory or pathname. If the spec file is outside the component library, it can be a relative or absolute pathname. The **.xml** suffix is assumed and not needed.

8.1.3 Language Attribute of the XYZWorker Element

This string attribute specifies the source code language used in this worker. The valid languages depends on the authoring model, and for each model there is a default language. Some authoring models have only one valid language in which case this attribute is not required.

8.1.4 Version Attribute of the XYZWorker Elements

This attribute has type **uchar** with a default value of zero if unspecified.

The interface between workers and their environment (control and data) evolves, usually backward-compatibly. When there is an incompatible change, the worker interface version is incremented, allowing the tools and runtime environment to continue to support workers written to the older interface. Actual interface changes are described in the authoring model documents (e.g. RCC, HDL etc.).

The worker interfaces were updated incompatibly in OpenCPI release 1.5, to version 2, which means that workers written to the new version 2 interface must explicitly set this attribute to 2. While most changes are specific to authoring models, a significant change to end-of-file indications and zero-length-message handling was made and is described in the [End-of-File \(EOF\) Indications](#) section. New workers should be written to version 2, and older workers that undergo significant revisions should too.

8.1.5 ControlOperations Attribute of the XYZWorker Element.

This attribute contains a comma-separated list of strings identifying the implemented control operations. For operations that are mandatory for the authoring model, they are assumed. The default implies a minimal implementation that only implements those operations required by the authoring model. The control operations are listed in the [LifeCycle Control](#) section. Control operations that are required by the authoring model do not need to be mentioned. When only mandatory operations are implemented, this attribute need not be specified.

8.1.6 OnlyPlatforms Attribute of the XYZWorker Element

This attribute can contain a list of platforms that are the only ones this worker should ever be built for. This should be used to indicate a permanent characteristic of the worker: that it is known to not be buildable or workable except for these. A (rare) example would be if the code had platform-specific code for a certain set of platforms.

To temporarily and frequently change the list of platforms to build for, use this same attribute in the worker's build configuration file in [Build Configuration Files](#).

8.1.7 ExcludePlatforms Attribute of the XYZWorker Element

This attribute can contain a list of platforms that this worker should never be built for. This should be used to indicate a permanent characteristic of the worker: that it is known to not be buildable or workable for these.

To temporarily and frequently change the list of platforms to build for, use this same attribute in the worker's build configuration file in [Build Configuration Files](#).

8.1.8 OnlyTargets Attribute of the XYZWorker Element

This attribute can contain a list of targets that are the only ones this worker should ever be built for. This should be used to indicate a permanent characteristic of the worker: that it is known to not be buildable or workable except for these. A (rare) example would be if the code had target-specific code for a certain set of platforms.

To temporarily and frequently change the list of targets to build for, use this same attribute in the worker's build configuration file in [Build Configuration Files](#).

8.1.9 ExcludeTargets Attribute of the XYZWorker Element

This attribute can contain a list of targets that this worker should never be built for. This should be used to indicate a permanent characteristic of the worker: that it is known to not be buildable or workable for these.

To temporarily and frequently change the list of targets to build for, use this same attribute in the worker's build configuration file in [Build Configuration Files](#).

8.1.10 SourceFiles Attribute of the XYZWorker Element

This attribute contains a list of subsidiary files that should also be compiled and linked together with the worker. For FPGA code, the files are guaranteed to be compiled

before the worker source file is compiled so that entities in these files can be directly instantiated in the worker source code file without redundant component declarations.

8.1.11 Libraries Attribute of the XYZWorker Element

This attribute contains a list of primitive libraries that the worker requires when it is built. RCC primitive libraries are not yet supported. When these libraries have slashes in their names they are directly naming specific primitive libraries (perhaps standalone), with either absolute or relative pathnames. With no slashes, they represent libraries that are subject to the search rules for primitive libraries, which is, in order:

- in the same project
- exported from projects the current project explicitly depends on
- exported by the `ocpi.core` project

8.1.12 IncludeDirs Attribute of the XYZWorker Element

This attribute contains a list of directories (via relative or absolute pathnames) that will be searched when including files during building.

8.1.13 ComponentLibraries Attribute of the XYZWorker Element

This attribute contains a list a component libraries to be searched when building this worker. Component libraries are needed when this worker refers to others that are not in the same library.

8.1.14 Endian Attribute of the XYZWorker Element.

This attribute specifies the endian behavior of the worker code. When workers are built, the build process may be run in three different modes to create three different types of binaries:

- Little endian
- Big endian
- Dynamic endian based on an input supplied at runtime

The third way, dynamic, is generally not relevant for software since compilers only generate code for a specific assumed endianness. But it is relevant to the FPGA build process to support FPGA bitstreams that can operate in both modes. This OWD attribute specifies how the worker's code will work in these three build scenarios ,as specified by the ***ocpi_endian*** parameter which is present for all workers of all types.

Table 6 – Worker Endian Attribute Settings

Endian Attribute	Description
Neutral	The worker code is unaffected by endian parameter settings and is correct regardless of the setting. This is the default value, and is generally correct for software workers.
Little	The worker code is unaffected by endian parameter settings and is will only operate correctly in a little endian mode.
Big	The worker code is unaffected by endian parameter settings and will only operate correctly in a little endian mode.
Static	The worker code will respect the endian parameter when set to “little” or “big” and the resulting binaries will operate correctly according to the compile-time parameter setting.
Dynamic	Worker code will respect all three values of the endian parameter. If the parameter is little or big , the resulting binaries work in the requested mode. If the parameter is dynamic , the resulting binaries will work in an endian mode specified by an input signal or variable as specified in the authoring model. Not all authoring models have this option.

8.2 *Property and SpecProperty Child Elements in the OWD*

Properties specified in the OCS indicate the external configuration interface for all implementations of the same spec. Properties specified in the OWD define additional worker-specific properties, beyond those in the OCS. The description of OCS **property** elements in [Property Elements in OCS](#) applies when defining worker-specific properties in the OWD, with some additional attributes described below that are valid *only* in OWD properties.

The OWD can also provide *additional* attributes to properties already defined in the OCS. E.g. the OWD might make an OCS property **writable** that was not writable in the component spec. The accessibility added would result in the worker supporting a superset of what was required by the component spec.

While a component spec can only contain **Property** elements, a worker description can contain both **Property** and **SpecProperty** elements. The **Property** elements in the OWD introduce new worker-specific properties unrelated to those defined in the OCS. The **SpecProperty** elements add worker-specific attributes to the properties *already defined* in the OCS.

There are a number of OWD-only attributes in the next section. They generally apply to both **Property** and **SpecProperty** elements in the OWD. Property elements in the OWD support all the attributes supported by Property elements in the OCS, plus the OWD-only attributes below. SpecProperty elements in the OWD support all the OWD-only attributes plus a small subset of the attributes supported in OCS Property elements.

SpecProperty elements can add OWD-only attributes to a property and augment the accessibility of the property, but cannot change the data type of the property.

Name Attribute for OWD Property or SpecProperty Elements

The **Name** attribute is the case insensitive name of the property. The **Name** attribute is used in **SpecProperty** elements to indicate which OCS property is being referenced. In the **Property** elements it indicates the name of the worker-specific property, which must not be the same as any **Property** element in the OCS.

Attributes only Allowed in OWD Properties and SpecProperties

These attributes only exist in OWDs and never in OCSs and apply to both **Property** and **SpecProperty** elements.

8.2.1 *ReadSync and WriteSync for OWD Property or SpecProperty Elements*

These Boolean attributes, defaulting to false, are used to indicate the properties that require the worker to be notified when they are read or written by control-software. The baseline behavior is that property accesses are directly made to property values in the worker, with no specific synchronization or notification implied. The worker accesses these values as local memory locations. When these attributes are true, the worker is notified when the access is made by control software.

The exact mechanism used for such worker notification is specific to the authoring model and is described in those documents. Some authoring models may not implement or require this attribute, but where needed, this definition is valid.

8.2.2 *ReadError/WriteError Attributes for OWD Property/SpecProperty Elements*

These Boolean attributes, default is **false**, indicate properties that may return errors when read, **ReadError**, or written, **WriteError**. If a worker does not return errors and always succeeds when property values are read or written, then leaving these values false allows control-software to avoid any error checking. In some models and systems error checking can carry significant overhead. Most workers simply accept new values using the default of false for this attribute.

The exact mechanism used for such worker error reporting is specific to the authoring model and is described in those documents.

Padding Attribute for OWD Property but not SpecProperty Elements

*Padding properties (its **padding** attribute is **true**), are those that only exist to force subsequent properties to proper alignment or offsets. No other accessibility attributes may be used with padding properties. The use case for this is when properties need to match a register map and may require exact offsets for each property. These padding properties are not accessible.*

8.2.3 *Readback attribute for OWD Property/SpecProperty Elements*

This attribute specifies that the worker supports reading back values written to non-volatile properties that are either **parameter**, **initial** or **writable**. It is never necessary for normal operation since control-software always caches such values when written, and provides those cached values when control-software reads the value.

This option allows such values to be read back directly from the worker when control-software requests an uncached read access, or when property values are being read directly from the worker and bypassing control-software altogether. An example of this is the **ocpihdl** utility for HDL/FPGA platforms (described in the HDL Development Guide) when a raw loaded bitstream is being examined directly.

8.3 Attributes Allowed in OWD SpecProperty Elements

In addition to OWD-only attributes described above, these are the attributes available in **SpecProperty** elements that are also used in OCS and OWD **property** elements. They are used to add accessibility or values beyond what is in the OCS.

8.3.1 Default Attribute for OWD SpecProperty Elements

In the **SpecProperty** element, this attribute is only valid for **parameter** properties and when a value is not already specified in the OCS (via **default** or **value** attributes). It might be used when the user provides its own default for the purposes of unit testing and build configurations where other values may be specified. See section [Build Configuration Files](#) for build configurations.

8.3.2 Value Attribute for OWD SpecProperty Elements

In the **SpecProperty** element, this attribute is only valid for **parameter** properties that are defined as parameters in the OCS. but do not specify a fixed value there.

An example would be where a property is defined in the OCS as a parameter (and possibly a default value) with the expectation that workers would provide a worker-specific fixed value.

Parameter Attribute for OWD SpecProperty Elements

This boolean attribute in a **SpecProperty** element is used to convert a property specified as **initial** in the OCS, to be a **parameter** property for this worker. The **initial** setting in the OCS allows the worker author to decide whether to implement it as a runtime property set at initialization time, or a parameter set at compile time. This attribute makes the latter choice.

When a worker has **initial** (in the OCS) properties that are **parameters** in the OWD, it means that the worker must be built for specific values of such properties. This has three implications:

- An application can only use the worker if the worker is built for a property value that matches what is requested as an “initial” value by the application.
- Binary component libraries may have multiple binary artifacts for the same worker, but with different combinations of parameter values.
- The worker developer must decide which combinations of parameter values to build, in order to make alternative settings of such parameters available.

This parameter feature allows workers to have compile time optimization for certain parameter values, and also allow a single worker source code module to be optimized for different values at compile time.

Parameter property values are applied to the build process as expected: e.g. by preprocessor symbol definitions for software, generics for VHDL, and parameters for Verilog.

8.4 Built-in Parameters of All Workers

OpenCPI automatically adds several properties to all workers. The values of these parameters are usually set during the build process in various ways. Some of these parameters are set to values by the build process and are not intended to be set manually at all. Others may be set or overridden manually.

Each authoring model may also specify additional built-in parameters for all workers using that authoring model. The built-in parameters that apply to all authoring models are described below.

All built-in parameters use the `ocpi_` prefix to avoid collisions with component and worker developers.

8.4.1 The `ocpi_debug` built-in parameter property

This boolean parameter property, indicates whether a debug build is being done. The default value is false. Setting this value to true indicates to worker source code that any debugging instrumentation or behavior should be enabled, at a potential cost of some resource usage and performance. This built-in parameter is always available, and should be used in worker code to enable things like extra logging or statistics keeping.

Setting this parameter to true will also in some cases enable some introspection or instrumentation capabilities of code that is in the OpenCPI infrastructure or is generated code used implicitly by the worker.

Properties can be defined with a debug attribute value of true, which indicates that those properties should only be present when the worker is built with this `ocpi_debug` parameter set to true.

These features allow debug behavior and debug properties to be permanently in the worker's source code and OWD while only being enabled as required.

8.4.2 The `ocpi_endian` built-in parameter property

[This feature is preliminary/untested, but mentioned here as a roadmap item]

This parameter property indicates to worker code which endian mode is being used when the worker is being compiled. Its type is an enumeration of three values:

- **little:** The build is intended to generate binaries for little endian systems
- **big:** The build is intended to generate binaries for big endian systems.
- **both:** The build is intended to generate binaries that can be used in either little or big endian mode, selected at initialization time in the runtime environment.

Software authoring models normally set this mode implicitly as compilers generate binary code for a specific endianness based on the processor being targeted , e.g. little for x86 and ARM, big for PPC.

However some authoring models, such as HDL, can support all three compilation modes. The ability of a worker's code to support various endian modes is specified in the worker's `endian` attribute at the top level of its OWD.

8.5 Port Elements of XYZWorker Elements

Ports are how workers communicate with each other. They define message-oriented, data-plane communication. Each authoring model defines how workers receive/consume and send/produce messages to or from other workers. This is independent of whether workers are collocated in the same container or executing elsewhere.

Each authoring model may have attributes and elements of this `Port` element specific to that authoring model, but there are a number of aspects common to all worker descriptions that are described here. A `Port` element in a worker description matches the `Port` element in the component spec by name, and adds worker-specific information about how the worker implements the port.

Some authoring models may use different names for `port` elements in order to imply additional information, but the attributes described below apply in any case.

8.5.1 Name attribute of Port Elements

This string attribute is required and must match one of the names of the port elements in the component spec. It indicates for which component port the worker is providing additional implementation-specific information.

8.5.2 Protocol Summary Override Attributes of Port Elements

The OCS file on which the worker is based defines protocols for ports, and various protocol summary attributes are derived from the protocol. These attributes may be overridden where the protocol is defined (see [Protocol Summary Attributes](#)), but they may also be overridden for a particular worker as attributes of the `port` element.

Care must be taken when doing such overrides since they may render a worker incompatible with other workers that implement the same spec. This could make unit testing of all such workers infeasible or difficult.

9 Worker Build Configuration XML Files

While the OWD file described above holds fundamental information about the source code, and information required to build and run the worker, the optional **build configuration file** describes the various ways it should be currently built.

Thus the build configuration file narrows down the perhaps large number of ways the worker *might* be built (for different combinations of parameter values or targeting different platforms), to what will actually get built. When it is not present, the worker is built using the default values of parameter properties, and the platforms being targeted is limited by the fundamental limitations on platforms expressed in the top-level attributes of the OWD.

For portable workers with no parameters, this file is not recommended or needed.

The use cases for this file include:

- specify useful combinations of compile-time parameters for unit testing
- specify known currently required combinations of compile-time parameters for applications
- limit platforms to build for, for temporary convenience rather than an expression of “what platforms are viable”, which is specified in the OWD.

The name of this optional file is **<worker>-build.xml**, and the top level element is **build**. The top level attributes controlling/limiting the build targets and platforms in the OWD are also valid here. Limitations in the OWD override those in this file since the OWD platform constraints are considered fundamental, while those in this file are considered for convenience to further limit what platforms or targets get built.

These platform/target constraint attributes that are defined for the OWD and also valid in the build file are: **OnlyPlatforms**, **ExcludePlatforms**, **OnlyTargets**, **ExcludeTargets**. There are no other top-level attributes for this file.

9.1 Build Configurations

We use the term **build configuration** to be a set of parameter property values to use when building the worker. Each build configuration has an **ID**. When the build configuration file does not exist, there is a single build configuration, with an ID of zero, with default values of all parameters (from OCS or OWD).

The **target-*** directories mentioned above are used to separate the files that result from building for different targets (e.g. **centos7** vs. **xilinx13_4**, or [Xilinx] **virtex6** vs [Altera] **stratix5**). A different **target-*** subdirectory is in fact created for each build configuration. If the ID is zero, the name of the directory is **target-*<platform>***. If it is not zero, then the directory name is **target-*<id>-<platform>***.

There are two valid child elements in the top level **build** XML element, which together determine the set of build configurations for the worker:

- **parameter** to specify parameter values for all configurations
- **configuration** to specify a build configuration.

If there are none of either element, a configuration with ID zero is generated using the default values of all parameter properties. The same thing happens when there is no file at all.

If there are **parameter** elements, but no **configuration** elements, then configurations are generated for all combinations of values mentioned in all the **parameter** elements, using default values for unmentioned parameter properties. Since **parameter** elements can provide multiple values (see below), this means that configurations are automatically generated for the cross-product of all parameter values. E.g. if property A is specified with values 1 and 2, and property B is specified with values 10, 11, and 12, and property C is unmentioned but having a default value of 77, then these build configurations will be generated:

ID	A	B	C
0	1	10	77
1	1	11	77
2	1	12	77
3	2	10	77
4	2	11	77
5	2	12	77

The file in this example would be:

```
<build>
  <parameter name='a' values='1,2' />
  <parameter name='b' values='10,11,12' />
</build>
```

When a **configuration** element is specified *with* an **id** attribute, all parameter properties unmentioned in the configuration element take on default values. Using the above example we might have:

```
<build>
  <configuration id='0'>
    <parameter name='a' value='2' />
    <parameter name='b' value='12' />
  </configuration>
</build>
```

Which would result in A being 2, B being 12 and C being 77.

When a configuration element is specified *without* an **id** attribute, configurations are generated that have the all the combinations mentioned in top-level parameter elements as in the above example, but with the values for parameter properties mentioned *inside* the configuration element taking those specific values. For example:

```
<build>
  <parameter name='a' values='1,2' />
  <parameter name='b' values='10,11,12' />
  <configuration>
    <parameter name='c' value='78' />
  </configuration>
</build>
```

This would create 6 build configurations (like the table above), but with property C having the value 78 rather than its default of 77.

9.2 Parameter elements in the `<worker>-build.xml` file

This element specifies one or more values for a parameter property. The **name** attribute specifies which property, and must match the name of the parameter property in the OWD or OCS (case insensitive). The values for the parameter can be specified using one of these attributes:

- The **value** attribute can specify the single value.
- The **values** attribute can specify a comma-separated list of values.
- The **valueFile** attribute can specify a value in a file.
- The **valuesFile** attribute can specify multiple values in a file.

When the **parameter** element is at the top level of the `<worker>-build` file, it is specifying values common to all configurations that do not have **id** attributes. If the same **parameter** element is mentioned as a child element of a **configuration** element, it overrides any top level value(s) for that parameter in that configuration. An example `<worker>-build` file is:

```
<build>
  <parameter name='debug' value='true' />
  <configuration id='1'>
    <parameter name='mode' value='lownoise' />
    <parameter name='taps' valueFile='taps.txt' />
  </configuration>
</build>
```

Since there is an **id** attribute present for the configurations, it would use the default value of the debug parameter, not the value specified in the **parameter** element (which in this case would be ignored).

9.3 Configuration Elements in the <worker>-build.xml File

The **configuration** element has a optional numeric **id** attribute which will appear as the build configuration suffix in the name of the target directory for that configuration. If the **id** attribute is zero, no suffix is added. The configuration element also has **parameter** subelements indicating the specific parameter values for that configuration.

When there is no **id** attribute, the globally defined values for parameters are used to generate build configurations.

10 Component Libraries

OpenCPI components are developed in libraries. OpenCPI component-based applications are defined as a composition of components, and the components are developed and built in component libraries.

A component library is a directory that contains:

- Component specifications, OCSs, OPSs in a **specs** subdirectory.
- Component implementations (workers), each in its own subdirectory.
- Component tests in ***.test** subdirectories.
- The **Makefile** for the component library.
- When built, a subdirectory call **lib**, which contains links the built binaries and metadata files required to use components in the library from outside the library.

A component library has two forms: source and binary. The source form is for component developers, and the binary form is for application developers and users. The binary form is the result of building the source library and exporting the results as a binary package. The package can then be installed onto a system such that the assets can be found and used by applications.

The exported version of a component library contains a combination of binary artifacts and XML files. The binary files can be a collection of heterogeneous built workers for various technologies.

Distribution of a library to an application developer or user requires the contents of the **lib** subdirectory. OpenCPI component libraries in binary form can contain compilations for different operating systems, FPGAs, or CPUs all in the same binary directory tree.

For most development efforts it is recommended to use the process described in the [Developing OpenCPI Component in Projects](#) section. This defines a larger directory structure containing a variety of OpenCPI assets, including component libraries and applications.

The basic directory structure of a component library is shown in the figure below.

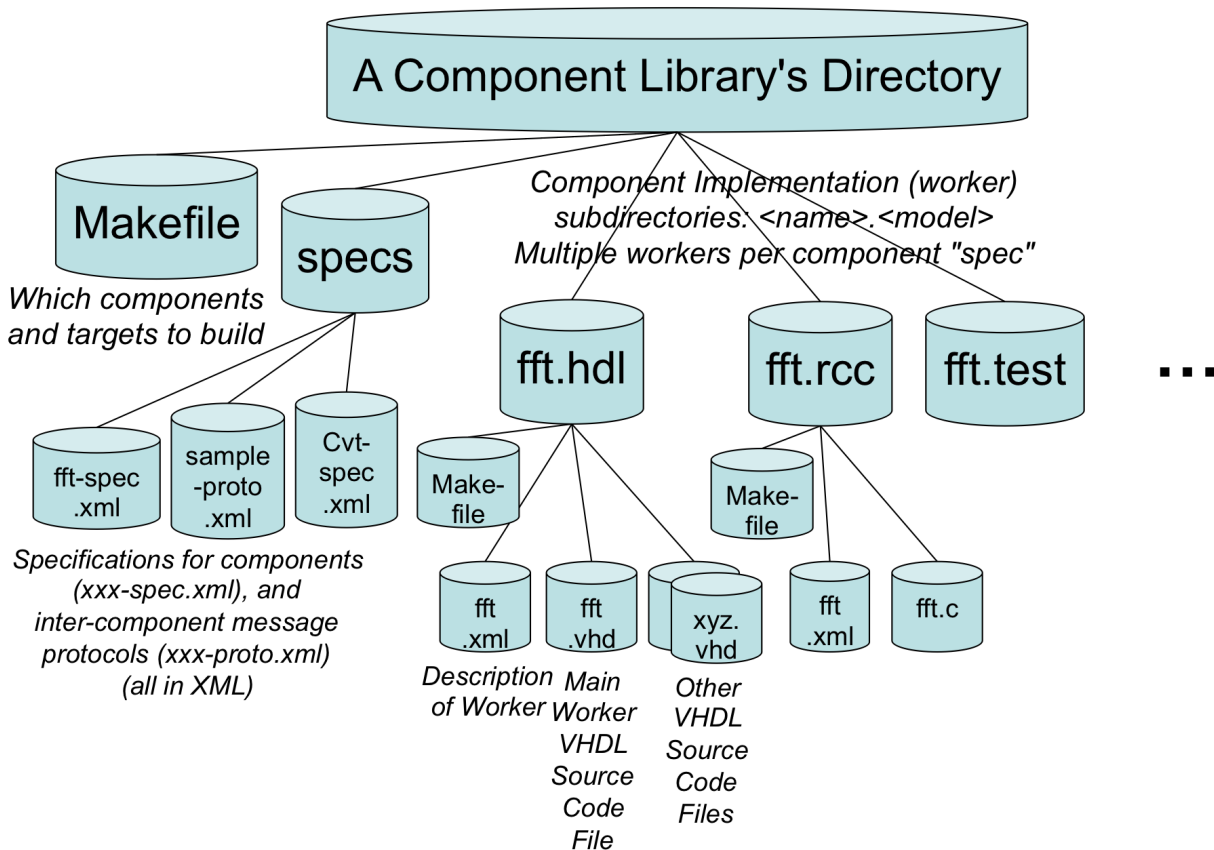


Figure 1: Component Library Directory Structure

10.1 The Component Library Makefile

The makefile in the top-level directory of the component library is usually generated automatically by the `ocpidev` command. The last line in the file establishes it as a makefile for an OpenCPI component library and should be:

```
include $(OCPI_CDK_DIR)/include/library.mk
```

The `OCPI_CDK_DIR` variable is set in the environment to point to the OpenCPI CDK installation. An important variable in this file is `Workers`, which is a list of which worker subdirectories to be built in this component library. When the `Workers` variable is not set at all, it indicates that all subdirectories of the component library that contain workers should be built. For example:

```
Workers=fft.rcc fft.hdl fft-for-xilinx.hdl fir.rcc
```

There are two reasons to set this variable at all rather than using the default.

1. If you want to temporarily avoid building some workers in the library, you can set this variable to only the ones you want to build, so any others are ignored.
2. If you want to specify the order in which the workers are built, you can set this variable to the workers you want to build, in the order you want them to be built. There are two situations where the order of building workers is important. First, if a worker is a proxy (see the **RCC Development Guide**) for another, the “slave” of the slave must be built before the proxy. Second, if a subdevice supports another device, the supported device must be built before the subdevice (see the **Platform Development Guide**).

The `ExcludeWorkers` variable can also be used to exclude certain workers from being built when the `Workers` variable is unset.

In order to avoid name space collisions when using multiple component libraries, there is also a `Package` variable that specifies what namespace should be used for the specs and workers in this library. See the [Package IDs](#) section for complete information about package IDs.

For each authoring model, there may be a (default) list of *platforms* to build for. I.e. for the RCC authoring model, the variable `RccPlatforms` would be set to a list of platforms to build all RCC workers for. For all software (not HDL) authoring models, the default RCC platform, if none is specified, is the machine and operating environment of the machine doing the building. Other software targets would use cross-compilers. These target variables can always be overridden on the command line or in the project's `Project.mk` file.

Other non-software authoring models (for processors that will never be the one running the tools), have other default platforms (described in documents for the authoring model).

Software targets use the software platform names. Examples are `centos6`, `centos7`, `macos10_13`, `xilinx13_4` (for Xilinx linux for Zynq).

HDL targets typically contain an architecturally compatible part family (e.g. **virtex6** or **stratix4**). See the ***HDL Development Guide*** for more details.

If all subdirectories containing workers should indeed be built, and the desired build targets are the default ones (or specified on the command line), and the package name is the default, then the single “include” line above is sufficient to build a component library.

Creating a new component library **myc1** is accomplished by using the **ocpidev** tool, using the **create library** command. This tool is described in the [ocpidev](#) section below.

```
ocpidev create library <library-name>
```

10.2 The Library.mk File

If present, the **Library.mk** file in the component library's directory contains settings that should apply to all the workers and tests in the library. This file is read as part of the **Makefile** of all workers in the library.

The table below lists the variables settable in this file. They are also usable at the project level in the project's **Project.mk** file which has a similar purpose.

Table 7: Variables in Library.mk files

Variable Name in Makefile	
Libraries	A list of primitive libraries built elsewhere. If a name has no slashes, it is searched in this project and in projects this project depends on.
OnlyTargets OnlyPlatforms	A list of the only targets/platforms for which this worker should be built
ExcludeTargets ExcludePlatforms	A list of targets/platforms for which this worker should NOT be built
XmlIncludeDirs	A list of directories elsewhere for searching for XML files.
IncludeDirs	A list of directories elsewhere for searching for files indicated by "include" directives in worker source files such as those in C, C++ or Verilog.
ComponentLibraries	A list of component libraries to search when the worker refers to another worker. The need for this is specific to the authoring model (e.g. proxies for RCC workers, or emulators for HDL device workers).

11 Developing Workers

This section describes the aspects of the worker development process that is common across all types of workers and authoring models. Previous sections above described the files involved in worker development, including:

OCS XML files: component specifications, usually in `../specs`

OPS XML files: protocol specifications, usually in `../specs`

OWD XML files: worker descriptions

Worker Build Configuration files: XML files for specifying build configurations

Worker Source files: Programming language source code for the worker.

This section describes the development process using these files.

The worker development details for each authoring model are described in the document for each authoring model. Some authoring models (e.g. RCC) support creating a single binary file **artifact** that implements multiple workers. However, usually a single worker implementation is in its own subdirectory, which when compiled results in a single binary file *for each build configuration* (combination target and parameter values).

11.1 Creating Workers

A worker is created, either standalone or in a component library, using the `ocpidev` tool, with the command:

```
ocpidev create worker <name> [-S <spec>] [-L <language>]
```

The authoring model is inferred from the `<name>`, using the suffix of the name as the authoring model. The optional `<spec>` argument specifies the name of the OCS file, normally without any directory indicated (expected to be in the `../specs` directory). If `<spec>` is not specified, it is assumed to be `<name>-spec.xml` in the library's `specs` directory. If the new worker will embed the component spec in its own OWD, then the `<spec>` argument can be set to `none`. While rare, some specialized workers will be the only implementation of a spec and there is no need for separate spec file.

The `<language>` is one of the programming languages allowed for the authoring model (e.g. `c`, or `c++` for RCC, `vhdl` or `verilog` for HDL). If not mentioned, the default language for the authoring model will be used.

The `ocpidev` command is usually executed in the directory where the new workers's directory will be created. Other options are fully described in the [ocpidev](#) section below.

The command:

```
ocpidev delete worker <name>
```

will remove the worker, and is essentially equivalent to (after asking for confirmation)

```
rm -r -f <name>
```

When a worker is created, all the worker's XML, **Makefile** and source language files are initially automatically generated by **ocpidev**. Several internal files (not for user editing) are also placed in a **gen** subdirectory of the worker's directory. When source files are compiled, the resulting binary files are placed in subdirectories named: **target-*<target>***, where ***<target>*** is the hardware the compilation is targeting. Cleaning (via **ocpidev clean**) a worker directory simply removes the **gen** and all **target-*** subdirectories. In almost all cases, files in the **gen** subdirectory should be considered read-only and not edited.

Creating a new worker creates initial versions of three files in the worker's directory;

1. the **Makefile**
2. the OWD file
3. the skeleton source file

These are the files the developer can edit as necessary. Although frequently the **Makefile** and the OWD XML file do not need any further editing.

The initial source file is termed the ***skeleton***, and is named

```
<worker-name>.<source-suffix>  
e.g.  
xyz.c
```

It can be compiled, but has empty logic. The skeletal code allows the worker to be test-built even before any editing is done. Each authoring model describes how and where this skeleton source file should be edited and “filled out” with the logic that makes it perform its intended function. A copy of this initial skeleton file is always put in the **gen** subdirectory, with the name:

```
gen/<worker-name>-skel.<source-suffix>  
e.g.  
gen/xyz-skel.c
```

This copy can always be examined to see what the skeleton was originally, before any editing. It can also be useful to examine, after the OCS, OPS, or OWD has changed, in case changes are require in the source file. After the initial skeleton (*not* the copy in “gen”) is edited by the developer, it will never be overwritten or removed by **ocpidev clean** or any other command.

11.2 The Worker Makefile

Each worker requires a **Makefile**. The default **Makefile** is simply:

```
include $(OCPI_CDK_DIR)/include/worker.mk
```

This makefile line indicates to the framework that this directory is for building a worker whose name and authoring model are derived from the name of the worker's directory. The directory must also contain a OWD and one or more source files for the worker's functional code. These will be generated when the worker is created using the `ocpidev` tool described in the [ocpidev](#) section.

For example if the name of the directory were `search.rcc`, then the simplest makefile would assume that the worker description file is in `search.xml`. If no language was specified in the OWD, the source code to compile for the worker would be `search.c`, since C is the default language for the RCC authoring model. The `ocpidev` tool will create the initial worker source file automatically as a skeleton of the implementation. This initial empty worker does nothing, but compiles as a valid worker of the given authoring model. This file can then be edited to add the logic to perform its function.

Several top level attributes in the OWD can also be set as variables in the Makefile. This is deprecated: it is preferable to put these values in the OWD as top-level attributes rather than as variables in the **Makefile**. The attributes in the OWD that may also be expressed as variables in the Makefile are: `Libraries`, `IncludeDirs`, `ComponentLibraries`, `OnlyPlatforms`, `ExcludePlatforms`, `OnlyTargets`, `ExcludeTargets`, and `SourceFiles`. See section [OWD Top-level Attributes](#).

If this worker directory is intended to build *multiple* workers into a single binary artifact file (when supported), the **Workers** variable specifies the workers to be built into a single artifact e.g.:

```
Workers=w1 w2
include $(OCPI_CDK_DIR)/include/worker.mk
```

The two files, `w1.c`, `w2.c` will be compiled together to form the worker binary file implementing `w1` and `w2`, as described by `w1.xml` and `w2.xml`.

In the normal case when there is only a single worker, the value of the **Workers** variable is inferred from the name of the directory and is not specified.

The worker **Makefile** is commonly left as created by `ocpidev`, but there are a few cases where some extra settings need to be present. The table below lists the variables settable in the makefile that are common to all authoring models.

11.2.1 Parameter Properties in Worker Makefiles

Parameter values for building workers are normally set if needed in the worker's **build configuraion file** described above. Setting them in the worker's **Makefile** as described here is deprecated.

Build configurations are sets of parameter property value settings, and these values may be set in the worker's **Makefile**.

To specify the value for a parameter, set a makefile variable whose name begins with **Param_** followed by the name of the parameter, such as:

```
Param_xxx=5
```

To specify multiple values for a parameter, makefile variables can be used to specify the values for any parameter for which the worker should be built. A makefile variable of the form:

```
ParamValues_xxx1=v1/v2/v3
```

specifies that the parameter whose name is **xxx1**, should be built for the values **v1**, **v2**, and **v3**. The forward slash character is used to separate the values rather than commas, since commas must be used within values when the parameter's data type is an array and/or sequence. This variable specifies these values independent of the values of any other parameter. In the case of the variables specified as:

```
ParamValues_p1=1/2/3  
ParamValues_p2=abc/xyz
```

then the worker will be built for each target specified, for all these combinations:

```
p1=1, p2=abc  
p1=1, p2=xyz  
p1=2, p2=abc  
p1=2, p2=xyz  
p1=3, p2=abc  
p1=3, p2=xyz
```

If the worker was built for two targets, this would result in the worker being built for twelve different build configurations. These are inferred from the settings of parameter variables and the target make command line settings and variables. These parameter variables can be specified on the command line when calling the makefile and at higher levels of makefiles. To build an entire component library with the **p1** parameter set to 2, simply specify this in the top level directory of a component library:

```
make Param_p1=2
```

The syntax of parameter values is the same as described for default values for properties in the Property Value Syntax section, with the limitation that string values cannot have spaces, quotes, commas or backslashes. If those are needed, the XML file method of specifying parameter values must be used.

There are a number of limitations when using makefile variables to specify parameter values:

- Different parameters are independent of each other creating the cross-product of all combinations.
- The value syntax is not generally usable for string values with embedded spaces, commas, quotes or backslashes.
- As with most makefile variable uses, misspelling the name of a makefile variable does not result in any error.

11.3 Editing Workers

Often it is useful to break the worker's logic into supporting code modules in other source files. Those files must be created manually and added to the **SourceFiles** make variable in the worker's **Makefile**. In some authoring models and languages, the files listed in the **SourceFiles** variable must be in dependency order, with lower level modules/files preceding those that depend on them. The primary source file is always considered the top level module for the worker and is essentially put at the end of the list automatically.

Some changes to the OWD, OCS, OPS and even makefiles can result in changes that require corresponding changes in the worker's primary source files, which was initially generated as a skeleton. Since the developer has likely manually edited the primary source file,, it is not touched when such changes are required. If it is clear to the developer when these changes are required, they can do it before any building. However, it is likely that the required changes will create build/compilation errors.

Examples would be such things are renaming ports or properties, adding or subtracting access attributes to properties, converting properties to parameters, etc.

When any changes are made, the skeleton is regenerated properly during the next build, and the result places in the file:

```
gen/<worker>-skel.<source-suffix>
```

The newly generated skeleton can be used as a guide when changes occur that might require changes in the edited worker source code.

The authoring model documents also list common changes to the OCS, OPS, and OWD files, and the corresponding changes required in the source file. An example is for VHDL workers using the HDL authoring model. The skeleton lists the lower level primitive libraries that the worker depends on. If such a primitive library is added to the worker makefile, the library needs to be added to the list of libraries in the skeleton.

12 The Worker Source Files

The worker source files must be written according to the authoring model. As a starting point OpenCPI provides the `ocpidev` tool to create an empty skeleton of a worker implementation that will in compile, build and execute, doing nothing.

The file hierarchy of a component library `mycl` is outlined below. The library contains a `search` component with RCC and HDL implementations, and a `transform` component with only an HDL implementation:

```
mycl/Makefile
  /specs/search-spec.xml
    /transform-spec.xml
  /search.rcc/Makefile
    /search.xml
    /search.c      (RCC C source file)
  /search.hdl/Makefile
    /search.xml
    /search.vhd    (HDL VHDL source file)
  /transform.hdl/Makefile
    /transform.xml
    /transform.vhd
```

12.1 How Parameter Value Settings Appear in Source Code.

Parameter values are compile-time constants in all authoring models. The precise way that parameters and their values appear in source code varies by authoring model and programming language. In most cases, there are standard data types, see [Data Types for Properties](#), for the OpenCPI properties, and constants are defined that specify these values. Examples are:

- C and C++: A `static const` variable is defined which is initialized to the parameter value. The name of the variable is the property name prefixed with `PARAM_`.
- VHDL: A generic with the parameter's name is set to the value.
- Verilog: A parameter with the parameter's name is set to the value.

12.2 Building Workers

Workers are normally built as part of building a whole component library, or as part of a whole project. To simply compile new code and locate syntax errors, a worker can be built in the worker's directory, by typing `ocpidev build`.

The target or platform of a worker build is specified in several ways. For software authoring models the default target is always the local development machine on which the building is taking place. For other authoring models, there is no default. On the `ocpidev` command line targets and platforms use options of the form:

```
--<model>-target <target>
--<model>-platform <platform>
```

Some examples are:

```
--rcc-platform centos6
--rcc-platform xilinx13_4
--hdl-target virtex6
--hdl-platform zed
```

Default target and platform settings can also be set in project's **Project.mk** file or the library's **Library.mk** file. In this case the syntax is setting Makefile variables of the form:

```
<Model>Target(s)
```

Some examples are:

```
HdlTarget=virtex6
HdlTargets=zynq spartan6
HdlPlatforms=zed ml605
RccPlatforms=centos7 xilinx13_4
```

Note that the plural variable name must be used when multiple targets or platforms are specified. These variables may also be set in the environment.

Workers can be built for multiple targets with one command. This can be useful to check whether the source code is acceptable to all the different compilers.

The worker build process has make-style dependencies such that rebuilds will only happen if any dependent files are changed, including the OCS/OPS/OWD XML files.

13 Unit Testing of Workers

OpenCPI supports unit testing where a `<component>.test` directory in a component library is created to hold a test suite for all the workers in the library that implement the same spec (OCS). The workers that are tested could be written to different authoring models or languages or simply be alternative source code implementations of the same spec.

E.g. if a library contained `fft.hdl` and `fft.rcc` and `fft_xilinx_dsp.hdl` workers that all implemented the `fft-spec.xml` OCS file in the library's `specs` directory, a single `fft.test` directory would be created to hold a test suite that tested them all. Each `<component>.test` directory is associated with a single OCS, has a `Makefile`, and has a test suite description XML file, called `<component>-test.xml`.

The OpenCPI unit test framework manages multiple dimensions of worker testing, with automation to minimize test design and preparation efforts. The dimensions are:

- Test cases (individual parameterized tests, possibly using different runtime values)
- Execution platforms (HDL hardware and simulation platforms, RCC/OCL Platforms)
- Workers (different source code implementations, different models)
- Worker build configurations (compiled in vs. runtime settable property values)

The unit test framework allows complex test scenarios while providing layered complexity to keep simple test cases very simple to define and execute. Test inputs and outputs can be pre-prepared data files (i.e. test vectors), or be developer-provided scripts for input data generation and output data verification.

Unit testing in this unit test framework proceeds in five phases:

1. **Generate** — generate testing artifacts after finding the spec and the workers
2. **Build** — building HDL bitstream/executable artifacts for testing (for HDL workers)
3. **Prepare** — examine available *built* workers and available platforms, creating execution scripts to use them all for executing feasible tests.
4. **Run** — execute tests for all workers, configurations, test cases and platforms
5. **Verify** — verify results from the execution of test cases on workers and platforms

13.1 The Phases of the Unit Test Process

The **generate(1)** phase performs the following tasks automatically without any developer involvement:

- Discovers the OCS associated with this test directory.
- Discovers workers in the same library that implement that OCS.
- Discovers the build configurations (parameter values) for each worker.
- Derives a baseline for test cases based on all the build configurations that have been used on any worker.
- Derives the actual tests appropriate for all parameter combinations vs. the actual worker build configurations they apply to.
- Generates XML applications (OAS files) that perform unit tests on all workers
- Generates HDL test assemblies (subdirectories, Makefiles and OHAD files) that can be built for HDL platforms (hardware and simulation).

Although there are many options described below, the default generated unit test applications feed specified input data to input ports of the worker being tested and cause a zero-length message to be send to input ports after the input data is sent.

The generated test application is considered finished when:

- The component being tested has a single output port and a zero length message is received from that port.
- The component being tested has multiple output ports and a zero length message is received from the first one.
- A specified time duration for the test has been exceeded.

After the test completes, the output data from all output ports as well as the final values of all properties are available to the verification process for assessment.

All the above tasks in the **generate(1)** phase are done “off-line”, without regard to

- which worker artifacts are built
- which platforms are available on which to execute tests
- the availability of any build-related tools (compilers or FPGA synthesis tools or simulators).

The **build(2)** phase is only necessary for testing HDL workers, and builds the *generated* HDL test assemblies for whichever platforms (including simulators) are specified. When building for hardware HDL platforms, this phase takes the longest. There must be tools available to build the specified platforms, but this phase does not require those platforms be available for execution.

The **prepare(3)** phase does the necessary work to prepare to actually execute test cases and perform associated verification. In this phase the unit test framework automatically does:

- Discovery of available execution platforms, local and remote (reachable via network)
- Discovery of available built artifacts that can be executed on available platforms (including generated HDL test assemblies generated in the **generate(1)** phase)
- Generation of test scripts to perform all feasible tests on all available platforms.

After preparation, the developer invokes the **run(4)** and **verify(5)** phases. These phases can be sequential (all executions followed by all verifications), or interleaved (each test subcase is executed and verified before executing the next one).

During execution and verification, there are various filtering capabilities to subset which tests are run, which platforms should be tested, and whether a test case failure should immediately stop the testing process.

The rest of this section will describe:

- what is required and possible in the test description XML file:
`<component>-test.xml`
- what can and should be specified in the **Makefile** for this directory
- how to provide data generation and verification scripts
- how to execute test cases and verify their results

A `<component>.test` directory can be created and initially populated using this command:

```
ocpidev create test <component>
```

This will create the two necessary files: **Makefile** and `<component>-test.xml`.

13.2 Unit Test Concepts and Terminology

Here are the terms used in the OpenCPI unit test automation framework:

Test suite – as embodied in the `<component>.test` directory, is a suite of test cases for testing all workers implementing a spec across all available platforms for which the workers have been built

Test matrix —the virtual multidimensional space of testing, across:

Workers — the different source code or authoring model implementations

Worker configurations — different parameter value sets for worker builds

Initial property values — runtime property value configurations

Platforms — possible runtime environments

User-defined test cases — with property values, inputs and outputs

Test case – a parameterized test

Using a defined set of inputs or generation scripts

Using a defined set of outputs or verification scripts

Using a defined matrix of property values

Test subcase – a very specific test

Defined by and generated from a test case

Using a specific worker build configuration

Using a specific set of property values

Not bound to a specific platform or artifact

Generator — script to create input data files for ports or property value files

Called for a subcase, with all property values supplied.

Verifier — script to verify test output data produced by output ports

Called for a subcase, supplied with each port's output data file, with all property values supplied, both initial and final (volatile)

Viewer — script to view the results of a subcase execution. (e.g. plot).

Default Test Case — the case that is automatically created when none are specified

Tests all parameter combinations as derived from all worker parameter/build configurations or all workers

Developer can supply runtime property settings with multiple values for each, resulting in the cross-product of subcases

One generation script and one verification script, per port, parameterized by subcase property values

13.3 Unit Test Description XML File

The `<component>-test.xml` file specifies test cases and the defaults that apply to all test cases. As with all OpenCPI XML files, element names and attribute names are case insensitive. The top-level element in the file is `<tests>`, with the possible child elements being:

`<input>` to define an input file or generator script usable by any test case

`<output>` to define an output file or verifier script usable by any test case

`<property>` to define property values for all test cases

`<case>` to define a non-default test case when needed

If no `<case>` element is defined, the default test case is used. This is a common situation since the default test matrix is based on the parameters that workers are built with, and the available workers that implement this component spec. Here is an example file using the default test case for a component:

```
<tests>
  <input port='in' script='generate.py 16' />
  <output port='out' script='verify.py 16384 16' view='view.sh' />
  <property name='phs_inc' values='-4096' />
  <property name='enable' values='0,1' />
</tests>
```

It specifies that the default test case should be used (no `<case>` elements), the “`generate.py 16`” command should be issued to generate test data for port `in`, the “`verify.py 16284 16`” command should be issued to verify output data from port `out`, the `phs_inc` property should always be tested set to `-4096`, and the `enable` property should be tested with values 0 and 1. All scripts are run per subcase and have access to the parameter properties as well as the runtime properties of the subcase being tested.

Several attributes described below refer to scripts that will be executed by the unit test framework. In all cases, scripts must properly return process/shell exit status, with zero indicating success and non-zero indicating failure. This is true regardless of the language used in the script.

13.3.1 Attributes for the Top-level Tests Element

The valid attributes for the top-level `tests` element apply to all test cases and are `Spec`, `UseHdlFileIO`, `ExcludeWorkers`, `OnlyWorkers`, `ExcludePlatforms`, `OnlyPlatforms`, `Timeout`, and `Duration`. All are optional and are specified in special situations.

13.3.1.1 Spec Attribute of the Top-level Tests Element

Normally the spec (OCS) for all the workers being tested is inferred from the name of the `<component>.test` directory, and found in the file:

```
../specs/<component>-spec.xml
```

When this is not the case, this **spec** attribute can specify the name of the spec file for this test suite, much like the same attribute can be used in a worker's OWD.

13.3.1.2 *UseHdlFileIO Attribute of the Top-level Tests Element*

This boolean attribute applies only when HDL workers are being tested on simulation platforms. When true, it indicates that file I/O between the worker being tested and the input and output test files is done in the simulator using VHDL/Verilog file operations directly. When false (the default), the file I/O is being done by file reading and writing RCC workers running *outside* the simulator, with the data flowing in and out of the simulator. Both settings can be useful, but the **true** setting generally results in faster simulation times since less logic is being simulated for this file I/O.

13.3.1.3 *ExcludeWorkers Attribute of the Top-level Tests Element*

This string attribute specifies a list of comma-separated workers (e.g. **fft.hdl**) that should *not* be tested, even if they implement the spec of this test suite.

13.3.1.4 *OnlyWorkers Attribute of the Top-level Tests Element*

This string attribute specifies a list of comma-separated workers that should be the *only* ones tested. Any others found to implement the same spec will be ignored.

13.3.1.5 *ExcludePlatforms Attribute of the Top-level Tests Element*

This string attribute specifies a comma-separated list of platforms that should *not* be tested. Wildcard patterns may be used for any name. E.g. “***sim**” would exclude any platform that ended with the letters **sim**. Any other available platforms that have built artifacts will be used.

13.3.1.6 *OnlyPlatforms Attribute of the Top-level Tests Element*

This string attribute specifies a comma-separated list of platforms that should be the *only* ones tested. Wildcard patterns may be used for any name. E.g. “***sim**” would test only platforms that ended with the letters **sim**. Any other available platforms will be ignored.

13.3.1.7 *Duration Attribute of the Top-level Tests Element*

Normally unit test cases execute until the test subcase application is **done**, which is usually when an EOF is asserted on the first or only output port of the worker being tested. See the [EOF Assertion](#) section.

When this default EOF behavior is not viable or desirable (e.g. when the workers being tested have no output ports), the **duration** attribute can be set to an amount of time the application should run before being considered successfully **done**.

The duration value is in seconds, and represents wall clock time for the execution. After this amount of clock time, the execution is stopped, and the execution is considered successful and any outputs or final property values are provided for verification.

This attribute is separate from the `timeout` attribute, which indicates when the execution should be considered **failed**. Only one of these attributes may be set.

13.3.1.8 *Timeout Attribute of the Top-level Tests Element*

The `timeout` attribute indicates an amount of wall clock time in seconds after which the execution of a test subcase is considered a failure. Setting this value prevents an execution from hanging and preventing completion (but failure) of the subcase. This is especially useful for non-interactive scripted regression testing. Only one of the `duration` or `timeout` attributes may be set.

13.3.2 *Input Element of Top-level Tests Element*

An `<input>` element as a direct child of the top-level `tests` element specifies a source of input data that can be used by test cases. It is not specific to a test case but may be used by any test case for any input port. Its allowable attributes are: `name`, `port`, `file`, `script`, `messageSize`, `messagesInFile`, `suppressEOF`, and `stressorMode`.

13.3.2.1 *Name Attribute of the Input Element*

This optional string attribute specifies the name of this input source, so it can be referenced by test cases that use it, by name. If it applies to all cases, it doesn't need a name. If it applies only to a specific `port`, the `port` attribute can be set, which is more common. One of `name` and `port` must be specified.

13.3.2.2 *Port Attribute of the Input Element*

This optional string attribute specifies the name of the port that this input source will always apply to. If there is only one input source for a port, it will be used for all cases. One of `name` and `port` must be specified.

13.3.2.3 *Script Attribute of the Input Element*

This string attribute indicates a command to execute to produce data. When data is generated for a subcase and for a port, this command will be issued. The attribute value is not just the name of a file to execute, but of a command, so it can have a command name followed by some command arguments. When the command is executed in order to produce data, it will be appended with the name of the file to be written into; i.e. the script's job is to write into the file whose name is at the end of the command. Thus if the value of this attribute was:

```
echo hello >
```

then the source of data would always be a line of text containing `hello` since the actual command executed would be:

```
echo hello > <output-file-supplied-by-unit-test-framework>
```

The way these scripts become more useful is that all parameter and initial runtime property values are supplied to the script as environment variables. Thus this script is parameterized by these values for the subcase being generated. Accessing

environment values is easy for the scripts, whether they are written as shell scripts, python, or C. When a script is executed for a subcase (and for a port), the value of each parameter and runtime property is the value of an environment variable named:

`OCPI_TEST_<prop>`

So, if the property's name was `myprop`:

In C or C++, the value (as a string) would be: `getenv("OCPI_TEST_myprop")`

In python, it would be: `os.environ.get("OCPI_TEST_myprop")`

In bash/shell, it would be: `$OCPI_TEST_myprop`

Only parameter properties or runtime properties that are *initial* or *writable* are present. Using scripts based on these values normally means one script can be applied to all test cases.

The command is executed by the shell in the `<component>.test` directory, and must have execute permissions.

13.3.2.4 *File Attribute of the Input Element*

This string attribute specifies the name of a file to be used as the source of data. It is not affected by any property values and is thus a “constant”. This is useful if the same input data should be used for a port for all test cases, or if the file is not easily generated by a script, but is used for one test case.

13.3.2.5 *MessageSize Attribute of the Input Element*

This positive integer attribute specifies the size of messages to be supplied to the port of the worker under test when this data source is being used. Since data flowing between ports always consists of messages, this determines their size. The data from this input source is split into messages of this size, in bytes. This attribute is the same as the property of the `file_read` worker.

13.3.2.6 *MessagesInFile Attribute of the Input Element*

This boolean attribute indicates that the data produced by this input source has message boundaries and opcodes embedded in the data. Each message in the input file is preceded by a header consisting of two 32 bit unsigned integers (little endian), with the first being the length of the message in bytes, and the second being the opcode (with only the low-order 8 bits used). This attribute is the same as the property of the `file_read` worker and is only valid for inputs.

13.3.2.7 *SuppressEOF Attribute of the Input Element*

This boolean attribute indicates no “EOF” message (zero-length message with opcode zero) should be provided to the input port after all the data (messages) has been provided as input. This attribute is the same as the property of the `file_read` worker and is only valid for inputs. The default behavior is to provide such EOF messages.

13.3.2.8 *StressorMode Attribute of the Input Element*

This attribute only applies when testing HDL workers and is only valid for input ports. The generated test benches for HDL workers include a **stressor** function at each input port which varies the range of input behavior to provide coverage of what the worker might see at its port in applications. This optional enumeration attribute specifies the stressor behavior at this port. The term **metadata** is used here to indicate the signals indicating message boundaries and whether the boundaries are separate from or coincident with, the first or last data word of a message. These signals are **som**, **eom**, and **valid**.

There are four modes available: **bypass**, **throttle**, **metadata**, and **full**. The meanings are:

- bypass** — all messages are passed through to the worker under test without change or delay.
- throttle** — the metadata of all received messages is passed through to the worker under test without change, but data within the message will be withheld at random intervals to create intra-message data starvation.
- metadata** — the **stressor** at this port will cycle through a pattern of all valid metadata combinations, but the data will not be throttled.
- full** — the **stressor** will cycle through a pattern of all valid metadata combinations, data within the message will be throttled at random intervals to create data starvation, and idle cycles will be inserted between messages.

The default mode is **bypass**.

13.3.3 *Output Element of Tests Top-level Element*

An **output** element as a direct child of the top-level **tests** element specifies how the output data from a port of the worker being tested may be verified for correctness. The valid attributes are: **name**, **port**, **file**, **script**, **stopOnEOF**, **view** and **disableBackpressure**. It may be applied to all test cases, be used as a default for test case that do not mention an **output** element for a port, or be referred to by name by some test cases.

It is very similar to the **input** element:

- The **name** attribute allows this element to be referred to in test cases.
- The **port** attribute specifies that this element should be used for a particular port.
- The **file** attribute specifies an existing file to compare the output data to for correctness.
- The **script** attribute specifies a command that takes a file name as the data to *verify*.

13.3.3.1 *Script Attribute of the Output Element*

This attribute is similar to the script attribute of **input** and **property** elements. The major difference is that there are multiple arguments appended to the command instead of one. The first is an input file that contains the output of the given output port as a result of executing the worker in a subcase. After that first file name argument there are file name arguments for each *input* port of the component that contain the input data supplied to that port, in the order the ports are declared in the OCS. This allows the script to not only access the resulting output data from an output port, but also access the data supplied to each input port (if needed for the verification).

For example, if the component had input ports **in1** and **in2**, and an output port name **out**, and a script command **<command>**, in the script attribute, the actual command executed would be:

```
<command> <output-from-port-out> <input-to-in1> <input-to-in2>
```

A second important difference for the output script vs. an input script is that the final values of writable and volatile properties are available in the environment in addition to the initial values of all other properties. For generated properties (those with a **script** attribute in its **property** element), the name of the generated file is placed in the environment variable named **OCPI_TESTFILE_<property-name>**, while the final value is still in the **OCPI_TEST_<property-name>** environment variable.

The name of the test case is in the **OCPI_TESTCASE** environment variable and the name (which is numeric) of the subcase is in the **OCPI_TESTSUBCASE** environment variable. E.g. if the subcase being run was **case43.03**, the case name is **case43** and the subcase name is **03**.

As with all other scripts, a process/shell exit status of zero indicates success, while a non-zero exit status indicates failure. The script may write other informational messages about the failure to **stderr** which will be logged. The script should not write simple success and failure (PASS/FAIL) messages since the unit test framework does that already, using green/red colors for PASS/FAIL, based on the exit status.

13.3.3.2 *View Attribute of the Output Element*

This optional string attribute operates similar to the **script** attribute, but has a different purpose. It provides a convenient way for the developer to ask for a “view” of the data for the port. Taking all the same arguments as the verification script (in the **script** attribute), it is expected to present the data in some useful way during test development, typically in some viewing or plotting window.

13.3.3.3 *StopOnEOF Attribute of the Output Element*

This boolean attribute indicates that reception of an “EOF” message (zero-length message with opcode zero) should stop the recording of output data from an output port. The default value is true (which is unusual for a boolean attribute). This attribute is the same as the property of the **file_write** worker and is only valid for outputs. The reception of EOF messages on the first or only output port is used to indicate the

end of execution of a test subcase unless the **duration** attribute is set for the case (or as a top-level default for all cases).

DisableBackpressure Attribute of the Output Element

This boolean attribute only applies to output ports of HDL workers under test and specifies the **backpressure** (i.e. flow control) behavior the worker will see at this output port during testing. Setting **disableBackpressure** to true will disable the backpressure applied to this port. When false (the default) the test bench will apply random backpressure to this port.

13.3.4 Property Element of Tests Top-level Element

This element specifies the default set of values for a property for all test cases, unless overridden in particular test cases. When multiple values are specified, the implication is that subcases should be generated that test each of the specified values.

For parameter properties, where the potential set of test values is normally derived from the values used to build the workers being tested, the values specified in this element act as a filter or subset of those values, since no tests can be performed for parameter values that are not used in any worker's build configuration.

The allowable attributes for the **property** element are: **name** (required), **test**, **value**, **values**, **valueFile**, **valuesFile**, **generate**, **add**, **only**, and **exclude**. Exactly one of the **value*** attributes must be specified. The textual syntax for property values is used, as described in [Property Value Syntax](#).

13.3.4.1 Name Attribute of the Property Element

This required string attribute identifies a property defined in the OCS of the test suite or in the OWD for a worker being tested. The values specified in other attributes are applied to this specified property during testing (except when the **test** attribute is true – see below).

13.3.4.2 Value Attribute of the Property Element

This attribute specifies a single value to be tested.

13.3.4.3 Values Attribute of the Property Element

This attribute specifies a comma-separated sequence of values to be tested.

13.3.4.4 ValueFile Attribute of the Property Element

This attribute specifies the name of a file containing a single value to be tested. Multiple lines in the file are considered elements of a sequence or array value.

13.3.4.5 ValuesFile Attribute of the Property Element

This attribute specifies the name of a file containing multiple values to be tested. Multiple lines in the file are considered separate values to be tested. Multiple values can also be specified on a single line in the syntax of a sequence of values of the type

of the property. E.g., if the type is `Ulong`, the `valuesFile` file could contain a single line of `1,2,3,4` or four lines containing the four values.

13.3.4.6 *Generate Attribute of the Property Element*

This attribute specifies a command to execute to create a file containing a value to be tested. An argument is added to the command for the name of the file to be written. All parameter and initial runtime property values are available to the script as environment variables. This feature is convenient when a property value depends on others in a complex way. Note that expressions can be used in the value attributes of `property` elements, so scripts are not necessary to perform simple arithmetic based on other parameters.

13.3.4.7 *Test Attribute of the Property Element*

This optional boolean attribute, when true, indicates that this property is a **test** property that is *not* a property of the workers being tested. It is a property whose value is available to all input generation, output verification, output viewing and property generation scripts. Its name must be different than all property names in the OCS or in any of the workers' OWDs.

Values assigned to a test property are used to generate other test cases not defined simply by the values of worker properties. When this attribute is true, other data type attributes used in an OCS `property` element, such as `type` or `arrayLength`, may be applied to this `property` element since it is in fact *defining a property*.

13.3.4.8 *Exclude/Only/Add Attributes of the Property Element*

These attributes limit certain values to certain platforms. The **exclude** and **only** attributes provide value-specific restrictions on platforms similar to the effect of the top-level **excludePlatforms** and **onlyPlatforms** attributes: they *exclude* certain platforms from testing certain values or specify that some values should *only* be tested on certain platforms. The **add** attribute specifies that the values should be *added* to the set of tested values only for certain platforms.

The syntax of these attributes is the same as mentioned in the [ExcludePlatforms Attribute](#) section above.

13.3.4.9 *The set child element of the Property Element for Delayed Values*

Normally any non-parameter initial property values are set before the test case application is started (after **initialize** and before **start**). This means that such values are stable during the life of the test. In some cases it is useful to set writable property values *during* execution.

The **set** child element can be used to “schedule” the setting of a property value at some time *after* the application is started. Delayed property settings in applications are described in the **Application Development Guide**, and the syntax here is similar: the **set** child element can occur more than once, and have **delay** and **value** attributes to specify the delay after start (in seconds, floating point), and a value to set at that time

(relative to when the application is started). A delay of zero causes the value to be set immediately after the application is started. The following example:

```
<property name='myctl' />
  <set delay='1' value='10' />
  <set delay='1.001' value='20' />
</property>
```

sets the **myctl** property to 10 one second after start, and sets the value to 20 one millisecond later.

13.3.5 Case Element of Tests Top-level Element

The **case** element defines a non-default test case when required. It is necessary when the automatic parameterization of the default test case is insufficient for testing the worker(s).

The allowed attributes of a case element are: **Name**, **OnlyWorkers**, **ExcludeWorkers**, **OnlyPlatforms**, **ExcludePlatforms**, **Duration** and **Timeout**. All but **Name** have the same function as previously defined for the top level **tests** elements, but only apply to this case.

If **ExcludePlatforms** is set at the top level, setting it again at the case level adds to the list of platforms that will not be tested. Setting **OnlyPlatforms** at the case level if **ExcludePlatforms** is set at the top level is valid, but **OnlyPlatforms** at the case level cannot list a platform that has been excluded at the top level.

If **OnlyPlatforms** is set at the top level, setting it again at the case level narrows the list of platforms that will be tested further, but it is not valid to set **OnlyPlatforms** at the case level to a platform not in the top level **OnlyPlatforms** list. It is valid to set **ExcludePlatforms** at the case level if **OnlyPlatforms** is set at the top level, but excluding a platforms not in the top level list has no effect.

ExcludeWorkers and **OnlyWorkers** at the top and case levels interoperate in the same way as **ExcludePlatforms** and **OnlyPlatforms**.

The allowed child elements under a case element are: **input**, **output**, and **property**.

Each case can override or use the default inputs and outputs for each port, and each case can override the property values tested for each property. If no **input** or **output** is defined for an input/output port, then the default **input/output** is used (the **input/output** defined for the port under the top-level **tests** element). If no **property** element is present for a property under a **case** element, then the values defined at the top level are used. For parameter properties, the default values tested are derived from the values defined in all the workers' build configurations, but this automatic default can still be overridden (limited) by a **property** element at the top level or under a **case** element.

13.3.5.1 *Name Attribute of Case Elements*

This optional string attribute specifies the name of the test case. If not present, the name of the case is **case** followed by a case number starting at zero, with at least 2 digits (i.e. the second case would be **case01**, and the 101st case would be **case100**). The name of a case is listed in various reports, and can be used when specifying that only certain cases (rather than all cases) should be executed or verified.

13.3.5.2 *Input Element under Case Elements*

This element specifies how input data is generated for a port, in a test case. If not specified, the default input source for the port specified at the top level is used. The **port** attribute of the **input** element specifies the **port** this input element applies to. The **name** attribute, when present, indicates that a specifically named input source defined at the top level should be used for this port. If the named input source at the top level already has a **port** attribute, no **port** attribute need be supplied for this **input** element.

When the **name** attribute is specified, none of the **file**, or **script** attributes are allowed. If there was a top-level **input** element like this:

```
<input name='pulsegen' port='in' script='mygen.py' />
```

then a **case** element could simply have:

```
<input name='pulsegen' />
```

Similarly, if the top-level **input** element was this (with no **port** attribute, allowing it to be used for different ports):

```
<input name='pulsegen' script='mygen.py' />
```

then a **case** element could have:

```
<input name='pulsegen' port='in' />
```

13.3.5.3 *Output Element under Case Elements*

The **output** element for a test case acts the same as the **input** elements. They refer to a named **output** element at the top level or override the default, per port.

13.3.5.4 *Property Element under Case Elements*

The **property** element for a test case acts the same as the **property** elements at the top level: it specifies values to be used for the named property for this test case. If a property is not mentioned in a **case** element, the default top level values are used.

A single test case can have multiple values for any property. Subcases are automatically generated for all combinations of property values specified for the test case whether specified at the top level as default sets of values, or specified for the test case in the **case** element.

Individual values for a property for a test case may be restricted to certain platforms as described above in the [Exclude/Only/Add Attributes](#) section.

13.4 Unit Test Makefile Contents

The Makefile in a `<component>.test` directory is normally untouched after being created with the `ocpidev create test` command. It is generated to contain only one line:

```
include $(OCPI_CDK_DIR)/include/test.mk
```

Several make variables can be used either on the make command line or specified in this Makefile to control the various phases of unit testing.

During the **build** phase, as with building projects, libraries and workers, these platform variables are applied: **HdlPlatform(s)**, **OnlyPlatform(s)**, **ExcludePlatform(s)**.

During later phases (**prepare**, **run**, **verify**) these platform variables apply: **OnlyPlatform(s)**, **ExcludePlatform(s)**.

The **view** variable can be set to 1, which will cause the “view” script to be run whenever verification is requested.

The **TestVerbose** variable can be set to 1, which will cause the execution and verification logs to be included in the console/shell output, rather than just placed in specific log files, per platform and per subcase.

The **KeepSimulations** variable can be set to 1 to cause the contents of the **simulations** directory to be retained after successful executions on simulation platforms. Successful verification for a platform normally causes the associated **simulations** directory to be removed immediately. Keeping simulation output may use lots of file system space (100s of GBs in some extreme cases).

The **TestAccumulateErrors** variable can be set to 1 to cause execution or verification errors to accumulate (i.e. not immediately stop the run and/or verify process) and simply be reported as they occur. The default is to stop on the first error either in execution or verification. When accumulating errors, if any test executions or verifications fail, the entire process will return an error condition after all cases have run on all platforms. There is no option that will accumulate errors across multiple test directories in a library or project.

The **Cases** variable is a wildcard pattern indicating which cases/subcases should be executed or verified. Subcases are named `<casename>.<subcase#>`, so this variable may be set to patterns that affect certain cases or subcases. Subcase numbers are listed in the report in `gen/cases.txt` report produced by the **generate** phase. The default case name is `case00`. For example, to only run subcases that end in 3, you could specify: `Cases='case*.*3'`. Multiple patterns are allowed, such as: `Cases='case00.01 case00.03'`. The quotes are not necessary in the Makefile, but are necessary on the command line.'

13.5 Preparing Unit Test Inputs

An `input` element must be specified for each input port of the component, either at the top level as a default for all cases, or for specific cases. It either specifies a pre-existing data file to use (using the `file` attributes) or a command to execute to generate the input data which can depend on the property settings for the specific subcase (using the `script` attribute). While a specified input file applies to all subcases regardless of property settings, a generator script can generate input data for each subcase that depends on all its property settings.

The format of the data in the (possibly generated) file is a series of message payloads as defined by the protocol for the port, as described in [Message Payloads on Data Ports](#). The data must be laid out according to the setting of the `ocpi_endian` built-in parameter property, whose value is available to all data generation scripts. All platforms currently supported use only little-endian data layout, but to test a worker that might be built for different endian systems, the layout of the data must match this parameter value.

The generator scripts for input ports are run for each subcase, with all property values for the subcase available to the script. The script is responsible for writing a file whose name is provided on the command line. Since the script command is executed as it would be on shell command line, it can be written in any language, such as python, bash, or even compiled C or C++. It is executed in the context of the development system (not the target, potentially embedded system), so it can depend on any tools installed on the development system. However, scripts that depend on tools not installed or required as part of OpenCPI will make the project as a whole less portable.

An example input generator script written in the python language is below for a FIR filter component. The script depends on two properties. The first `COEFF_WIDTH_p` is a parameter specifying the bit-width of samples. The second `NUM_TAPS_p` is the number of taps in the filter. The script generates an impulse, with a maximum value followed by zeroes. The file is binary 16 bit signed fixed point data.

```
#!/usr/bin/env python2
import sys, os, struct
max_tap = pow(2,int(os.environ.get("OCPI_TEST_COEFF_WIDTH_p"))-1)-1
num_taps = int(os.environ.get("OCPI_TEST_NUM_TAPS_p"))
fo = open(sys.argv[1], 'wb')
for j in range(num_taps):
    fo.write(struct.pack('h', max_tap))
    for i in range(1,num_taps*2):
        fo.write(struct.pack('h', 0))
```

If the script was in the local file `generate.py` and made executable (e.g. with `chmod a+x generate.py`), and the input port was named `in`, then the input specification that used this script would be:

```
<input port='in' script='generate.py' />
```

An input generation script must return exit status of zero/non-zero for success/failure.

13.6 Preparing for Unit Test Output Verification

An **output** element must be specified for each output port of the component, either at the top level as a default for all cases, or for specific cases. It either specifies a pre-existing file to compare test output data against (using the **file** attribute), or a command that examines the data produced at an output port to decide whether the execution of the subcase was successful (using the **script** attribute). Output verification scripts have access to the output data produced, the input data provided to all input ports, and the final values of all properties at the end of execution (in the environment).

If the component had input ports **in1** and **in2**, and an output port name **out**, and a script command **<command>**, in the script attribute, the actual command executed would be:

```
<command> <output-from-port-out> <input-to-in1> <input-to-in2>
```

The three filename arguments would be added by the unit test framework to run this output verification script for a given subcase, providing the file names for the data associated with the subcase (input and output). The script would run in the development environment and not in the environment of a potentially embedded target platform.

As with input data, the message payload formats must comply with the lay out as described in [Message Payloads on Data Ports](#), and also respect the value of the built-in **ocpi_endian** parameter property.

An output verification script must return exit status of zero/non-zero for success/failure.

13.7 Off-line One-time Tasks Prior to Test Execution and Verification

After creating the test suite in the `<component>.test` directory using the `ocpidev create test` command, the following steps are taken prior to running any tests.

- Making any necessary changes to the **Makefile** (rarely needed)
- Adding the elements (**input**, **output**, **property**, **case**) to the `<component>-test.xml` file.
- Prepare any input data files or input data generator scripts.
- Prepare any output data files (for comparison) or output verification scripts.
- Run the **generate** phase (see next section)
- Examine the report created in `gen/cases.txt` to see the generated subcases.
- If any workers are HDL workers, run the **build** phase to build the bitstream/executables.

After these steps, all applications, HDL assemblies, input data sets and verification scripts have been generated (in the **gen** directory) and any required HDL assemblies have been built for the desired HDL platforms. The **generate** phase does not depend on which platforms any of the workers being tested have been built for. Prior to the **build** phase no compilers or other build tools are required or used. The **build** phase does require any HDL workers to have been built for the desired platforms.

The **generate** phase is accomplished using the

```
make generate
```

command, and the build phase is accomplished using the

```
make build
```

command (or simply **make** with no goal specified). When the **build** phase is invoked, the **generate** phase may be re-invoked based on **make** dependencies. As mentioned above, the platform variables **HdlPlatform(s)**, **OnlyPlatform(s)**, and **ExcludePlatform(s)** may be used on the command line for building to modify defaults specified in the project, library or environment.

13.8 Testing on Remote Systems

The default behavior of the unit test framework is to test on locally available platforms. This usually means using the local CPU and any locally attached GPUs or FPGAs. However, there are two methods to including non-local, network-reachable systems in the set of available platforms for testing. The first is specific to unit testing and involves defining possible remote systems using the `OCPI_REMOTE_TEST_SYSTEMS` environment variable. This method is described in the next section.

The second method of using platforms on remote systems is to use the **remote container** capability that generally supports OpenCPI applications executing on a mix of local and remotely accessible platforms. This feature is fully described in the **Application Development Guide**. In order to have the unit test execution framework consider remote systems as test platforms, the remote and local systems have to be enabled as described in that document.

To enable the unit test framework to discover and use such systems, you set one environment variable: `OCPI_ENABLE_REMOTE_DISCOVERY=1`. This enables all such systems to be included in the set of available platforms for testing. This variable can also be set in the Makefile using: `export OCPI_ENABLE_REMOTE_DISCOVERY=1`.

Using remote containers requires no common network mounts between the local and remote systems and no ssh enablement, but has some firewall-related limitations in some environments.

13.8.1 Defining Remote Systems for Executing Tests

In order for the unit test framework to execute tests on platforms that are not available on the development system, remote systems with additional platforms must be specified. Such platforms are not discovered automatically but are specified in the `OCPI_REMOTE_TEST_SYSTEMS` environment variable.

Remote systems are accessed using the `ssh` remote execution command, with an `ssh` server capability required to be enabled on the remote system. This environment variable is a colon separated list of remote system specifications, and each remote system is specified by 5 fields separated by '='. The four fields are:

1. Remote Host name/IP address
2. SSH user name
3. SSH password (yes, this is not a secure solution)
4. Project directory mount path as seen on remote system
5. SSH Version (optional)

The remote system must meet the following requirements:

- SSH access from the development host, using the username and password in the first three (and optional fifth) fields. If the remote system is set up for public/private key access control from the development system, the password is not used, but must still be non-empty.

then, after a successful SSH login from the development system

- The project's directory on the development system must be mounted (NFS or equivalent) for access from the remote system, using the path of the fourth field.
- The OpenCPI kernel driver must be loaded on the remote system
- The `OCPI_CDK_DIR` environment must be set up properly consistent with the development host. An OpenCPI CDK installation is assumed, with the remote system and development system using the same OpenCPI release.

The project directory mount, the OpenCPI CDK installation, and the kernel driver may either be established at boot time or at SSH login time on the remote system.

A remote system provides may support multiple platforms (e.g. both HDL and RCC). Whatever platforms are available on the remote system will be discovered when the remote system is contacted. Like local platforms, these discovered platforms are subject to the filters specified by `OnlyPlatforms` and `ExcludePlatforms`.

Remote systems are frequently embedded systems which do not host a development environment, but any system can be a remote system. E.g. if the development system is running CentOS7, the remote system could be CentOS6 to run tests on that system also. Of course the CentOS6 system will only run RCC artifacts build for CentOS6.

Remote systems may be defined in the project's `Project.mk` file so that they are available for all test suites in the project, e.g.:

```
export OCPI_REMOTE_TEST_SYSTEMS:=10.0.1.16=root=root=/mnt/myproj
```

13.9 On-line Tasks for Test Execution and Verification

After the off-line steps described above, and any remote systems are defined, there are four steps that relate to actually executing tests and performing verification, and these steps are aware of available local and remote platforms on which to execute tests.

13.9.1 Preparing for Execution: Discovery and Execution Script Generation

The **prepare** phase is invoked by the **make prepare** command. It considers all available platforms, local and remote, including available RCC, HDL hardware and HDL simulators. The **OnlyPlatform(s)** and **ExcludePlatform(s)** make variables are used to filter the set of available platforms.

This step also considers which RCC built artifacts are available in the component library as well as which HDL bitstream/executable artifacts have been built locally from generated HDL test assemblies in the **build** phase. From the combination of available platforms and available artifacts it determines which subcases can be run on which platforms, and generates execution scripts accordingly. These execution scripts are generated in the **run** subdirectory of the test suite. The **gen** subdirectory captures the results of the **generate** and **build** phases, and the **run** subdirectory captures the results of the **prepare** and **run** phases. Both subdirectories are removed by **make clean**.

When determining available artifacts for test execution, it does *not* look at the prevailing setting of the **OCPI_LIBRARY_PATH**, but forces a new environment value that limits the artifact search to local artifacts in the component library, the **gen/assemblies** directory and the installed CDK on the local and remote systems.

After determining available and appropriate test execution platforms, this phase generates per-platform scripts that run all feasible subcases on that platform. These scripts are placed in a subdirectory of the **run** directory named after the platform.

13.9.2 Executing Tests on Available Platforms

The **run** phase executes, for each available platform, filtered by **OnlyPlatform(s)** and **ExcludePlatform(s)**, test applications for all possible cases and subcases.

Three different make goals can be used to invoke the **run** phase:

make runonly — perform the **prepare** phase and then the **run** phase.

make runnoprepare — manually rely on a **prepare** phase that has already be performed, and only perform the **run** phase. This has the benefit of avoiding some setup overhead and time, but runs the risk of something changing in the environment (e.g. a remote system becoming unavailable).

make run — perform the **prepare**, **run** and **verify** phases all together, with the **run** and **verify** steps interleaved by subcase with each subcase being verified immediately after it has been executed.

The **run** phase iterates through platforms (sequentially), executing all subcases on each platform in turn, and recording the results for later verification. The recorded results are:

- the output data from output ports
- the final values of all properties, including volatile ones
- a log of the actual execution of the tests.

All these results are recorded in the platform's subdirectory of the **run** directory with different files for different subcases. After the execution of each test on a given platform, a console message will indicate whether the execution succeeded or failed. If an execution fails, the execution of all tests for the platform is stopped, and the outputs can be examined for the cause of the failure. In this failure case other platforms will still be tested. If the **TestAccumulateErrors** variable is set to 1 (in the **Makefile** or on the **make** command line), such failures will be reported, but the execution of test cases will continue. This variable also applies to during verification of test cases. Setting this variable to 1 is useful to solve each problem as it occurs, while leaving it unset allows all errors to occur and be analyzed or examined later.

13.9.3 Verifying Test Results

The **verify** phase relies on the previous execution of appropriate subcases for all platforms, and performs verification using the results recorded previously in the **run** phase.

The **make verify** command performs verification for all platforms. The **make run** command will perform both **run** and **verify** phases in an interleaved mode where each subcase is verified after it is executed in order to show the results for subcases as they are executed.

The **verify** phase by itself does not involve any execution or access to local or remote execution platforms and thus can be performed offline. Whether the **verify** phase is executed with the **run** phase or by itself, the view option will run the defined view scripts with verification. The view option is enabled by specifying **view=1** on the **make** command line, and happens per subcase before each verification. When view scripts typically put up a window to display the data, they may wait for user input or simply return, allowing the data to be displayed while verification proceeds.

If the **TestAccumulateErrors** variable is set to 1 (in the **Makefile** or on the **make** command line) the **verify** phase will immediately stop and return an error if any subcase fails by returning non-zero exit status. If the variable is not set to 1, all subcases will be verified with failures reported as they occur. The output of the verify phase makes it clear whether each subcase **PASSED** or **FAILED**. The individual verification scripts are required to return a standard shell command exit status of zero for success and non-zero for failure.

13.9.4 Viewing Test Result Data

The **make view** command can be used to specifically run any view script defined for output ports. If the **view** variable in the **Makefile** (or on the **make** command line) is set to 1, this will always happen as a side effect of running the test, but when it is not, **make view** can be used to view results on demand (assuming a view script was defined).

13.10 Summary of Make Goals and Variables

The **make** command is used with the following goals to invoke the phases:

- build** — (the default goal if none is specified) build locally generated HDL assemblies, implies **generate** if needed
- run** — perform execution and verification, interleaved per subcase
- clean** — clean all generated, built and execution directories and files

The following goals control the process at a finer granularity:

- generate** — perform all offline generation tasks: scripted input data and property values, applications, HDL assemblies, generation report
- prepare** — discover platforms and artifacts and generate execution scripts per platform, assumes required artifacts are built
- runnoprepare** — execute tests assuming a previous prepare was done
- runonly** — prepare and execute tests, but do not perform any verification
- verify** — perform verification for whatever platforms have been executed
- view** — run the view scripts and put the results in separate windows
- cleanrun** — clean all run results
- cleansim** — clean all simulation output (but not run results)

The **View** and **KeepSimulations** variables apply to verification.

The **Case(s)** variables apply to both execution and verification.

The **OnlyPlatform(s)** and **ExcludePlatform(s)** applies to preparation, execution and verification.

The following make goals can be executed in a component library's directory to apply to all **<component>.test** directories in the library, or at the project level to apply to all component libraries in the project: **test**, **cleantest**, **runtest**, **verifytest**, **runonlytest**. They perform the same function as the corresponding goal in the **<component>.test** directory, without the **test** suffix. The **cleanrun** and **cleansim** goals can also be used at the library and project levels without a **test** suffix.

14 Developing OpenCPI Assets in *Projects*

In OpenCPI a **project** represents a work area in which a variety of assets are created and developed. Projects can contain all types of assets that are described fully either in this document or in others. A project can contain:

- Component libraries with specs and workers.
- Applications (described in the ***Application Development Guide***).
- HDL primitives and assemblies (described in the ***HDL Development Guide***)
- HDL devices, cards, slots (described in the ***HDL Development Guide***)
- Platform support assets (described in the ***Platform Development Guide***)

A project is a standard directory structure that holds the various OpenCPI assets in both source code form and built form, along with the makefiles that describe how they are built. The project structure provides a means to bundle a collection of assets which may have a logical relationship or be created for a specific application.

The ability to develop assets outside of a project (a.k.a. *standalone*) is also supported, but is preliminary and subject to change, and not discussed further in this section.

The `ocpidev` tool is used to create and then populate a project directory structure with the various asset types. The created skeleton directory structure is always buildable.

The structure of a project, and types of assets (shown enclosed in <>), is shown in the following diagram (with the makefiles and other files omitted except at the project level).

```
Project.mk
Project.exports
Makefile
applications/<applicationXYZ>/
specs/
components/<componentlibXYZ>/<workerXYZ>/
                                /specs/
---OR---
components/<workerXYZ>/
                                /specs/
hdl/primitives/<primitiveXYZ>/
hdl/assemblies/<assemblyXYZ>/
hdl/platforms/<platformXYZ>/
hdl/devices/<device-workerXYZ>/
                                /specs/
hdl/cards/<card-device-worker>/
                                /specs/
```

The optional top level `specs` directory is separate from the `specs` directory in any component library. It is a project-wide `specs` directory that is usable by all component libraries in the project. It can exist in the project, for use by other projects, even if there are no component libraries in the project.

Creation of a project (using `ocpidev`) creates a skeleton directory structure that is buildable, but it will build nothing initially as it contains no assets. All the intermediate directories are created by `ocpidev` as needed. If there are any component libraries created in the project (using `ocpidev create library <libname>`) a `components` directory is created, under which those component libraries will be created. Alternatively, for simpler projects which only have a single component library, the `components` directory *is* the single component library (created using `ocpidev create library components`) and workers are created directly under `components`. A project currently cannot easily be changed between the “flat”, single library structure and the “hierarchical”, multiple library structure.

When creating a worker, if no library is specified and the current directory is not in a component library already, it is placed directly in the `components` directory. It is an error to create such workers if component libraries already exist under `components`. Conversely, it is an error to create a component library under `components` if workers already exist there or `components` was already explicitly created.

In addition to the various directories, three required files are generated at the top level when the project is created by `ocpidev`:

Makefile: the top level Makefile for the project which supports building all assets from the top level project directory.

Project.mk: the make file fragment that defines make variables and options that are used *project-wide*, for all assets at all directory levels.

Project.exports: a file that specifies which assets and files should be visible from outside the project, i.e. visible to other projects which use some of the assets in this project.

These files are automatically created when the project is created, but may be edited later as necessary. The ***Project.mk*** and ***Makefile*** files must exist; the ***Project.exports*** is optional and created and edited manually.

14.1 Managing Project Assets.

The `ocpidev` tool described in detail later is used to manage all the asset types in a project. It is used to create, delete or test assets. Once created, the assets are based on text files that must be edited. Assets are created using the `ocpidev create` command and they are deleted using the `ocpidev delete` command. Most assets, including projects themselves, are based on a directory, with a **Makefile** in that directory. These include:

- Component libraries
- Workers
- Unit Test Suites
- HDL device workers
- HDL platforms
- HDL primitives (cores and libraries)
- HDL assemblies
- Applications (except the simplest ones)

When an asset is created, the appropriate directories are also created, an initial **Makefile** is created in the new directory, and in some cases other initial files are also created. The **Makefile** contents indicate which type of asset is in that directory.

Some assets are simply files and when created, an initial version of the file is created in the appropriate directory in the project. This type of asset includes:

- Component Specs
- Protocol Specs
- HDL card definitions
- HDL slot definitions
- Applications (simple XML based applications with no ancillary files)

When creating specs, protocols, workers and unit tests, a library option (`-l <library>`) may be supplied to `ocpidev` indicating which component library the asset should be added to. If the project has a single library in the `components` directory, this option is not used. For hardware-specific HDL workers, the `-h <library>` option specifies a directory under the project's `hdl/` subdirectory where the device worker should be created (usually `hdl/devices`).

When adding a platform-specific device worker or device proxy, a platform option (`-p <platform>`) may be supplied to indicate which platform-specific device library to add the device worker to. Portable device workers that are *not* platform-specific do not use this option.

14.2 Package IDs

A **package-ID** is the globally unique identifier of an OpenCPI asset. A project's package-ID is used when it is depended on by other projects. A component's package-ID is used to reference it in applications or workers. While all assets have package-IDs (either explicitly specified or inferred from the directory structure), only certain assets are actually identified by their package-IDs in the current OpenCPI release. Eventually, package-IDs will be used uniformly and universally.

A package-ID is a hierarchical sequence of names separated by periods. This OpenCPI package naming scheme follows the Java package naming conventions, which are roughly based on a reversed Internet domain name with the top-level domain first, and the more specific or local names after that. Thus if CNN had an OpenCPI project full of news-related components, its package-ID might be `com.cnn.news`.

Warning: When changing the `PackageName` or `PackagePrefix` of an existing project the project needs to be both unregistered and re-registered then cleaned and rebuilt. This also includes cleaning and rebuilding any projects that depend on this project.

14.2.1 Package-ID Usage and Structure

Package-IDs are divided into two parts, the **package-prefix** and the **package-name**. The package-prefix is the string before the last period, and the package-name is the name after the last period. Generally the package-ID of an OpenCPI asset consists of the package-prefix from its containing environment (e.g. the project of a library, or the library of a worker). Its package-name is simply the name given when the asset was created (using `ocpidev`). Thus the package-prefix acts as the name scope for the asset. This package naming scheme (from Java) provides for globally unique identifiers that are human readable (and type-able) and already administered by a globally known organization. A package-ID consists of `<package-prefix>.<package-name>`.

OpenCPI defines two reserved top level package prefixes: `ocpi` and `local`. The `ocpi` prefix is managed by the OpenCPI maintainers and is used for assets that are located at the OpenCPI github repository. Projects that are maintained there all have names with this prefix (currently the three projects in use are `ocpi.core`, `ocpi.assets`, and `ocpi.inactive`). The second reserved prefix (`local`) is the default packageID for all projects that are created with no explicit package-ID. This means that basic development projects (or tutorial examples) are not required to have a global identity to simply get started. As soon as other projects need to use or depend on a project, it should have a more explicitly assigned package-ID.

The package-ID of a component specification is generally prefixed with the package-ID of its containing library or project (followed by a period). This prefix is the name scope in which the component is defined. This allows components to be specified and implemented by different organizations, while still allowing any implementation found in a library to satisfy any (other) organization's component specifications. E.g., my project can have an additional, alternative implementation of a component specified in another organization's library, or can define its own specification for a component with the same package-name with a different package-prefix (name scope).

14.2.2 Package-ID attributes in XML files.

There are several cases where attributes are set in XML files to help specify package-IDs. In an application XML file, the top level **package** attribute provides a default package-prefix for all components mentioned in the file (that do not have periods in their name).

A component spec XML file can specify a package-prefix in a top-level **package** attribute. This would override the default, which is the package-ID of the library in which the spec is defined. If the spec is not in the library, but rather in the top level specs directory of a project, this would override the default prefix from the project's package-ID.

14.2.3 Package-ID Variables in Makefiles

The package-ID of an asset is normally inferred from the name it was created with serving as package-name with the package-ID of its containing asset serving as package-prefix. When this inferred package-ID is not what is needed, certain variables in the asset's **<Asset-type>.mk** file can override this default behavior. The **PackagePrefix** variable can override the default package-prefix supplied from the package-ID of the containing asset.

The default package-prefix for projects is **local**. As soon as this default is not what is wanted, the **PackagePrefix** variable can be set in the project's **Project.mk** file in the top level project directory. Similarly, a component library in a project could override its prefix (which would normally be the package-ID of its project) using the **PackagePrefix** variable in the library's **Library.mk** file.

In the unusual case where the package-name should be different than the name the asset was created with (which is also the name of its directory), the **PackageName** variable can be used to override that part of the package-ID.

When it is required to simply set the entire package-ID for an asset, the **Package** variable (which should be named **PackageID**) can be used. This overrides any logic for combining the package-prefix and the package-name.

14.2.4 How to Determine an Asset's Parent (and thus Default Package-Prefix)

The parent of a component library, any of:

```
components/  
components/<library>  
hdl/devices/  
hdl/cards/  
hdl/adapters/  
hdl/platforms
```

is the project itself.

An HDL platform's parent is the **hdl/platforms** directory.

A **hdl/platforms/<platform>/devices** library's parent is the containing platform.

A project has no parent, and so its package-prefix defaults to **local**.

The default package-prefix of a component (which does not have its own directory) is `<parent's package-ID>.<component-name>`.

14.2.5 Package-ID examples

Within the `ocpi.assets` project the `Project.mk` file contains:

```
PackagePrefix=ocpi
```

Thus the package-ID for the project is: `ocpi.assets`.

The `dsp_comps` library in the same project (located at `assets/components/dsp_comps`) has no variables set in its `Library.mk` file, so:

- The package-prefix defaults to package-ID of parent (project): `ocpi.assets`
- The package-name defaults to library's directory name: `dsp_comps`
- Thus the package-ID is: `ocpi.assets.dsp_comps`
- If we had `Package=full_package` in the file `dsp_comps/Library.mk`, the package-ID of the library would be, entirely, `full_package`

The assets project has an HDL platform (`matchstiq_z1`) with its own devices library in:

```
assets/hdl/platforms/matchstiq_z1/devices
```

- With no package variables set, the package-name defaults to the directory: `devices`
- The package-prefix defaults to the package-ID of parent (the `matchstiq_z1` platform), which is: `ocpi.assets.platforms.matchstiq_z1`
- Thus the package-ID for the library is:
`ocpi.assets.platforms.matchstiq_z1.devices`

The `fir_real_sse` component spec in assets project's `dsp_comps` library would be referenced in an application as:

```
ocpi.assets.dsp_comps.fir_real_sse
```

This is because a spec's parent is the library it is in.

14.3 The Project Registry: How Projects Depend on and Find Each Other

Every project depends on some other projects. At a minimum all projects depend on the core project provided as part of OpenCPI. Its package-ID is `ocpi.core`.

When assets in a project depend on assets in other projects, the project must declare that it depends on such other projects. Here are some examples of how assets can depend on other assets, and thus how an asset in one project may depend on assets in other projects.

- A worker may depend on a spec (OCS), whether in the same library as the worker, a different library in the same project or in a different project.
- An HDL assembly may depend on workers in any library in the same project or workers from libraries in other projects.
- An HDL worker may depend on an HDL primitive library in the same project or in a different project.

A project declares its dependency on another project by either specifying the dependency when the project is created (using `ocpidev`), or adding the dependency to its `Project.mk` file later. In either case the dependency is declared by providing the package-ID of the other project. As mentioned above, the dependency on the `ocpi.core` project is always assumed.

A project is visible to other projects by being **registered** under its package-ID. A project can only depend on another if that other project is registered with the OpenCPI installation. The default condition of a project, when created, is to be **unregistered** and thus not visible to others. When a project is in a suitable condition to be depended on by others (i.e. it contains assets that are ready to be used by other projects), it can be registered at that time (using the `ocpidev register` command). This action can be reversed using the `ocpidev unregister` command.

14.3.1 Projects That Implement Platforms

While platform assets (sometimes called *platform support packages*) are one of many types of OpenCPI assets, they are special in two ways:

- they enable other assets to be built (compiled, synthesized, etc.) targeting their platform
- when they are in a registered project, they are usable for asset building for *all other projects* regardless of whether those *other* projects declare a dependency on the project containing the platform asset

A project A does not have to explicitly depend on project B simply to build A's assets targeting a platform implemented in project B. This allows building for platforms without knowing which project they are implemented in. For any other dependency between assets in one project and another, the project dependency must be declared when the project is created or later in the project's `Project.mk` file.

14.3.2 Project Registries for Sandboxed Project Development

Projects depend on other projects, and make themselves visible to other projects via a **project-registry** that is part of the OpenCPI installation. The registry includes a mapping between package-IDs and the pathname of registered projects. This mapping is shown using the `ocpidev show registry --table` command.

There are two different relationships between projects and registries:

- Every project is automatically **associated with** a project registry upon creation and thus can view and depend on other projects registered there.
- A project can be **registered in** its associated project registry, and thus be visible to other projects associated there. Projects are not initially registered.

In cases when these three things are true:

- multiple versions of some projects are developed simultaneously
- those projects are depended-on by others
- the different project versions all use the same OpenCPI installation

it is useful to create an additional registry as a sandbox for that set of copies or versions of projects. To create or delete an alternative registry, the `ocpidev create|delete registry` command can be used. This establishes a registry in a user-specified directory. To use such an alternative registry as the new default or to use platforms defined and/or built in projects there (e.g. HDL simulators), the `OCPI_PROJECT_REGISTRY_DIR` environment variable must be set point to it. This overrides the default registry that is part of the OpenCPI installation.

Any projects created with this environment variable set will be associated with the alternative registry. This association of a project to its registry is persistent. This association does not depend on the current setting of the environment variable. To change an existing project's association with its registry, the `ocpidev set registry` command can be used. It can either specify an alternative registry directly or if none is specified, the default (possibly from the environment variable `OCPI_PROJECT_REGISTRY`) will be used. The association can be undone using the `ocpidev unset registry`, in which case the default registry will be used when needed (and that new association will persist).

So there are four uses of the relationship between projects and a registries:

1. A project's *associated* registry, set at creation or by `ocpidev set registry`, is the way a project *depends* on others. Normally this association is automatic and unchanged.
2. A project being *registered in* its registry, is the way for *other* projects, associated with the same registry, to depend on it.
3. A project depends on another *explicitly*, by the declaration in its `Project.mk` file, enabling assets in one project to depend on assets in other projects.
4. A project can be *built* for platforms implemented by any registered project in its associated registry, without an explicit dependency.

14.4 Project Makefiles

Most of the directories in a project contain a file named **Makefile**, which is created by **ocpidev** and used when building all the assets in the project. The **ocpidev** tool is used to actually build, clean and test assets, using the **build** or **clean** verb. The **Makefile** does not need to be examined or edited in most cases.

The generated **Makefile** has the same form for all asset types, setting optional **make** variables, and including a standard makefile fragment from the OpenCPI CDK. The **Makefile** can be left unmodified, with the default behavior usually being adequate and appropriate or can subsequently be edited by the developer to specify additional optional variable settings or add customized **make** targets etc in advanced cases. Each **Makefile** has an initial single line of content, and for the top-level **Makefile** in a project, it is:

```
include $(OCPI_CDK_DIR)/include/project.mk
```

This **project.mk** file included from inside the CDK is *not related* to the **Project.mk** file mentioned above.

When a **Makefile** is in a directory with a number of subdirectories for the same type of asset, a variable can be set which lists the assets to build. This variable is optional, and when not specified, all such assets are built. For example, in a component library where all the subdirectories contain workers, the default Makefile is simply:

```
include $(OCPI_CDK_DIR)/include/library.mk
```

This implies that all worker subdirectories should always be built. If there are workers that should *not* be built, or they should be built in a particular order, then the **Workers** variable can be specified to list the explicit set of workers that *should* be built, in order, e.g.:

```
Workers=fft.rcc fft.hdl  
include $(OCPI_CDK_DIR)/include/library.mk
```

Alternatively, if there are a few workers that should not be built, the **ExcludeWorkers** variable can list them, in which case all workers *except* those listed will be built.

This same idea applies to directories in a project that contain HDL assemblies, HDL platforms, HDL primitives, applications, devices, cards, and slots, etc. The exact name of this variable, and other optional variables, are described in the section for each asset.

The makefiles for directories that hold multiple assets of the same type are automatically created by **ocpidev** and have makefiles that indicate this by including the appropriate makefile fragment from the CDK:

- The **components** directory makefile, when there are multiple libraries, contains:

```
include $(OCPI_CDK_DIR)/include/libraries.mk
```
- The **applications** directory makefile contains:

```
include $(OCPI_CDK_DIR)/include/applications.mk
```
- The **hdl/platforms** directory makefile contains:

```
include $(OCPI_CDK_DIR)/include/hdl-platforms.mk
```

- The `hdl/assemblies` directory makefile contains:

```
include $(OCPI_CDK_DIR)/include/hdl-assemblies.mk
```

- The `hdl/primitives` directory makefile contains:

```
include $(OCPI_CDK_DIR)/include/hdl-primitives.mk
```

The variables in the top level project **Makefile** apply when make is invoked in that top-level project directory. The variables set in the **Project.mk** file in the top-level project directory apply anywhere in a project, when make is invoked in any of the project's directories.

14.5 The Project.mk File for Project-wide Variable Settings

This file is required in the top level directory of a project. It contains variable settings that apply to all levels of a project. Its existence indicates that the directory is in fact a project. *In all directories under a project, this file is found by looking in parent directories until the `Project.mk` file is found.* This is similar to how the `git` tool finds the top level of a git repository by searching for a directory named `.git`.

This top level project file is included in all the makefiles in the project automatically (by the included standard CDK makefile fragment at each level). It is *not* the (lower case) file in the CDK that is included via the directive at the end of the project's top-level **Makefile** using

```
include $(OCPI_CDK_DIR)/include/project.mk
```

Variable settings in the project's **Project.mk** file are used even when the developer runs **make** (or **ocpidev build|clean**) in subdirectories of the project (i.e. not running **make** in the top level project directory). Within an OpenCPI project, the **ocpidev build** command can be run directly in any asset's directory with the exception of HDL primitives, which must be built from the **hdl/primitives** directory or the top level project directory.

Variable settings that apply only when running **ocpidev build** from the top-level project directory can be put in the **Makefile** in the top-level directory rather than in **Project.mk**. Variable settings in this **Project.mk** file can either override settings made in a given Makefile, add to lists, or conditionally set the variables if not already set. For example putting the following lines in **Project.mk**:

```
ifndef HdlPlatforms
    HdlPlatforms=zed
endif
```

indicates that when any level of the project is built, if **HdlPlatforms** has no value, use this definition.

The project variables that may be set in a **Project.mk** file are in the following table and are all optional. This file must be present, but may be empty.

Table 8: Variables set in the *Project.mk* file

Variable Name in <i>Project.mk</i>	Default	Description
PackageName	Project directory's name	The name used for this project.
PackagePrefix	local	The package-prefix for all assets in the library. The default, local , is appropriate when the assets are intended to be used only in the local organization or prior to registration.
Package	<prefix>.<name>	The package-ID of the project, overriding PackageName and PackagePrefix
ProjectDependencies	" "	A list of the package-IDs of other registered projects that this project depends on.

In order to avoid name space collisions when using multiple projects or component libraries (e.g. spec names and worker names), the project's package-ID is the default namespace for all named assets in the project.

The **ProjectDependencies** variable should be used to declare other projects that this project depends on. This will automatically use these other projects when searching for assets that are subject to search paths, such as:

- OCS files (when building workers)
- HDL primitives (when building HDL workers)
- Component libraries when building workers (for slaves of proxies, or devices of emulators)

While project dependencies should generally be the package-IDs of registered projects, they may also be relative or absolute pathnames of other projects. However using such direct pathnames is deprecated and may not be supported in the future.

Other useful *make* variables that can be specified in the *Project.mk* file include variables providing default lists of build targets and platforms:

HdlTargets
HdlPlatforms
RccPlatforms
RccHdlPlatforms

Search list variables may also be put in this file, although care should be taken to not unduly pollute search spaces for *all* assets in the project. These variables include:

ComponentLibraries
HdlLibraries

Such variables may be better placed in `Library.mk` files in component library directories.

14.6 Project Exports

When a project's assets are used by assets *outside* the project, they use the project's assets via its **exports**. **Exports** of a project are the files within the project that are explicitly made visible and usable from outside the project. I.e., without exports, nothing in a project is intended to be visible outside the project. A project's complete directory structure contains source files and artifacts of the build process. The **exports** are the files needed by users of the project, and can be thought of as the installable and deliverable subset of the files in the project after it is built.

When other projects depend on a project, as specified in those projects' **Project.mk** file **ProjectDependencies** setting, that means they use project exports from those projects they depend on. The projects they depend on must be registered.

The **exports tree** is a directory containing the project's exports, and is constructed as a tree of symbolic links, under the directory named **exports** at the top level of a project. The structure of the exports tree is not necessarily the directory structure of the project itself, but is a structure appropriate and convenient for users of the project's assets. By constructing the **exports** tree using symbolic links, the exported view of a project uses no extra space (no copies). The assets are therefore used externally exactly where they exist in source form or where they are built (although indirectly via the symbolic links in the exports tree).

Much like there is a standard directory structure for OpenCPI projects, there is an implied standard exports tree based on the contents of a project. At the top level of a project, the make target **exports** creates the exports tree. i.e.:

```
make exports
```

The exports tree has two different uses. One is to allow the project's intended deliverable results to be used in-place, without any copying or “installing”. The other is to provide an implicit recipe or bill-of-materials for creating an installable package for users of the project. In this latter case a simple single file deliverable package can be created (see below).

14.6.1 The Exports Tree

The top level directory **exports**, is created and populated automatically based on the file **Project.exports**. This **exports** directory can always be deleted and recreated. It is never manually constructed or changed. If the **Project.exports** file is empty or does not exist no exports are created. The default export tree is created based on the assets in the project. It is the set of exports when the **Project.exports** contains the single line containing the single word: **all**.

The next sections describe the default **exports** tree and the format of the **Project.exports** file that can be used to add or subtract from the default exports.

Here are the rules used to populate the default exports tree when **make exports** is invoked at the top level of a project and the **Project.mk** file contains the single line: **all**.

For an example of how the exports tree works:

Note: This structure is subject to change between Major releases of OpenCPI and should not be depended on by any end user developed code

component library deliverables are made available in the exports tree under the `lib` directory, using the name of the library. If there is a component library in the project in the directory `components/util_comps`, where its own locally built deliverables are in its `lib` subdirectory (`components/util_comps/lib`), then these deliverables are available in the exports tree using `lib/util_comps`, which is a symbolic link to `components/util_comps/lib`, i.e.:

```
exports/lib/util_comps -> ../../../../components/util_comps/lib
```

This directory structure is subject to change but creates an export view of the deliverables of the project using a sparse set of symbolic links. Users of the project, seeing only the exports tree, see `lib/util_comps` for this library.

All assets have a default place in the exports tree where other projects that depend on them can find them, if the project owner decides to export them.

14.6.2 The *Project.exports* file

The `Project.exports` file specifies which assets and files in the project should be made visible, usable, and accessible from outside the project (usually by other projects). If the file is missing or empty (or only contains blank lines or `#` comments), the project is not intended to be used by others outside the project.

This *nothing-is-exported* condition is appropriate for projects in development before anything is ready for use by others (and before a package-ID is assigned). It is also appropriate for projects that contain end products like applications for test purposes.

At the other end of the “visibility” spectrum, the `Project.exports` file can contain a single line containing the word `all`, which indicates that all assets in the project should be visible/usable from outside the project.

Between these two extremes we can selectively export assets and/or files in the project.

14.6.2.1 Exporting Assets by Type or Name

To export assets, we simply provide their type and name, or in some cases the plural form of their type to indicate all assets of that type. Here is a list of the lines you can use to export assets by listing their types and names. When `<name>` is missing, it indicates that all the assets of the indicated type should be exported.


```

all
spec <name>          # a spec in the top level specs/ directory
specs               # all specs in the top level specs/ directory
libraries           # all component libraries under "components"
library <name>       # <name> can also be a path within the project
hdl platforms
hdl platform <name>
platforms           # all possible platforms
rcc platforms
rcc platform <name>
hdl primitives
hdl primitive library <name>
hdl primitive core <name>
hdl assemblies
hdl devices

```

14.6.2.2 *Exporting Individual Files and Directories*

The `Project.exports` file may also contain lines that add and subtract files from the asset exports of a project. This capability is rarely used but covers some edge cases to augment or prune the exports tree created based on asset exports. File additions are lines that start with a plus sign (+), and subtractions are lines that start with a minus sign (-). White space (spaces or tabs) can precede the -, +, or # characters.

The format of addition (+) lines is two fields separated by white space. The first field is the relative pathname within the project for the file to be exported, and the second field is the location in the export tree where the file should be linked. If the second field ends in a slash, then the filename part of the first field is used as the file name in the exports tree. Pathnames or other names with embedded spaces are not currently supported.

The line:

```
+special_dir/special_file include/ # this exports my special file
```

would export the file in the project named `special_dir/special_file` as:

```
include/special_file
```

If the name of the exported file should be different, it can be included in the second field, e.g.

```
+special_dir/special_file include/different-file
```

If the second field is blank (doesn't exist), then the project file or directory is exported in the same place as it exists,

```
+special_dir # export this directory where it is in the project
```

would simply make `special_dir` a top level directory in the exports tree.

Any directory that is exported implicitly exports all files underneath it.

The first field can also have wildcard patterns using normal sh/bash wildcard patterns. A special string, `<target>`, indicates the software target that is currently being exported. When “`make exports`” is executed in the project directory, it is done in the context of a particular software target. The line below

```
+applications/myapp/target-<target>/myspecial_exe bin/<target>/
```

exports a secondary executable in the application.

Subtraction lines start with a minus sign (-). As with addition lines, the first field is the pathname within the project, with possible wildcards and <target> strings.

In both additions and subtractions the first field can actually be an extended regular expression if the field starts with a vertical bar (|). For example, an addition line might be:

```
+|^ (abc|xyz) special_dir/
```

which would export anything starting with **abc** or **xyz** in the directory **special_dir**.

14.7 Exporting a Project to be used Elsewhere

A project can be exported and used elsewhere in source form or in built form without source files. In both cases a tar file is created from the current project.

From the top level directory of the project to be exported, create the single tar file using one of the two commands below. To export the project in source form, use the command:

```
$ tar -czf ../my-project.tar.gz --exclude=exports --exclude=imports
```

This creates the tar file with every file in the project regardless of what is exported. To export the project in built form, which only includes exported assets and files, use the command:

```
$ make exports
$ tar -czf ../my-project.tar.gz -C exports --exclude=imports
```

The **make exports** command ensures that exports are ready, and the second makes the tar file consisting only of exported assets and files.

The above commands will create the tar file **my-project.tar.gz** in the directory above and outside of the project's directory. At the destination where the project should be installed, make a directory for the project and extract the tar file there:

```
$ mkdir my-project
$ cd my-project
$ tar -xvf ../my-project.tar.gz
```

If you are using the default registry (or setting the default in the environment) you can set the registry for the project using the **ocpidev set registry** command. Make sure that the registry contains any projects this exported project depends on.

If you want other projects to use or depend on *this* project, you can register it using the **ocpidev register project** command.

14.8 Using Other Projects that Exist Outside the Project Being Developed

The convenient and recommended way to use a project **A** from another project **B** is to add the package-ID of **A** to the **ProjectDependencies** variable in the project **B**'s **Project.mk** file. This declares the dependency without using any environment settings or pathnames. The dependency is persistent (in **Project.mk**) and stays with the project even if it copied or moved. Since the other project is specified in **ProjectDependencies** by its package-ID, it can also be moved (and re-registered) and the dependency is still valid.

The actual directories searched within a project depends on the type of asset being sought. E.g., if a component library is being sought, the search will look for that component library within all projects mentioned in **ProjectDependencies**. E.g.:

```
ProjectDependencies=local.utils ocpi.assets
```

would indicate that this project depends first on the **local.utils** project and then also depends on assets in the **ocpi.assets** project (as well as **ocpi.core** which is added to the dependency list). Any component library being sought might be in either one, and if found in **local.utils**, the **ocpi.assets** project will not be used.

When a more dynamic setting is needed, e.g. for temporarily using one version of a project vs. another, the **OCPI_PROJECT_PATH** environment variable can be used. This variable specifies a colon-separated set of other projects (by pathname) to be searched in order. These are searched *before* projects mentioned in the **ProjectDependencies** variable to allow the environment variable to temporarily override what is declared persistently in the project. Projects in **OCPI_PROJECT_PATH**, are examined by looking in their **exports** subdirectory if it exists.

To summarize searching at the project level, the order is:

1. Use assets in projects in **OCPI_PROJECT_PATH** to temporarily override anything later in this list.
2. Use assets in the local project.
3. Use assets in **ProjectDependencies**, in the order given in that variable (with **ocpi.core** automatically added at the end, as previously mentioned)

In cases where a standalone component library or application is being developed outside the project structure, the **OCPI_PROJECT_PATH** variable can also be used when these non-project assets depend on projects other than **ocpi.core**.

15 The `ocpidev` Tool for Managing Assets

The `ocpidev` command line tool is used to perform various development-related tasks inside projects as well as retrieving information about the environment. When used in projects, it may be invoked at the top level of a project, or in lower level directories of the project as appropriate to the particular command being used. The `ocpidev` command has full tab completion for its options and arguments.

The general usage of the `ocpidev` command is:

```
ocpidev [<options>] <verb> [<noun> [<name> [<arguments>]]]
```

The options can in fact occur anywhere in the command for the user's convenience.

The general usage concept is:

*perform the **<verb>** operation on the **<noun>** asset type whose name is **<name>**.*

The verbs are:

create — create the named asset, creating files and directories as required, and creating any skeleton files for future editing.

delete — remove all directories and files associated with this named asset

build — build the asset(s), running appropriate tools to create the binary files

clean — remove all the generated and compiled files for the asset(s)

show — display information about assets in registered projects and the current build environment (preliminary)

[un]register — register/unregister a project in its registry

[un]set — set/unset the registry used by the current project

run — execute the unit-test or application

refresh — manually regenerate the metadata associated with the project

Most options are only valid for specific verbs or nouns, but some are general purpose and described in the table just below. As each option is described, it is also indicates whether it takes a value from the following argument (under the **V?** column), and whether the option may appear more than once (under the **M?** column). Options are either single letters following a hyphen or complete words or acronyms following two hyphens and separated by hyphens.

*Table 9: General Purpose Options for the **ocpidev** command*

Option	V?	M?	Description
--help	?	N	Display help information; may be used in conjunction with a verb to display verb-specific information, e.g.: ocpidev --help show
-v --verbose	N	N	Be verbose, describing what is happening in more detail
-d <dir>	Y	N	Specify the directory in which this command should be run. Analogous to the -c option in the POSIX make command.

15.1 Create/Delete Assets

The **create** and **delete** verbs create/delete OpenCPI assets in/from a project specified by the noun arguments that follow the verb.

```
ocpidev [options] create|delete <noun> <name>
```

15.1.1 Asset Types for Create and Delete

project — create/delete a project which holds and organizes OpenCPI assets

registry — create/delete a registry that enable projects to depend on each other

application — create/delete an application; XML or ACI as specified by options

spec — create/delete a component specification in a library or project

protocol — create/delete a protocol in a library or project

test — create/delete a component unit test in a library

library — create/delete a component library

worker — create/delete a worker in a library based on a component specification

hdl — a prefix to denote an HDL asset, in the **hdl** subdirectory of a project

assembly — create/delete HDL assembly

card — create/delete a HDL card definition

slot — create/delete a HDL slot definition

device — create/delete an HDL device worker, possibly in a specified platform

platform — create/delete a HDL platform

primitive — create/delete an HDL primitive; type needs to be follow

library — a set of HDL primitive modules to use in workers as a library

core — a primitive a single module, perhaps presynthesized (i.e. from coregen or megawizard)

All HDL assets should avoid having names that are Verilog or VHDL reserved identifiers, such as “reg”, etc.

15.1.2 Options for Create and Delete

The options for these verbs are usually specific to the asset type that is being created or deleted. When deleting an asset, **ocpidev** will ask for confirmation. This can be overridden by using the **-f** option to force the deletion without confirmation. When creating an asset, if creation fails all partially created directories and/or files are removed unless the **-k** (keep) option is specified.

The following table shows all creation and deletion options and which asset types the options apply to.

Table 10: Creation/Deletion Options for the *ocpidev* command

Option	V?	M?	Description
<i>Option when creating any asset type</i>			
-k	N	N	Keep files and directories created after a creation fails. Normally all such files and directories are removed upon any failure.
<i>Option when deleting any asset type</i>			
-f	N	N	Force deletion, do not ask for confirmation when deleting an asset. Normally, the user is asked to confirm a deletion.
<i>Options when creating projects</i>			
-D <pkg-ID>	Y	Y	Indicate another project that the project being created will depend on. The value is the package-ID of that other project.
--register	N	N	Register the project as it is created.
<i>Options when creating projects or libraries</i>			
-N <pkg-name>	Y	N	Specify the package name of the created project or library. The default is the name argument after the create verb.
-F <pkg-prefix>	Y	N	Specify the package prefix of the created project or library. The default is local for projects, and <i><package-ID of parent></i> for libraries.
-K <pkg-ID>	Y	N	Specify the package-ID directly. The default is: <package-prefix>.<package-name> .
<i>Options when creating projects, libraries, workers, hdl devices, or hdl platforms</i>			
-A <dir>	Y	Y	Specify a directory to search for XML include files.
-y <comp-lib>	Y	Y	Specify a component library to search for workers/devices/specs that this asset (or assets it contains) references.
-I <dir>	Y	Y	Specify a directory to search for include-files (C, C++, Verilog).
-Y <prim-lib>	Y	Y	Specify a primitive library this asset (or assets it contains) depends on.
<i>Options when creating specs</i>			
-t	N	N	Indicate that when creating a spec ,also create the unit test for it.
-n	N	N	Indicates that the created spec should have no control interface. This is rarely used or recommended, but is required for certain infrastructure specs. Only applies to specs with only HDL implementations.
<i>Option when creating or deleting specs or protocols</i>			
-p	N	N	Indicates the protocol or spec will be (create) and is (delete) in the specs directory of the project, and not in any library's specs directory.

Option	V?	M?	Description
<i>Options when creating or deleting specs, protocols, workers or hdl devices</i>			
-l <library>	Y	N	Indicate that the asset should be created in or deleted from the specified library.
--hdl-library <hdl-library>	N	N	Indicates that the asset should be created in or deleted from the library under the hdl directory of the project. Valid values are only one of: devices , cards , adapters .
<i>Options when creating workers of all types, as well as when creating tests</i>			
-s <spec>	Y	N	Indicate the spec to be implemented by the created worker. The default is <name>-spec or <name>_spec depending on what is found in the specs directory of the library or project (or libraries indicated by the -y option or other projects indicated by the -d option).
<i>Options when creating workers of all types</i>			
-l <language>	Y	N	Specifies the source language for the worker being created. Defaults to the default language for the authoring model. It must be C or C++ for RCC workers and VHDL for HDL workers.
-p <platform>	Y	N	Indicates that the worker being created should be created in the devices library for the specified HDL platform (in this project).
-o <file>	Y	Y	Specify a source code file to compile when building this worker that is not included by default (i.e. in addition to the <worker>.[c cc vhd] file).
<i>Options when creating RCC workers</i>			
-v <slave-wkr>	Y	N	Indicates that the RCC worker being created is a proxy for the slave worker identified by the value of this option.
-w <worker>	Y	Y	Specify one of multiple workers implemented in this RCC worker's directory when a single RCC worker directory is creating a multi-worker artifact. This is supported but rarely required or recommended.
-R <prereq-lib>	Y	Y	Specifies that the worker being created depends on and should be statically linked with a prerequisite library. See RccStaticPrereqLibs .
-r <prereq-lib>	Y	Y	Specifies that the worker being created depends on and should be dynamically linked with a prerequisite library. See RccDynamicPrereqLibs .
<i>Option when creating HDL workers, devices and platforms</i>			
-C <core>	Y	Y	Specifies a HDL primitive core that this worker depends on and should be built with.
<i>Options when creating HDL devices</i>			
-E <hdl-device>	Y	N	Indicates that the worker being created should be an emulator worker, and specifies the HDL device worker associated with the device being emulated.

Option	V?	M?	Description
-U <hdl-device>	Y	Y	Indicates that the worker being created is a subdevice, and specifies one of the HDL device workers that this subdevice supports.
<i>Options when creating HDL platforms</i>			
-g <part>	Y	N	Specify the part (die-speed-package, e.g. xc7z020-1-c1g484) for HDL platform.
-q <frequency>	Y	N	Specify the time server frequency for this platform.
-u	N	N	Indicates the platform will not support the SDP (the scalable data plane). This is for legacy platforms and not recommended.
<i>Options when creating any buildable asset</i>			
-T <target>	Y	Y	Only build the asset for the specified target (see OnlyTargets)
-Z <target>	Y	Y	Do not build the asset for the specified target (see ExcludeTargets)
-G <platform>	Y	Y	Only build the asset for the specified platform (see OnlyPlatforms)
-Q <platform>	Y	Y	Do not build the asset for the specified platform (see ExcludePlatforms)
<i>Options when creating HDL primitives libraries</i>			
-H	N	N	Indicates that this primitive does not depend on any other primitive libraries. By default hdl primitive libraries are assumed to depend on the built-in primitive libraries (currently bsv , fixed_float , ocpi , and util).
-J	N	N	Indicates that this primitive library should not be elaborated when building, which results in a slight time savings at the cost of catching elaboration errors earlier in the build process.
<i>Options when creating HDL primitives cores</i>			
-M <module>	Y	N	Specifies the name of the top module of the core. The default is the name of the core specified in the <name> argument after create .
-B <file>	Y	N	Indicates that the core is <i>prebuilt</i> (e.g. from coregen) and specifies the name of the prebuilt core file (e.g. an .ngc or .qxp file).
<i>Options when creating applications</i>			
-X	N	N	The application being created will simply be an XML file in the applications directory of the project.
-x	N	N	The application has its own directory but it will only be an XML file. The default creates an application with an ACI C++ file. Cannot use with the -X option.

15.1.3 Examples of Create/Delete

Create a project with name **my-project**. This project's package-ID will be **local.my-project**.

```
ocpidev create project my-project
```

Create a project with package-prefix **org.my_organization** and name **my_project**. This project's package-ID will be **org.my_organization.my_project**.

```
ocpidev create project my-project -F org.my_organization
```

Create a project-registry **~/workspace/my-registry**. Note that a new registry will not have a core project registered, so the next thing to do is locate and register a core project. Finally, **ocpidev set registry ~/workspace/my-registry** can be used from within a project to tell that project to use the new registry. None of this is necessary if using the default project-registry in the OpenCPI installation.

```
ocpidev -d ~/workspace create registry my-registry
```

Create the default **components** library for the project. If a name other than **components** is provided, it will be created underneath the **components** directory.

```
ocpidev create library components
```

Create a component spec with name **mycomp**. It will be placed in the default location, which is **components/specs**. Provide **-l <library-name>** if there are sub-libraries underneath **components**.

```
ocpidev create spec mycomp
```

Create an HDL worker named **myworker** that implements the **mycomp** spec. If the worker was instead named **mycomp.hdl**, the **-S mycomp-spec** argument can be omitted because the default spec is **<worker-name>-spec.xml**.

```
ocpidev create worker myworker.hdl -S mycomp-spec
```

Create a C++ RCC worker named **mycomp** that implements the **mycomp** spec.

```
ocpidev create worker myworker.rcc -L c++
```

Create an HDL assembly named **myassy** that will only be used on the **isim** platform.

```
ocpidev create hdl assembly myassy -G isim
```

15.2 Build/Clean Assets

These verbs compile/clean the OpenCPI assets specified by the noun used in the command. If no nouns are provided, `ocpidev` will compile/clean the current directory. When the plural version of the nouns are used, no name is specified and all assets of that type are built/cleaned. Usage is as follows:

```
ocpidev [options] build|clean [<noun>] [<name>]
```

15.2.1 Asset Types for Build and Clean

The asset types that are valid with the **build** and **clean** verbs are:

project – build/clean all assets in a project

application(s) — build/clean all (**applications**) or a specified ACI **application**. This includes any custom building/cleaning specified by a customized application **Makefile**.

test — build/clean (but do not run) the unit test with the specified *<name>* in the specified library

library — build/clean all workers in a library

worker – build/clean a specific worker in a specified library (using **-l** *<library>*)

hdl – specifies the HDL subdirectory of a project

assembl(y|ies) — build/clean all or a specified assembly

device – build/clean a specified device in a specified library similar to worker

platform(s) — build/clean all or a specified platform (also builds primitives and devices for specified platform)

primitive(s) — build/clean all or a specified primitive. If specifying a primitive you also need to specify if it is a library or a core

library — build/clean a primitive that results in a library of modules

core — build/clean a primitive that results in a synthesized core/netlist

15.2.2 Options for Build and Clean

The valid options with the **build** and **clean** verbs are in the following table. The options are primarily used to limit what is built (focus the build process on specific assets of interest), or specify the platforms and targets to build the assets for.

Table 11: Build and Clean Options for the **ocpidev** command

Option	V?	M?	Description
<i>Options when operating on projects</i>			
--no-assemblies	N	N	Do not build or clean HDL assemblies in the project. Building HDL assemblies (and their containers) can be slow.
--hdl-assembly <assembly>	Y	Y	Limit the assemblies being built/cleaned to assemblies specified by this option. May also be used in the hdl/assemblies directory.
<i>Options when operating on projects or libraries</i>			
--rcc	N	N	Limit the assets being built/cleaned to RCC workers.
--hdl	N	N	Limit the assets being built/cleaned to HDL primitives, workers, platforms or assemblies.
--worker <wkr>	Y	Y	Only operate on workers whose names are specified by this option.
<i>Options when building projects, libraries, tests, RCC workers or ACI applications</i>			
--rcc-platform <platform>	Y	Y	Build RCC/ACI assets for the RCC platforms specified by this option. If this option is not used at all (and --hdl-rcc-platform is also not used), the current development software platform is used as the single RCC platform used for building.
--hdl-rcc-platform <hdl-platform>	Y	Y	Build RCC/ACI assets for the RCC platforms associated with the HDL platforms specified by this option. If this option is not used at all (and --rcc-platform is also not used), the current development software platform is used as the single RCC platform used for building.
<i>Options when building projects, libraries, tests (with HDL workers) or other HDL assets</i>			
--hdl-target <target>	Y	Y	Build HDL assets for the HDL targets specified by this option. If only HDL targets are specified (and no HDL platforms), containers will not be build for assemblies or unit tests.
--hdl-platform <hdl-platform>	Y	Y	Build HDL assets for the HDL platforms specified by this option.

15.2.3 Examples of Build and Clean

15.2.3.1 Build Examples

Build the current asset/directory and those underneath it for the **zed** HDL platform and the **xilinx13_3** RCC platform.

```
ocpidev build --hdl-platform zed --rcc-platform xilinx13_3
```

Build the assets project for the **zed** HDL platform and **xilinx13_3** RCC platform. Omit the name **assets** if inside the project.

```
ocpidev build project assets --hdl-platform zed
--rcc-platform xilinx13_3
```

Inside the **assets** project, build the **complex_mixer.hdl** worker in the **dsp_comps** library for the **zed** HDL platform and **xilinx13_3** RCC platform. The **-l dsp_comps** can be omitted if operating in a project with only a single top-level **components** library.

```
ocpidev -l dsp_comps build worker complex_mixer.hdl
--hdl-platform zed
--rcc-platform xilinx13_3
```

Inside the **assets** project, build the **dsp_comps** library located at **components/dsp_comps** for the **zed** HDL platform and **xilinx13_3** RCC platform. (Note: that the **library noun** differs from the **-l option** in that the noun is used when building the library itself, and the option is used when building an asset inside the library.)

```
ocpidev build library dsp_comps --hdl-platform zed
--rcc-platform xilinx13_3
```

Inside the **assets** project, build the **dsp_comps** library, but limit the build to the **complex_mixer.hdl** and **cic_dec.hdl** workers. Note that the **--worker option** differs from the **worker noun** in that the noun is used to build a single worker, while the **--worker** option is used when building a *library* to limit the build to an individual worker (or a list of workers if **--worker** is specified multiple times). If building a library without the **--worker** option, all workers will be built.

```
ocpidev build library dsp_comps --worker complex_mixer.hdl
--worker cic_dec.hdl
--hdl-platform zed
```

15.2.3.2 Clean Examples

Clean the current asset/directory and those underneath it.

```
ocpidev clean
```

Cleans the **assets** project. Omit the name **assets** if inside the project.

```
ocpidev clean project assets
```

Inside the **assets** project, this will clean the **fir_real_sse.hdl** worker in the **dsp_comps** library.

```
ocpidev -l dsp_comps clean worker fir_real_sse.hdl
```

Inside the **assets** project, clean the **dsp_comps** library located at **components/dsp_comps**.

```
ocpidev clean library dsp_comps
```

15.3 Show Assets

The **show** verb is slightly different than other verbs since the command currently is not associated with a project, but with the operating environment by default. There are two scoping options to associate with the command by default the **--global-scope** option is set and it will look in *all* registered projects (and the current project if applicable). The registry used is the one in the OpenCPI installation unless overridden using the **OCPI_PROJECT_REGISTRY_DIR** environment variable. The second mode to run the command in to use with just the current project use the **--local-scope** option.

Usage is as follows:

```
ocpidev show [options] <noun>
```

15.3.1 Asset Types for Show

The nouns that are valid with this verb are:

- registry** — display information on the registry
- project** — display information on the registry
- tests** — display information on the unit tests in the current scope
- libraries** — display information on all the libraries in the current scope
- projects** — display all registered projects
- worker** — display information on the specified worker
- workers** — display all workers in registered projects
- components** — display all components in registered projects
- component** — display information on the specified component
- platforms** — display all platforms that assets that can be built for
- targets** — display all targets that assets that can be built for

The keywords **rcc** and **hdl** can be used to scope what is displayed for the **platforms** and **targets** nouns. E.g.:

```
ocpidev show rcc platforms
```

only shows the available rcc platforms.

15.3.2 Options for Show

The options that are valid with this verb are:

- table** — default display format used if **--simple** or **--json** are not set, the output in an easy-to-read table
- simple** — format the output as simple as possible
- json** — format the output in json format, for integration with other software
(see [Appendix A](#) for more information)
- global-scope** — default scope used if **--local-scope** is not set, show assets from all registered projects and the current project if applicable
- local-scope** — Only show assets in the local project
- very-verbose** —Be very verbose with output
- very-very-verbose** —Be very very verbose with output
- l LIBRARY, --library LIBRARY** —
Specify the component library in which this operation will be performed.
- hdl-library HDL_LIBRARY** —
Specify the hdl library in which this operation will be performed.
- P HDL_PLAT_DIR** —Specify the hdl platform subdirectory to operate in.

15.3.3 Examples of Show

List all RCC and HDL platforms that assets can be built for.

```
ocpidev show platforms
```

List all HDL platforms that HDL assets can be built for.

```
ocpidev show hdl platforms
```

List all components in any registered projects (omits components at the top level of project due to a defect)

```
ocpidev show components
```

Print a table of the registered projects and their locations.

```
ocpidev show --table registry
```

15.4 Run

Utility for running component unit-tests and applications. Applications that are simply a application xml are not yet implemented using ocpidev.

Component unit-tests have 5 phases:

- Generate: generate testing artifacts after finding the spec and the workers
- Build: building HDL bitstream/executable artifacts for testing
- Prepare: examine available built workers and available platforms, creating execution scripts to use them all for executing feasible tests.
- Run: execute tests for all workers, configurations, test cases and platforms

Verify: verify results from the execution of tests on workers and platforms

Based on these 5 phases there are modes for a run command that can be set with the `--modes` option as follows:

all- execute all 5 phases

gen- execute generate phase

gen_build- execute generate and build phases

prep_run_verify- execute prepare, run, and verify phases

prep- execute prepare phase

run- execute run phase

prep_run- execute prepare and run phases

verify- execute verify phase

view- execute the view script (view.sh) on a already executed run

clean_all- clean all generated files

clean_run- clean all files generated during the run phase

clean_sim- clean all simulation files

There are 2 types of variables for setting platforms. The first type (`--rcc-platform` and `--hdl-platform`) that causes the command to build (generate and build phases) for the specified platforms. The second type (`--only-platform` `--exclude-platform`) are used to limit the run time available platform (prepare, run, and verify phases) where the command is executed.

Usage is as follows:

```
ocpidev [options] run <noun>
```

15.4.1 Associated nouns

test — run the unit test that is specified

tests — run all the tests in the given project or library

library — run all the unitests in the given library

application — run the application that is specified

applications — run all the applications in a project

project — run all unit tests and applications in the specified project

15.4.2 Associated options

Valid for both Unit tests and applications:

--rcc-platform —

Specify which RCC platform from the list of buildable platforms to use with unit test. Only valid in generate and build phases. For application specifies which RCC platform to build ACI applications.

Valid for only Unit tests:

--keep-simulations —

Keep HDL simulation files regardless of verification results. By default, simulation files are removed if the verification is successful. Warning: Simulation files can become large!

--accumulate-errors —

When enabled, execution or verification errors are accumulated (i.e. not stop the process) and simply be reported as they occur, rather than ending the test upon first failure detected.

--only-platform —

Specify which platforms to use with a unit test from the list of runtime platforms. (-G is another way to specify this functionality)

--exclude-platform —

Specify which platforms not to use with a unit test from the list of runtime platforms. (-Q is another way to specify this functionality)

--mode —

Specify which phase(s) of the unit test to execute. Valid modes are : all, gen, gen_build, prep_run_verify, prep, run, prep_run, verify, view, clean_all, clean_run, clean_sim

--remotes —

Specify remote systems to run the test(s) by setting the OCPI_REMOTE_TEST_SYS variable. See section [OCPI_REMOTE_TEST_SYS](#) for more information.

--view —

When set the view script (view.sh) for this test is run at the conclusion of the test's execution.

--hdl-platform —

Specify which HDL platform from the list of buildable platforms to use with unit test. only valid in generate and build phases.

--case —

Specify which test case(s) that will be run/verified. Wildcards are valid, ex. case*., case0.0*, case00.01

Valid for only Applications:

--before —

Argument(s) to insert before the ACI executable or ocpirun, such as environment settings or prefix commands.

--after —

Argument(s) to insert at the end of the execution command line

--run-arg —

Argument(s) to insert immediately after the ACI executable or ocpirun.

15.4.3 Examples

run a single application:

ocpidev run application <app_name>

run all applications in a project:

ocpidev run applications

run a single test:

ocpidev run -l <library_dir> test <test_name>

run all tests in a library:

ocpidev run library <library_name>

run generate and build stages of a single test:

ocpidev run -l <library_dir> --mode gen_build test <test_name>

15.5 Refresh

Projects have some metadata associated with them. This metadata is used by ocpidev show and can be used by other programs can use it to interface with projects. If ocpidev is used throughout usage of the project metadata is automatically updated this verb is only required if something outside of ocpidev operates on a project. Usage is as follows:

ocpidev [options] refresh project

15.5.1 Associated nouns

project — the only valid noun, indicating the the current projects metadata to be refreshed.

15.5.2 Associated options

No verb specific options

15.5.3 Examples

Manually regenerate the project's metadata.

ocpidev refresh project

15.6 Utilization

Utility for displaying/recording FPGA resource utilization for HDL OpenCPI assets. In order to display any utilization information about a hdl asset for a given platform the asset needs to be built for that platform. Usage is as follows:

```
ocpidev [options] utilization [<noun> [<name>]]
```

15.6.1 Associated nouns

project — Indicating the the current project to display all utilization information for all built containing workers, platforms, and assemblies.

worker — Display all utilization information for the indicated hdl worker

workers — Indicating the the current project to display all utilization information for all built containing workers.

library — Indicating a component library to display all utilization information for all built containing workers.

hdl — a prefix to denote an HDL asset, in the **hdl** subdirectory of a project

platform — Display all utilization information for the indicated hdl platform

platforms — Indicating the the current project to display all utilization information for all built containing platforms.

assembly — Display all utilization information for the indicated hdl assembly

assemblies — Indicating the the current project to display all utilization information for all built containing assemblies.

15.6.2 Associated options

--format {table,latex}

Format to output utilization information. "latex" results in silent stdout, and all output goes to "utilization.inc" files in the directories for the assets acted on.

--hdl-platform HDL_PLAT

Specify which HDL platform from the list of buildable platforms to show utilization for.

--hdl-target HDL_TGT

Specify which HDL target from the list of buildable targets to show utilization for. Only valid for workers (not assemblies)

-I LIBRARY, --library LIBRARY

Specify the component library in which this operation will be performed.

--hdl-library HDL_LIBRARY

Specify the hdl library in which this operation will be performed.

-P HDL_PLAT_DIR

Specify the hdl platform subdirectory to operate in.

15.6.3 Examples

show utilization for a single worker (using build results from all platforms):

ocpidev utilization worker <worker-name>

show utilization for a single worker (using build results from a single platform):

ocpidev utilization worker <worker-name> --hdl-platform <hdl-platform>

show utilization for a single worker (using build results from a single target):

ocpidev utilization worker <worker-name> --hdl-target <hdl-target>

show utilization for all workers (in the current project/library/etc):

ocpidev utilization workers

show utilization for a single worker in a named library:

ocpidev utilization worker <worker-name> -l <library>

show utilization for a single HDL Platform:

ocpidev utilization hdl platform <platform-name>

show utilization for all HDL Platforms (in the current project):

ocpidev utilization hdl platforms

show utilization for a single HDL Assembly:

ocpidev utilization hdl assembly <assembly-name>

show utilization for all HDL Assemblies (in the current project):

ocpidev utilization hdl assemblies

record utilization for a single HDL Assembly in LaTeX format:

ocpidev utilization hdl assembly <assembly-name> --format=latex

show utilization for all supported assets in a project:

ocpidev utilization project

record utilization for all supported assets in a project in LaTeX format:

ocpidev utilization project --format=latex

15.7 Register/Unregister

These verbs are used to add or remove the current project to/from its registry. If the OpenCPI installation is RPM-based, and the current registry is the one in the OpenCPI installation, the user must be in the `opencpi` group to use these commands.

When registering a project, if the project has no associated registry, the default one is used and becomes persistently associated with the project. The default registry is the one in the OpenCPI installation unless overridden by the `OCPI_PROJECT_REGISTRY_DIR` environment variable. Registering a project makes its exports visible and usable to other projects using the same registry. Unregistering a project removes that visibility.

The usage of the command is the following:

```
ocpidev [options] [un]register project
```

When registering, this command also creates exports for the currently available and built assets in the project. If the project is already registered, the exports are still recreated.

15.7.1 Asset Types for Register/Unregister

project – only valid noun which is the project to add or remove from its registry

15.7.2 Options for Register/Unregister

The general purpose options (`-v` and `-d <dir>`) options may be used, and in particular, the `-d <dir>` option may be used to specify the pathname of the project being registered or unregistered.

15.7.3 Examples

Register the current project to its registry (or the default if it doesn't have one yet):

```
ocpidev register project
```

Unregister the current project.

```
ocpidev unregister project
```

15.8 Set/Unset

Projects are automatically associated with the current registry when they are created or used. A project can be *manually* associated with or disassociated from a particular registry. This command must be run from within a project. The `<dir>` argument only applies to the `set` verb (not `unset`), and specifies the registry's directory path. Usage is as follows:

```
ocpidev [options] [un]set registry [<dir>]
```

15.8.1 Associated nouns

registry — the only valid noun, indicating that the current project's registry is being set/unset.

15.8.2 Associated options

No verb specific options

15.8.3 Examples

Set the current project's registry to the default.

```
ocpidev set registry
```

Set the current project's registry to the directory called **my-registry** located at **~/workspace/**.

```
ocpidev -d ~/workspace set registry my-registry
```

Unset the current project's registry. The next time this project is built, the default registry will be used.

```
ocpidev unset registry
```

15.9 Assets Managed by *ocpidev*

Below are two tables summarizing all the types of assets that the *ocpidev* command operates on, and which nouns apply to it. The first table is for assets that are not HDL specific. The second is for HDL specific assets located under the *hdl* sub directory.

Table 12: Non-HDL Asset Types (Nouns) for the *ocpidev* Command

Name/ Noun	In Library	Description
application	no	A component application, specified either in XML or C++
applications	no	All applications in a project, in its applications directory
library	no	A component library.
project	no	A project containing all other asset types.
properties	yes	A properties XML file, which can be at the project or library level.
protocol	yes	A protocol specification XML file, which can be at the project or library level.
registry	no	A registry for projects to import/export dependencies from/to.
component	yes	A component specification XML file (OCS), which can be at the project or library level.
test	yes	A unit test suite for a component specification.
worker	yes	A worker, that implements a component spec.

The following table describe HDL asset types. When these assets are described as being in an HDL library, it means in one of the fixed HDL libraries in a project (**hdl/adapters**, **hdl/cards**, and **hdl/devices**) or the **devices** library within an HDL platform's directory.

*Table 13: HDL Asset Types (Nouns) for the **ocpidev** Command*

Name/ Noun	In Library	Description
hdl assembly	no	An assembly of HDL workers, used to build HDL containers
hdl assemblies	no	All the assemblies in the project or in an HDL assemblies directory.
hdl card	yes	A card specification XML file, which can be at the project or HDL library level.
hdl device	yes	A HDL device worker, in an HDL component library (adapters , cards , devices or a HDL platform's devices).
hdl platform	no	An HDL platform worker, in an HDL platforms directory, including its platform configurations.
hdl platforms	no	All HDL platforms in a project or in its hdl/platforms directory.
hdl primitive core	no	A HDL primitive core in a project or in its hdl/primitives directory
hdl primitive library	no	A HDL primitive library in a project or in its hdl/primitives directory
hdl primitives	no	All HDL primitives (cores or libraries) in a project.
hdl signals	no	A signals specification XML file, which can be at the project or HDL library level.
hdl slot	no	A slot type specification XML file, which can be at the project or HDL library level.
hdl subdevice	yes	An HDL subdevice worker, that supports other HDL device workers, in an HDL devices library.

16 Environment Variables used in Component Development

While any **make** variable can be set in the environment, only those specifically mentioned as usable in the environment should be set, since setting arbitrary **make** variables in the environment can lead to unpredictable/undefined behavior.

All OpenCPI environment variables start with the prefix **OCPI_**. While some defined variables apply only to certain authoring models or targets, a master list is kept here. Some may be described as documented elsewhere.

There also may be environment variables starting with **OCPI_** that are used internal to OpenCPI and not documented for users.

In the table below, environment variables that are boolean options have the value 0 or 1. When they are unset, or set to the empty string, it is equivalent to the value 0.

In the table below a **list** means a white-space-separated list of items. A **path** means a colon-separated list of directories. All these variables have commonly used defaults so that most users have very few settings.

Most of the items are set at installation time, several are simply convenience variables derived from others. The ones that a developer might actually set *during* development are shaded.

Name	When set?	Data type	Description
OCPI_ALTERA_DIR	install	directory	Top level directory for all Altera tools. Default: <code>~/altera</code> or <code>/opt/Altera</code>
OCPI_ALTERA_VERSION	install	directory	Directory under <code>OCPI_ALTERA_DIR</code> for current/desired Altera tools version. Default is highest numeric version present under <code>\$OCPI_ALTERA_DIR</code>
OCPI_ALTERA_TOOLS_DIR	install	directory	Directory for using Altera tools. Default: <code>\$OCPI_ALTERA_DIR/\$OCPI_ALTERA_VERSION</code>
OCPI_ALTERA_LICENSE_FILE	install	file	Altera license file. Default: <code>\$OCPI_ALTERA_DIR/Altera-License.lic</code>
OCPI_ALTERA_KITS_DIR	install	directory	Directory for Altera development kits. Default is <code>\$OCPI_ALTERA_TOOLS_DIR/kits</code>
OCPI_ASSERT	build	boolean	Enable assertions in any language (when != 1, for C and C++, sets NDEBUG). Default: 1
OCPI_CDK_DIR	install	directory	The root of the installed CDK. Default: <code>/opt/opencpi/cdk</code>
OCPI_CFLAGS	build	list	Flags when compiling C code. Default set per target, but may be overridden.
OCPI_CXXFLAGS	build	list	Flags when compiling C++ code. Default set per target, but may be overridden.
OCPI_DEBUG	build	boolean	Controls debug logging, etc. For C and C++, enables “-g” also, or if not set, -O2. Is defined as a preprocessor macro for core software or executables. Default: 1
OCPI_DYNAMIC	build	boolean	For main programs and software libraries: use dynamic linking and dynamic libraries. For software components, build for use in dynamic executables. Currently only 0 is supported.
OCPI_LIBRARY_PATH	run	path	Runtime search path for binary artifacts. Default: <code>\$OCPI_CDK_DIR/lib/components/rcc</code>
OCPI_PREREQUISITES_DIR	install	directory	Location for using prerequisites, default locations for building and installing prerequisites. Default: <code>/opt/opencpi/prerequisites</code>
OCPI_PROJECT_REGISTRY_DIR	build	directory	Location of current project registry, overriding the default in the OpenCPI installation.
OCPI_SMB_SIZE	run	number	To override size in bytes of data plane endpoints. Default: 100000000.

Name	When set?	Data type	Description
OCPI_TOOL_PLATFORM	build	string	The platform on which the current environment is executing. Automatically determined by the <code>platforms/getPlatforms.sh</code> script.
OCPI_TARGET_PLATFORM	build	String	The default target software platform to build for.
OCPI_XILINX_DIR	install	directory	Top level directory for all Xilinx tools. Default is from Xilinx: <code>/opt/Xilinx</code>
OCPI_XILINX_VERSION	install	directory	Directory under <code>OCPI_XILINX_DIR</code> for current/desired Xilinx tools version. Default is highest numeric version present under <code>\$OCPI_XILINX_DIR</code> .
OCPI_XILINX_TOOLS_DIR	install	directory	Directory for using Xilinx tools. Default: <code>\$OCPI_XILINX_DIR/\$OCPI_XILINX_VERSION</code>
OCPI_XILINX_LICENSE_FILE	install	file	Xilinx tool license file. Default is: <code>\$OCPI_XILINX_DIR/License.lic</code>
OCPI_XILINX_LABTOOLS_DIR	Install	directory	Where a Xilinx Lab Tools installation is, default is <code>\$OCPI_XILINX_TOOLS_DIR</code> .

17 Tools Used in Component Development

The most commonly used tools used during OpenCPI component development are **make**, **ocpidev**, and **ocpirun**. The latter is used for executing component-based applications, and is fully described in the ***Application Development Guide***. The **make** tool invokes a variety of other tools to build workers, all behind the scenes. When there are errors or anomalies during the build process, the logs from the various tools used can be examined to understand what went wrong. The **ocpidev** tool is used to create and delete various OpenCPI assets, usually within a project. The **make** command should only be used for the small number of development tasks not yet supported by **ocpidev**, such as executing unit test suites.

18 Appendices

18.1 Appendix A - *ocpidev* JSON

All nouns (asset types) used with the **ocpidev show** command have the option to output information in a JSON format and the format of the JSON output will be different for each noun. This appendix details the format of JSON output for each noun.

If the data type of an asset in the JSON file is a JSON String this will be the directory of the asset, otherwise it will be a dictionary with an entry of **directory** that is the directory of the asset. (It is possible that Strings may be changed to dictionaries as needed as the interface matures.)

Libraries can override the default package-id and this can cause multiple libraries within the same project to have the same package-id. In order to express this in a JSON file, numbers are added to the string of the package-id if there are more than one library with the same package-id. E.g. the core project has two libraries (**hdl/adapters** and **components**) with the package-id of **ocpi.core**. The libraries are named **ocpi.core** and **ocpi.core:1** in the JSON file for the project. It is up to the JSON data consumer to parse these package-ids and use them in the correct way.

Levels of verbosity are used to ask **ocpidev** for more information, the default being less information. The rule of thumb for what is output when showing the minimal information for an asset is to not include anything about the assets that are contained by it (i.e. a library contains components etc.). Each verbosity level that is added will show the information for one more level of contained assets. E.g. very verbose (-vv) for a project will show information on the project, libraries in the project, and workers/components/test in each library but verbose (-v) will not show workers/components/test in each library.

18.1.1 *ocpidev show project --json*

```
{
  "project": {
    "package": <String - Project package-id>,
    "dependencies": <Array of strings - package-ids of projects to
                    depend on>,
    "directory": <String - Project directory>
  }
}
```

18.1.2 *ocpidev show project -v --json*

```
{
  "project": {
    "directory": <String - Project directory>,
    "package": <String - Project package-id>,
    "dependencies": <Array of strings - package-ids of projects to
                      depend on>,
    "libraries": {
      <String- package id of library>: <String- directory of
library>,
      ... more libraries
    },
    "components": {
      <String- name of component>: <String - directory of library>,
      ... more components
    },
    "platforms": {
      "rcc": {
        <String- name of rcc platform>: <String - directory of rcc
platform>,
        ... more rcc platforms
      },
      "hdl": {
        <String- name of rcc platform>: <String - directory of hdl
platforms>,
        ... more hdl platforms
      }
    }
  }
}
```


18.1.3 *ocpidev show project -vv --json*

```
{
  "project": {
    "directory": <String - Project directory>,
    "package": <String - Project package-id>,
    "dependencies": <Array of strings - package-ids of projects to
depend on>,
    "libraries": {
      <String- package-id of Library>: {
        "package": <String- package-id of Library>,
        "directory": <String - Library directory>,
        "components": {
          <String- component name>: <String - Components directory>,
          ... more components
        },
        "tests": {
          <String- test name>: <String - Test directory>,
          ... more tests
        },
        "workers": {
          <String- worker name>: <String - Worker directory>,
          ... more workers
        }
      },
      ... more libraries
    },
    "primitives": {
      <String - name of primitive>: <String - Primitive directory>,
      ... more primitives
    },
    "components": {
      <String- name of component>: <String - directory of library>,
      ... more components
    },
    (Continued on next page)
  }
  (Continued from previous page)
  "platforms": {
    "rcc": {
      <String- name of rcc platform>: <String - directory of rcc
platform>,
      ... more rcc platforms
    },
    "hdl": {
      <String- name of rcc platform>: <String - directory of hdl
platforms>,
      ... more hdl platforms
    }
  }
}
```

18.1.4 *ocpidev show tests --local-scope --json*

The `--local-scope` option will act upon the project that the command is executed in so this command shows all the tests in the specified project.

```
project": {
  "libraries": {
    <String-Library name>: {
      "directory": String - Library directory>,
      "package": String - Library package-id>,
      "tests": {
        <String - Test name>:<String - Test directory>,
        ... more tests
      }
    }
    ... more libraries
  }
}
```

18.1.5 *ocpidev show platforms/registry --json*

```
{
  "rcc": {
    <String-platform name>: {
      "target": <String-target associated with this platform>,
      "package_id": <String - platform package-id>,
      "directory": <String - platform directory>
    },
    ... more rcc platforms
  },
  "hdl": {
    "xsim": {
      "target": <String - target associated with this platform>,
      "package_id": <String - platform package-id>,
      "directory": <String - platform directory>,
      "part": <String-part associated with platform>,
      "built": <Boolean - if this platform is built>,
      "tool": <String - FPGA tool associated with this platform>,
      "vendor": <String - FPGA vendor associated with this platform>
    },
    ... more hdl platforms
  }
}
```

18.1.6 *ocpidev show targets --json*

```
{
  "rcc": {
    <String - platform name>: {
      "target": <String-target name>
    },
    ... more rcc targets
  },
  "hdl": {
    <String - vendor name>: {
      <String - target name>: {
        "tool": <String - FPGA tool associated with this target>,
        "parts": <Array of Strings - FPGA parts associated with this
                    target>
      },
      ... more hdl targets
    },
    ... more hdl vendors
  }
}
```

18.1.7 *ocpidev show component --json*

```
{
  "directory": <String - directory of component>,
  "name": <String - name of component>,
  "package_id": <String - package-id of component>,
  "ports": {
    <String- name of port>: {
      "protocol": <String - name of protocol for port>,
      "producer": <String - "0" or "1" for direction of port>
    },
    ... more ports
  },
  "properties": {
    <String - name of property>: {
      "accessibility": {
        "volatile": <String - "0" or "1">,
        "readback": <String - "0" or "1">,
        "parameter": <String - "0" or "1">,
        "initial": <String - "0" or "1">,
        "writable": <String - "0" or "1">,
        "padding": <String - "0" or "1">
      },
      "name": <String - name of property>
      "type": <String - data type of property>
    },
    ... more properties
  }
}
```

18.1.8 *ocpidev show worker --json*

```
{
  "directory": <String - directory of worker>,
  "name": <String - name of worker>,
  "package_id": <String - package-id of worker>,
  "ports": {
    <String- name of port>: {
      "protocol": <String - name of protocol for port>,
      "producer": <String - "0" or "1" for direction of port>
    },
    ... more ports
  },
  "properties": {
    <String - name of property>: {
      "accessibility": {
        "volatile": <String - "0" or "1">,
        "readback": <String - "0" or "1">,
        "parameter": <String - "0" or "1">,
        "initial": <String - "0" or "1">,
        "writable": <String - "0" or "1">,
        "padding": <String - "0" or "1">
      },
      "name": <String - name of property>
      "type": <String - Datatype of property>
    },
    ... more properties
  }
  "slaves": {
    <String - name of slave worker>: {
      "name": <String - name of slave worker>
    }
    ... more slaves
  }
}
```

18.1.9 *ocpidev show registry/projects --json*

```
{
  "registry_location": <String - directory of registry>,
  "projects": {
    <String - package-id of project>: {
      "exists": <Boolean - if project exists on the file system >,
      "real_path": <String - directory of project>
    },
    ... more projects
  }
}
```

18.1.10 *ocpidev show workers --global-scope --json*

```
{
  "projects": {
    "ocpi.assets": {
      "package_id": <String - package-id of project>,
      "directory": <String - directory of project>,
      "libraries": {
        <String - package-id of library>: {
          "package_id": <String - package-id of library>,
          "directory": <String - directory of library>,
          "workers": {
            <String - name of worker>:<String - directory of
                                worker>,
            ... more workers
          }
          ... more libraries
        },
      },
      ... more projects
    }
  }
```

18.1.11 *ocpidev show components --global-scope*

```
{
  "projects": {
    "ocpi.assets": {
      "package_id": <String - package-id of project>,
      "directory": <String - directory of project>,
      "components": {
        <String - name of component>: <String - directory of
                                    component>,
        ... more components
      }
      "libraries": {
        <String - package-id of library>: {
          "package_id": <String - package-id of library>,
          "directory": <String - directory of library>,
          "components": {
            <String - name of component>: <String - directory of
                                    component>,
            ... more components
          }
          ... more libraries
        },
      },
      ... more projects
    }
  }
```